

Berkeley DB

Berkeley DB Installation and Build Guide

Version 2026.09.7

Copyright (c) 1990, 2013 Oracle and/or its affiliates. All rights reserved.

Preface

Table of Contents

[Conventions Used in this Book](#)

[For More Information](#)

[Contact Us](#)

Welcome to Berkeley DB (DB). This document describes how to build, install and upgrade Berkeley DB

This document reflects Berkeley DB 11g Release 2, which provides DB library version 11.2.5.3.

Conventions Used in this Book

The following typographical conventions are used within in this manual:

Variable or non-literal text is presented in italics. For example: "Go to your *DB_INSTALL* directory."

Program examples are displayed in a monospaced font on a shaded background. For example:

Note

Finally, notes of interest are represented using a note block such as this.

For More Information

[Contact Us](#)

Beyond this manual, you may also find the following sources of information useful when building a DB application:

- [Getting Started with Transaction Processing for C](#)
- [Berkeley DB Getting Started with Replicated Applications for C](#)
- [Berkeley DB C API Reference Guide](#)
- [Berkeley DB C++ API Reference Guide](#)

- [Berkeley DB STL API Reference Guide](#)
- [Berkeley DB TCL API Reference Guide](#)
- [Berkeley DB Programmer's Reference Guide](#)
- [Berkeley DB Upgrade Guide](#)
- [Berkeley DB Getting Started with the SQL APIs](#)

To download the latest documentation along with white papers and other collateral, visit <http://www.oracle.com/technetwork/indexes/documentation/index.html>.

For the latest version of the Oracle downloads, visit <http://www.oracle.com/technetwork/database/berkeleydb/downloads/index.html>.

Contact Us

You can post your comments and questions at the Oracle Technology (OTN) forum for

For sales or support information, email to: berkeleydb-info_us@oracle.com You can subscribe to a low-volume email announcement list for the Berkeley DB product family by sending email to: bdb-join@oss.oracle.com

Chapter 1. Introduction

Table of Contents

[Installation Overview](#)

[Supported Platforms](#)

Welcome to the Berkeley DB. This manual describes how to configure, build and install Berkeley DB. Installation of DB for all of the platforms it officially supports is described in this manual. Upgrade instructions and release notes for newer versions of this product are described here. For information on upgrading from historical versions, see the [Berkeley DB Upgrade Guide](#).

Note that some operating systems and distributions might provide DB, either by default or as part of an installation option. If so, those platforms will have installation instructions for DB specific to them. In this situation, you should see the documentation for your operating system or distribution provider for information on how to get DB on your platform.

Installation Overview

Berkeley DB is an open-source product, and as such it is usually offered in source-code format. This means that placing DB on your platform requires you to configure the build scripts, compile it, and then install the product onto your host system. The exception to this are Microsoft Windows platforms for which a binary installer is available. Note that for Windows platforms, you can still compile the product from source if you desire.

For *nix systems, including the BSD and Linux systems, the usual `configure`, `make` and `make install` installation process is used to place DB on your platform.

For information on building and installing Berkeley DB on:

- Microsoft Windows, see [Building Berkeley DB for Windows](#) or [Building Berkeley DB for Windows Mobile](#) .
- Unix/POSIX — including Linux, BSD, Apple iOS (known as iPhone OS previously), and Mac OS X — see [Building Berkeley DB for UNIX/POSIX](#) .
- VxWorks, see [Building Berkeley DB for VxWorks](#) .

Supported Platforms

You can install Berkeley DB on the following platforms:

- Most versions of Linux (x86-64 and x86) including:
 - Oracle Linux 4, 5, and 6
 - Red Hat
 - Ubuntu
 - Wind River
 - MontaVista Embedded Linux version 6.0
- Oracle Solaris versions 9 and 10 on x86_64, x86, and SPARC.
- FreeBSD
- Microsoft Windows (x86-64 and x86).
 - XP (SP2, SP3)
 - Vista
 - Windows 7

- Server 2008
- Windows Mobile (6.x)
- Apple Mac OS X 10.5 and 10.6.
- IBM AIX version 5 and 6.
- VxWorks 6.x
- QNX Neutrino/POSIX version 6
- Android
- Apple iOS (previously known as iPhone OS)

Apart from those mentioned in the list above, you can install Berkeley DB on most other systems which are POSIX-compliant. When there is a need to run Berkeley DB on a platform that is currently not supported, DB is distributed in source code form that you can use as base source to port Berkeley DB to that platform. For more information on porting to other platforms, see the Berkeley DB Porting Guide.

Chapter 2. System Installation Notes

Table of Contents

[File utility /etc/magic information](#)

[Magic information](#)

[Big-endian magic information](#)

[Little-endian magic information](#)

[Building with multiple versions of Berkeley DB](#)

File utility /etc/magic information

[Magic information](#)

[Big-endian magic information](#)

[Little-endian magic information](#)

The `file(1)` utility is a UNIX utility that examines and classifies files, based on information found in its database of file types, the `/etc/magic` file. The following information may be

added to your system's `/etc/magic` file to enable `file(1)` to correctly identify Berkeley DB database files.

The `file(1)` utility `magic(5)` information for the standard System V UNIX implementation of the `file(1)` utility is included in the Berkeley DB distribution for both big-endian (for example, Sparc) and little-endian (for example, x86) architectures. See [Big-endian magic information](#) and [Little-endian magic information](#) respectively for this information.

The `file(1)` utility `magic(5)` information for Release 3.X of Ian Darwin's implementation of the `file` utility (as distributed by FreeBSD and most Linux distributions) is included in the Berkeley DB distribution. This `magic.txt` information is correct for both big-endian and little-endian architectures. See the next section for this information.

Magic information

```
# Berkeley DB

# Ian Darwin's file /etc/magic files: big/little-endian version.

# Hash 1.85/1.86 databases store metadata in network byte order.
# Btree 1.85/1.86 databases store the metadata in host byte order.
# Hash and Btree 2.X and later databases store the metadata in
# host byte order.

0 long 0x00061561 Berkeley DB
>8 belong 4321
>>4 belong >2 1.86
>>4 belong <3 1.85
>>4 belong >0 (Hash, version %d, native byte-order)
>8 belong 1234
>>4 belong >2 1.86
>>4 belong <3 1.85
>>4 belong >0 (Hash, version %d, little-endian)

0 belong 0x00061561 Berkeley DB
>8 belong 4321
>>4 belong >2 1.86
>>4 belong <3 1.85
>>4 belong >0 (Hash, version %d, big-endian)
>8 belong 1234
>>4 belong >2 1.86
>>4 belong <3 1.85
>>4 belong >0 (Hash, version %d, native byte-order)

0 long 0x00053162 Berkeley DB 1.85/1.86
>4 long >0 (Btree, version %d, native byte-order)
0 belong 0x00053162 Berkeley DB 1.85/1.86
>4 belong >0 (Btree, version %d, big-endian)
0 lelong 0x00053162 Berkeley DB 1.85/1.86
```

```

>4  llong  >0    (Btree, version %d, little-endian)

12 long  0x00061561 Berkeley DB
>16 long  >0    (Hash, version %d, native byte-order)
12 belong 0x00061561 Berkeley DB
>16 belong >0    (Hash, version %d, big-endian)
12 llong  0x00061561 Berkeley DB
>16 llong  >0    (Hash, version %d, little-endian)

12 long  0x00053162 Berkeley DB
>16 long  >0    (Btree, version %d, native byte-order)
12 belong 0x00053162 Berkeley DB
>16 belong >0    (Btree, version %d, big-endian)
12 llong  0x00053162 Berkeley DB
>16 llong  >0    (Btree, version %d, little-endian)

12 long  0x00042253 Berkeley DB
>16 long  >0    (Queue, version %d, native byte-order)
12 belong 0x00042253 Berkeley DB
>16 belong >0    (Queue, version %d, big-endian)
12 llong  0x00042253 Berkeley DB
>16 llong  >0    (Queue, version %d, little-endian)

12 long  0x00040988 Berkeley DB
>16 long  >0    (Log, version %d, native byte-order)
12 belong 0x00040988 Berkeley DB
>16 belong >0    (Log, version %d, big-endian)
12 llong  0x00040988 Berkeley DB
>16 llong  >0    (Log, version %d, little-endian)

```

Big-endian magic information

```

# Berkeley DB

# System V /etc/magic files: big-endian version.

# Hash 1.85/1.86 databases store metadata in network byte order.
# Btree 1.85/1.86 databases store the metadata in host byte order.
# Hash and Btree 2.X and later databases store the metadata in
# host byte order.

0  long  0x00053162 Berkeley DB 1.85/1.86 (Btree,
>4  long  0x00000002 version 2,
>4  long  0x00000003 version 3,
>0  long  0x00053162 native byte-order)

0  long  0x62310500 Berkeley DB 1.85/1.86 (Btree,
>4  long  0x02000000 version 2,
>4  long  0x03000000 version 3,
>0  long  0x62310500 little-endian)

```

```

12 long 0x00053162 Berkeley DB (Btree,
>16 long 0x00000004 version 4,
>16 long 0x00000005 version 5,
>16 long 0x00000006 version 6,
>16 long 0x00000007 version 7,
>16 long 0x00000008 version 8,
>16 long 0x00000009 version 9,
>12 long 0x00053162 native byte-order)

```

```

12 long 0x62310500 Berkeley DB (Btree,
>16 long 0x04000000 version 4,
>16 long 0x05000000 version 5,
>16 long 0x06000000 version 6,
>16 long 0x07000000 version 7,
>16 long 0x08000000 version 8,
>16 long 0x09000000 version 9,
>12 long 0x62310500 little-endian)

```

```

0 long 0x00061561 Berkeley DB
>4 long >2 1.86
>4 long <3 1.85
>0 long 0x00061561 (Hash,
>4 long 2 version 2,
>4 long 3 version 3,
>8 long 0x000004D2 little-endian)
>8 long 0x000010E1 native byte-order)

```

```

12 long 0x00061561 Berkeley DB (Hash,
>16 long 0x00000004 version 4,
>16 long 0x00000005 version 5,
>16 long 0x00000006 version 6,
>16 long 0x00000007 version 7,
>16 long 0x00000008 version 8,
>16 long 0x00000009 version 9,
>12 long 0x00061561 native byte-order)

```

```

12 long 0x61150600 Berkeley DB (Hash,
>16 long 0x04000000 version 4,
>16 long 0x05000000 version 5,
>16 long 0x06000000 version 6,
>16 long 0x07000000 version 7,
>16 long 0x08000000 version 8,
>16 long 0x09000000 version 9,
>12 long 0x61150600 little-endian)

```

```

12 long 0x00042253 Berkeley DB (Queue,
>16 long 0x00000001 version 1,
>16 long 0x00000002 version 2,
>16 long 0x00000003 version 3,
>16 long 0x00000004 version 4,

```

```
>16 long 0x00000005 version 5,
>16 long 0x00000006 version 6,
>16 long 0x00000007 version 7,
>16 long 0x00000008 version 8,
>16 long 0x00000009 version 9,
>12 long 0x00042253 native byte-order)
```

```
12 long 0x53220400 Berkeley DB (Queue,
>16 long 0x01000000 version 1,
>16 long 0x02000000 version 2,
>16 long 0x03000000 version 3,
>16 long 0x04000000 version 4,
>16 long 0x05000000 version 5,
>16 long 0x06000000 version 6,
>16 long 0x07000000 version 7,
>16 long 0x08000000 version 8,
>16 long 0x09000000 version 9,
>12 long 0x53220400 little-endian)
```

```
12 long 0x00040988 Berkeley DB (Log,
>16 long 0x00000001 version 1,
>16 long 0x00000002 version 2,
>16 long 0x00000003 version 3,
>16 long 0x00000004 version 4,
>16 long 0x00000005 version 5,
>16 long 0x00000006 version 6,
>16 long 0x00000007 version 7,
>16 long 0x00000008 version 8,
>16 long 0x00000009 version 9,
>16 long 0x0000000a version 10,
>16 long 0x0000000b version 11,
>16 long 0x0000000c version 12,
>16 long 0x0000000d version 13,
>16 long 0x0000000e version 14,
>16 long 0x0000000f version 15,
>12 long 0x00040988 native byte-order)
```

```
12 long 0x88090400 Berkeley DB (Log,
>16 long 0x01000000 version 1,
>16 long 0x02000000 version 2,
>16 long 0x03000000 version 3,
>16 long 0x04000000 version 4,
>16 long 0x05000000 version 5,
>16 long 0x06000000 version 6,
>16 long 0x07000000 version 7,
>16 long 0x08000000 version 8,
>16 long 0x09000000 version 9,
>16 long 0x0a000000 version 10,
>16 long 0x0b000000 version 11,
>16 long 0x0c000000 version 12,
>16 long 0x0d000000 version 13,
```

```
>16 long  0x0e000000 version 14,
>16 long  0x0f000000 version 15,
>12 long  0x88090400 little-endian)
```

Little-endian magic information

```
# Berkeley DB

# System V /etc/magic files: little-endian version.

# Hash 1.85/1.86 databases store metadata in network byte order.
# Btree 1.85/1.86 databases store the metadata in host byte order.
# Hash and Btree 2.X and later databases store the metadata in
# host byte order.

0 long  0x00053162 Berkeley DB 1.85/1.86 (Btree,
>4 long  0x00000002 version 2,
>4 long  0x00000003 version 3,
>0 long  0x00053162 native byte-order)

0 long  0x62310500 Berkeley DB 1.85/1.86 (Btree,
>4 long  0x02000000 version 2,
>4 long  0x03000000 version 3,
>0 long  0x62310500 big-endian)

12 long  0x00053162 Berkeley DB (Btree,
>16 long  0x00000004 version 4,
>16 long  0x00000005 version 5,
>16 long  0x00000006 version 6,
>16 long  0x00000007 version 7,
>16 long  0x00000008 version 8,
>16 long  0x00000009 version 9,
>12 long  0x00053162 native byte-order)

12 long  0x62310500 Berkeley DB (Btree,
>16 long  0x04000000 version 4,
>16 long  0x05000000 version 5,
>16 long  0x06000000 version 6,
>16 long  0x07000000 version 7,
>16 long  0x08000000 version 8,
>16 long  0x09000000 version 9,
>12 long  0x62310500 big-endian)

0 long  0x61150600 Berkeley DB
>4 long  >0x02000000 1.86
>4 long  <0x03000000 1.85
>0 long  0x00061561 (Hash,
>4 long  0x02000000 version 2,
>4 long  0x03000000 version 3,
>8 long  0xD2040000 native byte-order)
```

```

>8 long 0xE1100000 big-endian)

12 long 0x00061561 Berkeley DB (Hash,
>16 long 0x00000004 version 4,
>16 long 0x00000005 version 5,
>16 long 0x00000006 version 6,
>16 long 0x00000007 version 7,
>16 long 0x00000008 version 8,
>16 long 0x00000009 version 9,
>12 long 0x00061561 native byte-order)

12 long 0x61150600 Berkeley DB (Hash,
>16 long 0x04000000 version 4,
>16 long 0x05000000 version 5,
>16 long 0x06000000 version 6,
>16 long 0x07000000 version 7,
>16 long 0x08000000 version 8,
>16 long 0x09000000 version 9,
>12 long 0x61150600 big-endian)

12 long 0x00042253 Berkeley DB (Queue,
>16 long 0x00000001 version 1,
>16 long 0x00000002 version 2,
>16 long 0x00000003 version 3,
>16 long 0x00000004 version 4,
>16 long 0x00000005 version 5,
>16 long 0x00000006 version 6,
>16 long 0x00000007 version 7,
>16 long 0x00000008 version 8,
>16 long 0x00000009 version 9,
>12 long 0x00042253 native byte-order)

12 long 0x53220400 Berkeley DB (Queue,
>16 long 0x01000000 version 1,
>16 long 0x02000000 version 2,
>16 long 0x03000000 version 3,
>16 long 0x04000000 version 4,
>16 long 0x05000000 version 5,
>16 long 0x06000000 version 6,
>16 long 0x07000000 version 7,
>16 long 0x08000000 version 8,
>16 long 0x09000000 version 9,
>12 long 0x53220400 big-endian)

12 long 0x00040988 Berkeley DB (Log,
>16 long 0x00000001 version 1,
>16 long 0x00000002 version 2,
>16 long 0x00000003 version 3,
>16 long 0x00000004 version 4,
>16 long 0x00000005 version 5,
>16 long 0x00000006 version 6,

```

```

>16 long 0x00000007 version 7,
>16 long 0x00000008 version 8,
>16 long 0x00000009 version 9,
>16 long 0x0000000a version 10,
>16 long 0x0000000b version 11,
>16 long 0x0000000c version 12,
>16 long 0x0000000d version 13,
>16 long 0x0000000e version 14,
>16 long 0x0000000f version 15,
>12 long 0x00040988 native byte-order)

12 long 0x88090400 Berkeley DB (Log,
>16 long 0x01000000 version 1,
>16 long 0x02000000 version 2,
>16 long 0x03000000 version 3,
>16 long 0x04000000 version 4,
>16 long 0x05000000 version 5,
>16 long 0x06000000 version 6,
>16 long 0x07000000 version 7,
>16 long 0x08000000 version 8,
>16 long 0x09000000 version 9,
>16 long 0x0a000000 version 10,
>16 long 0x0b000000 version 11,
>16 long 0x0c000000 version 12,
>16 long 0x0d000000 version 13,
>16 long 0x0e000000 version 14,
>16 long 0x0f000000 version 15,
>12 long 0x88090400 big-endian)

```

Building with multiple versions of Berkeley DB

In some cases it may be necessary to build applications which include multiple versions of Berkeley DB. Examples include applications which include software from other vendors, or applications running on a system where the system C library itself uses Berkeley DB. In such cases, the two versions of Berkeley DB may be incompatible, that is, they may have different external and internal interfaces, and may even have different underlying database formats.

To create a Berkeley DB library whose symbols won't collide with other Berkeley DB libraries (or other application or library modules, for that matter), configure Berkeley DB using the `--with-uniquename=NAME` configuration option, and then build Berkeley DB as usual. (Note that `--with-uniquename=NAME` only affects the Berkeley DB C language library build; loading multiple versions of the C++ or Java APIs will require additional work.) The modified symbol names are hidden from the application in the Berkeley DB header files, that is, there is no need for the application to be aware that it is using a special library build as long as it includes the appropriate Berkeley DB header file.

If "NAME" is not specified when configuring with `--with-uniquename=NAME`, a default value built from the major and minor numbers of the Berkeley DB release will be used. It is rarely necessary to specify NAME; using the major and minor release numbers will ensure that only one copy of the library will be loaded into the application unless two distinct versions really are necessary.

When distributing any library software that uses Berkeley DB, or any software which will be recompiled by users for their systems, we recommend two things: First, include the Berkeley DB release as part of your release. This will insulate your software from potential Berkeley DB API changes as well as simplifying your coding because you will only have to code to a single version of the Berkeley DB API instead of adapting at compile time to whatever version of Berkeley DB happens to be installed on the target system. Second, use `--with-uniquename=NAME` when configuring Berkeley DB, because that will insure that you do not unexpectedly collide with other application code or a library already installed on the target system.

Chapter 3. Debugging Applications

Table of Contents

[Introduction to debugging](#)

[Compile-time configuration](#)

[Run-time error information](#)

[Reviewing Berkeley DB log files](#)

[Augmenting the Log for Debugging](#)

[Extracting Committed Transactions and Transaction Status](#)

[Extracting Transaction Histories](#)

[Extracting File Histories](#)

[Extracting Page Histories](#)

[Other log processing tools](#)

Introduction to debugging

Because Berkeley DB is an embedded library, debugging applications that use Berkeley DB is both harder and easier than debugging a separate server. Debugging can be harder because when a problem arises, it is not always readily apparent whether the problem is

in the application, is in the database library, or is a result of an unexpected interaction between the two. Debugging can be easier because it is easier to track down a problem when you can review a stack trace rather than deciphering interprocess communication messages. This chapter is intended to assist you with debugging applications and reporting bugs to us so that we can provide you with the correct answer or fix as quickly as possible.

When you encounter a problem, there are a few general actions you can take:

Review the Berkeley DB error output:

If an error output mechanism has been configured in the Berkeley DB environment, additional run-time error messages are made available to the applications. If you are not using an environment, it is well worth modifying your application to create one so that you can get more detailed error messages. See [Run-time error information](#) for more information on configuring Berkeley DB to output these error messages.

Review the options available for the `DB_ENV->set_verbose()` method:

Look to see if it offers any additional informational and/or debugging messages that might help you understand the problem.

Add run-time diagnostics:

You can configure and build Berkeley DB to perform run-time diagnostics. (By default, these checks are not done because they can seriously impact performance.) See [Compile-time configuration](#) for more information.

Apply all available patches:

Before reporting a problem in Berkeley DB, please upgrade to the latest Berkeley DB release, if possible, or at least make sure you have applied any updates available for your release from the [Berkeley DB web site](#).

Run the test suite:

If you see repeated failures or failures of simple test cases, run the Berkeley DB test suite to determine whether the distribution of Berkeley DB you are using was built and configured correctly.

Compile-time configuration

There are three compile-time configuration options that assist in debugging Berkeley DB and Berkeley DB applications:

`--enable-debug`

If you want to build Berkeley DB with `-g` as the C and C++ compiler flag, enter `--enable-debug` as an argument to `configure`. This will create Berkeley DB with debugging symbols, as well as load various Berkeley DB routines that can be called directly from a debugger

to display database page content, cursor queues, and so forth. (Note that the -O optimization flag will still be specified. To compile with only the -g, explicitly set the CFLAGS environment variable before configuring.)

`--enable-diagnostic`

If you want to build Berkeley DB with debugging run-time sanity checks and with `DIAGNOSTIC #defined` during compilation, enter `--enable-diagnostic` as an argument to `configure`. This will cause a number of special checks to be performed when Berkeley DB is running. This flag should not be defined when configuring to build production binaries because it degrades performance.

`--enable-umrw`

When compiling Berkeley DB for use in run-time memory consistency checkers (in particular, programs that look for reads and writes of uninitialized memory), use `--enable-umrw` as an argument to `configure`. This guarantees, among other things, that Berkeley DB will completely initialize allocated pages rather than initializing only the minimum necessary amount.

Run-time error information

Normally, when an error occurs in the Berkeley DB library, an integer value (either a Berkeley DB specific value or a system `errno` value) is returned by Berkeley DB. In some cases, however, this value may be insufficient to completely describe the cause of the error, especially during initial application debugging.

Most Berkeley DB errors will result in additional information being written to a standard file descriptor or output stream. Additionally, Berkeley DB can be configured to pass these verbose error messages to an application function. There are four methods intended to provide applications with additional error information: `DB_ENV->set_errcall()`, `DB_ENV->set_errfile()`, `DB_ENV->set_errpfx()` and `DB_ENV->set_verbose()`.

The Berkeley DB error-reporting facilities do not slow performance or significantly increase application size, and may be run during normal operation as well as during debugging. Where possible, we recommend these options always be configured and the output saved in the filesystem. We have found that this often saves time when debugging installation or other system-integration problems.

In addition, there are three methods to assist applications in displaying their own error messages: `db_strerror()`, `DB_ENV->err()`, and `DB_ENV->errx()`. The first is a superset of the ANSI C `strerror` function, and returns a descriptive string for any error return from the Berkeley DB library. The `DB_ENV->err()` and `DB_ENV->errx()` methods use the error message configuration options described previously to format and display error messages to appropriate output devices.

Reviewing Berkeley DB log files

[Augmenting the Log for Debugging](#)

[Extracting Committed Transactions and Transaction Status](#)

[Extracting Transaction Histories](#)

[Extracting File Histories](#)

[Extracting Page Histories](#)

[Other log processing tools](#)

If you are running with transactions and logging, the [db_printlog](#) utility can be a useful debugging aid. The [db_printlog](#) utility will display the contents of your log files in a human readable (and machine-readable) format.

The [db_printlog](#) utility will attempt to display any and all log files present in a designated db_home directory. For each log record, the [db_printlog](#) utility will display a line of the form:

```
[22][28]db_big: rec: 43 txnid 80000963 prevlsn [21][10483281]
```

The opening numbers in square brackets are the log sequence number (LSN) of the log record being displayed. The first number indicates the log file in which the record appears, and the second number indicates the offset in that file of the record.

The first character string identifies the particular log operation being reported. The log records corresponding to particular operations are described following. The rest of the line consists of name/value pairs.

The rec field indicates the record type (this is used to dispatch records in the log to appropriate recovery functions).

The txnid field identifies the transaction for which this record was written. A txnid of 0 means that the record was written outside the context of any transaction. You will see these most frequently for checkpoints.

Finally, the prevlsn contains the LSN of the last record for this transaction. By following prevlsn fields, you can accumulate all the updates for a particular transaction. During normal abort processing, this field is used to quickly access all the records for a particular transaction.

After the initial line identifying the record type, each field of the log record is displayed, one item per line. There are several fields that appear in many different records and a few fields that appear only in some records.

The following table presents each currently written log record type with a brief description of the operation it describes. Any of these record types may have the string "_debug" appended if they were written because [DB_TXN_NOT_DURABLE](#) was specified and the system was configured with [--enable-diagnostic](#).

Log Record Type	Description
bam_adj	Used when we insert/remove an index into/from the page header of a Btree page.
bam_cadjust	Keeps track of record counts in a Btree or Recno database.
bam_cdel	Used to mark a record on a page as deleted.
bam_curadj	Used to adjust a cursor location when a nearby record changes in a Btree database.
bam_merge	Used to merge two Btree database pages during compaction.
bam_pgno	Used to replace a page number in a Btree record.
bam_rcuradj	Used to adjust a cursor location when a nearby record changes in a Recno database.
bam_relink	Fix leaf page prev/next chain when a page is removed.
bam_repl	Describes a replace operation on a record.
bam_root	Describes an assignment of a root page.
bam_rsplite	Describes a reverse page split.
bam_split	Describes a page split.
crdel_inmem_create	Record the creation of an in-memory named database.
crdel_inmem_remove	Record the removal of an in-memory named database.
crdel_inmem_rename	Record the rename of an in-memory named database.
crdel metasub	Describes the creation of a metadata page for a subdatabase.
db_addrem	Add or remove an item from a page of duplicates.
db_big	Add an item to an overflow page (overflow pages contain items too large to place on the main page)
db_cksum	Unable to checksum a page.
db_debug	Log debugging message.
db_noop	This marks an operation that did nothing but update the LSN on a page.
db_ovref	Increment or decrement the reference count for a big item.
db_pg_alloc	Indicates we allocated a page to a database.
db_pg_free	Indicates we freed a page (freed pages are added to a freelist and reused).

Log Record Type	Description
db_pg_freedata	Indicates we freed a page that still contained data entries (freed pages are added to a freelist and reused.)
db_pg_init	Indicates we reinitialized a page during a truncate.
db_pg_sort	Sort the free page list and free pages at the end of the file.
dbreg_register	Records an open of a file (mapping the filename to a log-id that is used in subsequent log operations).
fop_create	Create a file in the file system.
fop_file_remove	Remove a name in the file system.
fop_remove	Remove a file in the file system.
fop_rename	Rename a file in the file system.
fop_write	Write bytes to an object in the file system.
ham_chgpg	Used to adjust a cursor location when a Hash page is removed, and its elements are moved to a different Hash page.
ham_copypage	Used when we empty a bucket page, but there are overflow pages for the bucket; one needs to be copied back into the actual bucket.
ham_curadj	Used to adjust a cursor location when a nearby record changes in a Hash database.
ham_groupalloc	Allocate some number of contiguous pages to the Hash database.
ham_insdel	Insert/delete an item on a Hash page.
ham_metagroup	Update the metadata page to reflect the allocation of a sequence of contiguous pages.
ham_newpage	Adds or removes overflow pages from a Hash bucket.
ham_replace	Handle updates to records that are on the main page.
ham_splitdata	Record the page data for a split.
heap_addrem	Add or remove an entry from a Heap database.
heap_pg_alloc	Indicates we allocated a page to a Heap database.
heap_trunc_meta	Records the truncation of the meta page in a Heap database.
heap_trunc_page	Records the truncation of a data page in a Heap database.
qam_add	Describes the actual addition of a new record to a Queue.
qam_del	Delete a record in a Queue.
qam_delext	Delete a record in a Queue with extents.
qam_incfirst	Increments the record number that refers to the first record in the database.
qam_mvptr	Indicates we changed the reference to either or both of the first and current records in the file.
txn_child	Commit a child transaction.

Log Record Type	Description
txn_ckp	Transaction checkpoint.
txn_recycle	Transaction IDs wrapped.
txn_regop	Logs a regular (non-child) transaction commit.
txn_xa_regop	Logs a prepare message.

Augmenting the Log for Debugging

When debugging applications, it is sometimes useful to log not only the actual operations that modify pages, but also the underlying Berkeley DB functions being executed. This form of logging can add significant bulk to your log, but can permit debugging application errors that are almost impossible to find any other way. To turn on these log messages, specify the `--enable-debug_rop` and `--enable-debug_wop` configuration options when configuring Berkeley DB. See [Configuring Berkeley DB](#) for more information.

Extracting Committed Transactions and Transaction Status

Sometimes, it is helpful to use the human-readable log output to determine which transactions committed and aborted. The awk script, `commit.awk`, (found in the `db_printlog` directory of the Berkeley DB distribution) allows you to do just that. The following command, where `log_output` is the output of `db_printlog`, will display a list of the transaction IDs of all committed transactions found in the log:

```
awk -f commit.awk log_output
```

If you need a complete list of both committed and aborted transactions, then the script `status.awk` will produce it. The syntax is as follows:

```
awk -f status.awk log_output
```

Extracting Transaction Histories

Another useful debugging aid is to print out the complete history of a transaction. The awk script `txn.awk` allows you to do that. The following command line, where `log_output` is the output of the [db_printlog](#) utility and `txnlist` is a comma-separated list of transaction IDs, will display all log records associated with the designated transaction ids:

```
awk -f txn.awk TXN=txnlist log_output
```

Extracting File Histories

The awk script `fileid.awk` allows you to extract all log records that refer to a designated file. The syntax for the `fileid.awk` script is the following, where `log_output` is the output of `db_printlog` and `fids` is a comma-separated list of fileids:

```
awk -f fileid.awk PGNO=fids log_output
```

Extracting Page Histories

The awk script `pgno.awk` allows you to extract all log records that refer to designated page numbers. However, because this script will extract records with the designated page numbers for all files, it is most useful in conjunction with the `fileid` script. The syntax for the `pgno.awk` script is the following, where `log_output` is the output of `db_printlog` and `pgnolist` is a comma-separated list of page numbers:

```
awk -f pgno.awk PGNO=pgnolist log_output
```

Other log processing tools

The awk script `count.awk` prints out the number of log records encountered that belonged to some transaction (that is, the number of log records excluding those for checkpoints and non-transaction-protected operations).

The script `range.awk` will extract a subset of a log. This is useful when the output of [db_printlog](#) utility is too large to be reasonably manipulated with an editor or other tool. The syntax for `range.awk` is the following, where `sf` and `so` represent the LSN of the beginning of the sublog you want to extract, and `ef` and `eo` represent the LSN of the end of the sublog you want to extract:

```
awk -f range.awk START_FILE=sf START_OFFSET=so END_FILE=ef \
END_OFFSET=eo log_output
```

Chapter 4. Building Berkeley DB for Android

Table of Contents

[Building the Drop-In Replacement for Android](#)

[Migrating from SQLite to Berkeley DB](#)

[Building the Android JDBC Driver](#)

Android Configuration Options

Berkeley DB provides support for the Android platform enabling you to develop and deploy a wide range of mobile applications and services. Android provides SQLite as the default database for developing applications that need database support. Berkeley DB SQL API is fully compatible with SQLite and can be used as a replacement. The `build_android` directory in the Berkeley DB distribution contains a makefile, `Android.mk`, for building a drop-in replacement for SQLite.

Oracle offers two different solutions for building the BDB SQL API for Android. The first creates a library that can be used as a drop-in replacement for SQLite on Android. The second creates a JDBC driver for Android.

Building the Drop-In Replacement for Android

Migrating from SQLite to Berkeley DB

This section describes how to build a library that can be used as a drop-in replacement for SQLite on Android.

1. Download and compile the Android source tree.

The compiling process takes time but is a one time activity. For information on downloading and compiling the Android source code, see <http://source.android.com/source/download.html>.

2. Copy the Berkeley DB code into the Android build tree.

```
$ cd <root>/external/sqlite/dist
$ tar xzf ${DB_PATH}
```

where `<root>` is the root of the Android source tree and `${DB_PATH}` is the path where you saved the `db-xx.tar.gz` version of the Berkeley DB distribution.

3. Update the Android build file to identify Berkeley DB.

Replace the `Android.mk` file with the one from the Berkeley DB source tree by doing the following:

```
$ cd <root>/external/sqlite/dist
$ mv Android.mk Android.mk.sqlite
$ cp ${DB_INSTALL}/build_android/Android.mk ./
```

where `${DB_INSTALL}` is the directory into which you installed the Berkeley DB library.

4. Tuning parameters.

The configuration options for performance tuning can be added/edited in the `Android.mk` file by modifying `LOCAL_CFLAGS` located in the `build/libsqlite` replacement section. For more information, see [Android Configuration Options](#).

It is also possible to change these settings using `PRAGMA` commands or through the `DB_CONFIG` file.

5. Build the new Android image.

To build the Android image with Berkeley DB SQL included, do the following:

```
$ cd <root>
$ . build/envsetup.sh
$ make clean-libsqlite
$ mmm -B external/sqlite/dist
$ make snod
```

You can locate the new image in `<root>/out/target/product/generic`.

Migrating from SQLite to Berkeley DB

This section describes how to enable automatic conversion of SQLite format databases to Berkeley DB SQL when they are opened. To do this, you must first make sure that the `-DBDBSQL_CONVERT_SQLITE` option is added to `LOCAL_CFLAGS` when you configure your Berkeley DB database build.

1. Build a static SQLite shell for Android platform.

Create a script, `build_sqlite3_shell.sh`, in the `<root>/external/sqlite/dist` directory.

```
#!/bin/bash
# This script shows how to use built-in toolchain to build
# sqlite3 shell, which is required by Berkeley DB SQL
# on-the-fly migration feature.

# Note: these variables should be set per active Android source tree
# We assume $PWD=$ROOT/external/sqlite/dist
ROOT=${PWD}/../..
TOOLCHAIN=${ROOT}/prebuilt/linux-x86/toolchain/arm-eabi-4.4.0
CC=${TOOLCHAIN}/bin/arm-eabi-gcc
LIB="${ROOT}/out/target/product/generic/obj/lib"
INCLUDE="${ROOT}/ndk/build/platforms/android-8/arch-arm/usr/include"

# CFLAGS should be set per Android.mk.sqlite (the original
# version of SQLite's Android.mk)
CFLAGS="-DHAVE_USLEEP=1 -DSQLITE_THREADSAFE=1 -DNDEBUG=1 \
```

```
-DSQLITE_DEFAULT_JOURNAL_SIZE_LIMIT=1048576 \
-DSQLITE_ENABLE_MEMORY_MANAGEMENT=1 \
-DSQLITE_DEFAULT_AUTOVACUUM=1 \
-DSQLITE_TEMP_STORE=3 -DSQLITE_ENABLE_FTS3 \
-DSQLITE_ENABLE_FTS3_BACKWARDS -DTHREADSAFE=1"
CFLAGS="${CFLAGS} -I${INCLUDE}"

LDFLAGS="-ldl -nostdlib -Wl,--gc-sections -lc -llog -lgcc \
-Wl,--no-undefined,-z,nocopyreloc ${LIB}/crtend_android.o \
${LIB}/crtbegin_dynamic.o -L${LIB} -Wl,-rpath,${LIB}"

${CC} -DANDROID -DOS_ANDROID --sysroot="${SYSROOT}" -mandroid \
-fvisibility=hidden -ffunction-sections -fdata-sections \
-fPIC ${LDFLAGS} ${CFLAGS} \
sqlite3.c shell.c -o sqlite3orig
```

Ensure you adjust the variables as per your actual Android environment. This script is suited for Android 2.2.

2. Execute the `build_sqlite3_shell.sh` script and to get the static sqlite3 shell utility - `sqlite3orig`.
3. Change the system image file.

Use the `xyaffs2` utility to decompress the `system.img` and get the directory `system`.

```
$ xyaffs2 ./system.img system
```

Add static sqlite3 shell utility.

```
$ cp <root>/external/sqlite/dist/sqlite3orig \
system/xbin/sqlite3orig
```

Use the `mkyaffs2image` utility to rebuild `system.img` from the changed directory `system`.

```
$ mkyaffs2image -f $PWD/system system.img
```

Note

To open the database in the `SQLite` format use the `sqlite3orig` command.

Building the Android JDBC Driver

This section describes how to build and use the BDB JDBC driver for Android. Note that the BDB JDBC driver cannot currently be built on a Windows platform.

1. Download and install the Android SDK. The installation instructions can be found here:

<http://developer.android.com/sdk/installing.html>

2. Download and install the Android NDK. It can be found here:

<http://developer.android.com/sdk/ndk/index.html>

3. Build the BDB JDBC libraries.

1. If you do not already have it, download the Berkeley DB package from here:

<http://www.oracle.com/technetwork/database/berkeleydb/downloads/index.html>

Note that you must be using a 5.3.x or higher version of the product in order for these instructions to work. Once you have the package, unpack it:

```
$ tar zxvf db-x.y.z.tar.gz
$ cd db-x.y.z
```

Where `x.y.z` the major, minor, and point release version of the Berkeley DB distribution which you are using.

Also, note that in the following instructions, the directory denoted by `db-x.y.z`, above, is referred to as `<db>`.

2. Build an x86/x64 JDBC package. This is required because the building process will generate target files which are required to build Android NDK. Also, the built JAR file can be imported by eclipse, which will then convert it to the Android Dalvik JAR format.

To do this, edit `<db>/lang/sql/jdbc/SQLite/Database.java` and replace all instances of `System.loadLibrary("sqlite_jni")` with `System.loadLibrary("oracle-jdbc")`.

Once you have done this, configure and make the library. The following example shows the minimum configuration options that you need to use in order to configure the Berkeley DB JDBC driver. For your particular installation, other configuration options might be interesting to you. See [Configuring Berkeley DB](#) and [Android Configuration Options](#) for more information.

```
cd <db>/build_unix
../dist/configure --enable-jdbc && make
```

4. Build the Android NDK:

```
$ cd <db>/build_android/jdbc/jni
$ <ndk-install-directory>/ndk-build
```

This results in the following required files:

<db>/build_android/jdbc/libs/armeabi/liboracle-jdbc.so
<db>/build_android/jdbc/libs/armeabi/dbsql
<db>/build_unix/jdbc/sqlite.jar

Having built the JDBC driver, you can now use it with your project. You can do this using Eclipse and the ADT plugin, which you can get from here:

<http://developer.android.com/sdk/eclipse-adt.html>

To make sure everything is working:

1. Start Eclipse and create an Android project. Use:
 - `test_jdbc` as the Android project name.
 - Create it as an Android 3.2 project.
 - For the package name, use `example.jdbc`.
2. This results in an empty code file. Copy and paste the following example code into that file:

```
package example.testjdbc;

import SQLite.*;
import java.io.*;

import android.app.Activity;
import android.os.Bundle;
import android.widget.TextView;
import java.sql.*;

public class Test_jdbcActivity extends Activity {

    /*
     * This is the main entrance/body of our sample program. This
     * example illustrates all of the major API usages.
```

```

*/
public void onCreate(Bundle savedInstanceState) {
    super.onCreate(savedInstanceState);

    TextView tv = new TextView(this);
    tv.setText("App Started");
    setContentView(tv);

    System.out.println("Appstart: ");

    String url =
        "jdbc:sqlite://data/data/example.testjdbc/example.db";
    Connection con;
    String dropString = "drop table if exists COFFEES";
    String createString;
    createString = "create table COFFEES "
        + "(COF_NAME varchar(32), "
        + "SUP_ID int, " + "PRICE float, " + "SALES int, "
        + "TOTAL int)";
    String insertString = "drop table COFFEES if exists";
    String query = "select COF_NAME, PRICE from COFFEES";
    Statement stmt;

    try {
        Class.forName("SQLite.JDBCdriver");

    } catch (java.lang.ClassNotFoundException e) {
        System.err.print("ClassNotFoundException: ");
        System.err.println(e.getMessage());
    }

    try {
        con =
            DriverManager.getConnection(url, "myLogin", "myPW");

        stmt = con.createStatement();
        stmt.executeUpdate(dropString);
        stmt.executeUpdate(createString);
        stmt.close();

        stmt = con.createStatement();
        stmt.executeUpdate("insert into COFFEES "
            + "values('Colombian', 00101, 7.99, 0, 0)");

        stmt.executeUpdate("insert into COFFEES "
            + "values('French_Roast', 00049, 8.99, 0, 0)");

        stmt.executeUpdate("insert into COFFEES "
            + "values('Espresso', 00150, 9.99, 0, 0)");

        stmt.executeUpdate("insert into COFFEES "

```

```

        + "values('Colombian_Decaf', 00101, 8.99, 0, 0)");

stmt.executeUpdate("insert into COFFEES "
    + "values('French_Roast_Decaf', 00049, 9.99, 0, 0)");

ResultSet rs = stmt.executeQuery(query);

System.out.println("Coffee Break Coffees and Prices:");
while (rs.next()) {
    String s = rs.getString("COF_NAME");
    float f = rs.getFloat("PRICE");
    System.out.println(s + " " + f);
}
stmt.close();
con.close();

} catch (SQLException ex) {
    System.err.println("SQLException: " + ex.getMessage());
}
}
}

```

3. Copy the following files into place:

```

$ cd <workspace>/test_jdbc
$ mkdir -p libs/armeabi
$ cp -r <db>/build_android/jdbc/libs/armeabi/liboracle-jdbc.so \
libs/armeabi
$ cp -r <db>/build_unix/jdbc/sqlite.jar libs

```

4. Back in Eclipse, right click the project name, and select the `refresh` option to reload the project from the directory. The two new files that were copied into place in the previous step are now included in the project view.

5. Convert the JAR file to the Android Dalvik format:

1. Right-click on your project.
2. Choose `Build Path -> Configure Build Path`
3. Click the `Libraries` tab.
4. Click `Add JARS`.

6. Run the project:

1. Choose `Property -> Android` and select any one of the usable build targets.
2. Right click the project. Choose `Run As -> Android`

7. Verify your installation. After a short pause (depending on the speed of your system), the application logo is displayed. Use the Android adb command line application to make sure the application is running as expected:

```
$ cd <android-sdk>/platform-tools
$ ./adb logcat
I/System.out( 539): Appstart:
I/System.out( 539): Coffee Break Coffees and Prices:
I/System.out( 539): Colombian 7.99
I/System.out( 539): French_Roast 8.99
I/System.out( 539): Espresso 9.99
I/System.out( 539): Colombian_Decaf 8.99
I/System.out( 539): French_Roast_Decaf 9.99
```

You can also check if the database (example.db) exists in the emulator:

```
$ ./adb shell ls /data/data/example.testjdbc
example.db
example.db-journal
lib
```

Finally, check the database using the BDB SQL shell:

```
$ ./adb push <db>/build_android/jdbc/libs/armeabi/dbsql \
/data/data/example.testjdbc
326 KB/s (1293760 bytes in 3.865s)
$ ./adb shell
root@android:/ # cd /data/data/example.testjdbc
root@android:/data/data/example.testjdbc # ./dbsql example.db
Berkeley DB 11g Release 2, library version 11.2.5.2.36
Enter ".help" for instructions
Enter SQL statements terminated with a ";"
dbsql> .tables
COFFEES
dbsql> select * from COFFEES;
Colombian|101|7.99|0|0
French_Roast|49|8.99|0|0
Espresso|150|9.99|0|0
Colombian_Decaf|101|8.99|0|0
French_Roast_Decaf|49|9.99|0|0
dbsql> .quit
```

Android Configuration Options

There are several configuration options you can specify in `LOCAL_CFLAGS` located in the `Android.mk` file.

- `BDBSQL_CONVERT_SQLITE`

This option enables to convert SQLite database to BDB SQL database format. See [Migrating from SQLite to Berkeley DB](#) for more information.

- `BDBSQL_SHARE_PRIVATE`

This flag is enabled by default and keeps all the region files in memory instead of the disk. This flag also implements database-level locking.

- `SQLITE_DEFAULT_CACHE_SIZE`

SQLite provides an in-memory cache which you size according to the maximum number of database pages that you want to hold in memory at any given time. See [Changing Compile Options](#).

- `SQLITE_DEFAULT_JOURNAL_SIZE_LIMIT`

For SQLite, this pragma identifies the maximum size that the journal file is allowed to be. Berkeley DB does not have a journal file, but it writes and uses log files. A new log file is created when the current log file has reached the defined maximum size. This flag defines this maximum size for a log file. Default value is 10 MB for Berkeley DB SQL interface.

Hard-coded numbers in the build can be adjusted using the following SQLite PRAGMA commands:

- `PRAGMA cache_size`
- `PRAGMA journal_size_limit`

You can configure most aspects of your Berkeley DB environment by using the [DB_CONFIG](#) file.

Chapter 5. Building Berkeley DB for Windows

Table of Contents

[Building Berkeley DB for 32 bit Windows](#)

[Visual C++ .NET 2010](#)

[Visual C++ .NET 2008](#)

[Visual C++ .NET 2005](#)

[Build results](#)

[Building Berkeley DB for 64-bit Windows](#)

[x64 build with Visual Studio 2005 or newer](#)

[Building Berkeley DB for Windows on ARM64](#)

[Building Berkeley DB with Cygwin](#)

[Building the C++ API](#)

[Building the C++ STL API](#)

[Building the Java API](#)

[Building the SQL API](#)

[Binary Compatibility With SQLite](#)

[Setting Preprocessor Flags](#)

[Enabling Extensions](#)

[Disabling Log Checksums](#)

[Building the JDBC Driver](#)

[Using the JDBC Driver](#)

[Building the ODBC Driver](#)

[Using the ADO.NET Driver](#)

[Building the Tcl API](#)

[Distributing DLLs](#)

[Additional build options](#)

[Building a small memory footprint library](#)

[Running the test suite under Windows](#)

[Building the software needed by the tests](#)

[Running the test suite under Windows](#)

[Building the software needed by the SQL tests](#)

[Windows notes](#)

[Windows FAQ](#)

This chapter contains general instructions on building Berkeley DB for specific windows platforms using specific compilers. The [Windows FAQ](#) also contains helpful information.

The `build_windows` directory in the Berkeley DB distribution contains project files for Microsoft Visual Studio:

Project File	Description
Berkeley_DB.sln	Visual Studio 2005 (8.0) workspace
*.vcproj	Visual Studio 2005 (8.0) projects
Berkeley_DB_vs2010.sln	Visual Studio 2010 workspace
*.vcxproj	Visual Studio 2010 projects

These project files can be used to build Berkeley DB for the following platforms: Windows NT/2K/XP/2003/Vista and Windows7; and 64-bit Windows XP/2003/Vista and Windows7.

Building Berkeley DB for 32 bit Windows

[Visual C++ .NET 2010](#)

[Visual C++ .NET 2008](#)

[Visual C++ .NET 2005](#)

[Build results](#)

Visual C++ .NET 2010

1. Choose File -> Open -> Project/Solution.... In the `build_windows` directory, select `Berkeley_DB_vs2010.sln` and click Open.
2. Choose the desired project configuration from the drop-down menu on the tool bar (either Debug or Release).
3. Choose the desired platform configuration from the drop-down menu on the tool bar (usually Win32 or x64).
4. To build, right-click on the `Berkeley_DB_vs2010` solution and select Build Solution.

Visual C++ .NET 2008

1. Choose File -> Open -> Project/Solution.... In the `build_windows` directory, select `Berkeley_DB.sln` and click Open.
2. The Visual Studio Conversion Wizard will open automatically. Click the Finish button.
3. On the next screen click the Close button.
4. Choose the desired project configuration from the drop-down menu on the tool bar (either Debug or Release).
5. Choose the desired platform configuration from the drop-down menu on the tool bar (usually Win32 or x64).
6. To build, right-click on the Berkeley_DB solution and select Build Solution.

Visual C++ .NET 2005

1. Choose File -> Open -> Project/Solution.... In the `build_windows` directory, select `Berkeley_DB.sln` and click Open
2. Choose the desired project configuration from the drop-down menu on the tool bar (either Debug or Release).
3. Choose the desired platform configuration from the drop-down menu on the tool bar (usually Win32 or x64).
4. To build, right-click on the Berkeley_DB solution and select Build Solution.

Build results

The results of your build will be placed in one of the following Berkeley DB subdirectories, depending on the configuration that you chose:

<code>build_windows\Win32\Debug</code>
<code>build_windows\Win32\Release</code>
<code>build_windows\Win32\Debug_static</code>
<code>build_windows\Win32\Release_static</code>

When building your application during development, you should normally use compile options "Debug Multithreaded DLL" and link against `build_windows\Debug\libdb53d.lib`. You can also build using a release version of the Berkeley DB libraries and tools, which will be placed in `build_windows\Win32\Release\libdb53.lib`. When linking against the release build, you should compile your code with the "Release Multithreaded DLL" compile option. You will also need to add the `build_windows` directory to the list of include directories of your application's project, or copy the Berkeley DB include files to another location.

Building Berkeley DB for 64-bit Windows

[x64 build with Visual Studio 2005 or newer](#)

The following procedure can be used to build natively on a 64-bit system or to cross-compile from a 32-bit system.

When building 64-bit binaries, the output directory will be one of the following Berkeley DB subdirectories, depending upon the configuration that you chose:

build_windows\x64\Debug
build_windows\x64\Release
build_windows\x64\Debug_static
build_windows\x64\Release_static

x64 build with Visual Studio 2005 or newer

1. Follow the build instructions for your version of Visual Studio, as described in [Building Berkeley DB for 32 bit Windows](#).
2. Select x64 from the Platform Configuration dropdown.
3. Right click on Solution 'Berkeley_DB' in the Solution Explorer, and select Build Solution

Building Berkeley DB with Cygwin

To build Berkeley DB with Cygwin, follow the instructions in [Building for UNIX/POSIX](#).

Building the C++ API

C++ support is built automatically on Windows.

Building the C++ STL API

In the project list of the Berkeley_DB.sln solution, build the "db_stl" project and "db_stl_static" project to build STL API as a dynamic or static library respectively. And in your application, you should link this library file as well as the Berkeley DB library file to your application. The STL API library file is by default always put at the same location as the Berkeley DB library file.

And you need to include the STL API header files in your application code. If you are using the Berkeley DB source tree, the header files are in <Berkeley DB Source Root >/stl

directory; If you are using the pre-built installed version, these header files are in < Berkeley DB Installed Directory>/include, as well as the db.h and db_cxx.h header files.

Building the Java API

Java support is not built automatically. The following instructions assume that you have installed the Sun Java Development Kit in d:\java. Of course, if you installed elsewhere or have different Java software, you will need to adjust the pathnames accordingly.

1. Set your include directories.

- In Visual Studio 2005/Visual Studio 2008 - Choose Tools -> Options -> Projects -> VC++ Directories. Under the "Show directories for" pull-down, select "Include files". Add the full pathnames for the d:\java\include and d:\java\include\win32 directories. Then click OK.
- In Visual Studio 2010 - Right-click db_java project, choose Properties->Configuration Properties-> VC++ Directories->Include Directories. Add the full pathnames for the d:\java\include and d:\java\include\win32 directories. Then click OK.

These are the directories needed when including jni.h.

2. Set the executable files directories.

- In Visual Studio 2005/Visual Studio 2008 - Choose Tools -> Options -> Projects -> VC++ Directories. Under the "Show directories for" pull-down, select "Executable files". Add the full pathname for the d:\java\bin directory, then click OK.
- In Visual Studio 2010 - Right-click db_java project, choose Properties->Configuration Properties-> VC++ Directories->Executable Directories. Add the full pathnames for the d:\java\bin directories. Then click OK.

This is the directory needed to find javac.

3. Set the build type to Release or Debug in the drop-down on the tool bar.

4. To build, right-click on db_java and select Build. This builds the Java support library for Berkeley DB and compiles all the java files, placing the resulting db.jar and dbexamples.jar files in one of the following Berkeley DB subdirectories, depending on the configuration that you chose:

build_windows\Win32\Debug
build_windows\Win32\Release

Building the SQL API

[Binary Compatibility With SQLite](#)

[Setting Preprocessor Flags](#)

[Enabling Extensions](#)

[Disabling Log Checksums](#)

[Building the JDBC Driver](#)

[Using the JDBC Driver](#)

[Building the ODBC Driver](#)

[Using the ADO.NET Driver](#)

SQL support is built as part of the default build on Windows. For information on the build instructions, see [Building Berkeley DB for Windows](#) .

The SQL library is built as `libdb_sql53.dll` in the Release mode or `libdb_sql53d.dll` in the Debug mode. An SQL command line interpreter called `dbsql.exe` is also built.

Binary Compatibility With SQLite

`libdb_sql53.dll` is compatible with `sqlite3.dll`. You can copy `libdb_sql53.dll` to `sqlite3.dll` and `dbsql.exe` to `sqlite3.exe`, and use these applications as a replacement for the standard SQLite binaries with the same names. However, if you want to do this, then any legacy data in use by those tools must be migrated from the standard SQLite database to a Berkeley DB SQL database before you replace the standard SQLite dll and executable with the Berkeley DB equivalent. For information on migrating data from standard SQLite databases to a Berkeley DB SQL database, see the Berkeley DB Getting Started with the SQL APIs guide.

Warning

Rename your dlls and executables to the standard SQLite names with extreme care. Doing this will cause all existing tools to break that currently have data stored in a standard SQLite database.

For best results, rename your dlls and command line tool to use the standard SQLite names only if you know there are no other tools on your production platform that rely on standard SQLite.

Setting Preprocessor Flags

By default, Berkeley DB SQL generates each table as a subdatabase in a single file. To generate each table in a separate file, specify `BDBSQL_FILE_PER_TABLE` in Preprocessor Definitions of the `db_sql` project.

When this option is enabled, the SQL database name is used as a directory name. This directory contains one file for the metadata and one file each for every table created by the SQL API. Do not add or delete files from the database directory. Adding or deleting files may corrupt the database. To backup just the metadata (schema), make a copy of the `metadata` and `table00001` files from the database directory. Make a new copy whenever the schema is changed.

Enabling Extensions

The Berkeley DB SQL API provides extensions such as full text search and R-Tree index. To enable these extensions, do the following:

1. Open the Berkeley DB solution in Visual Studio.
2. Specify `SQLITE_ENABLE_FTS3` or `SQLITE_ENABLE_RTREE` in Preprocessor Definitions of the `db_sql` project.
3. Re-build the `db_sql` project.

See the SQLite Documentation for more information on [full text search](#) and [R-Tree](#).

Disabling Log Checksums

You can disable checksums in log records so as to provide a boost to database performance. However, this comes at the risk of having undetectable log file corruption that could prevent data recovery in the event of database corruption.

Note

Note that disabling log record checksums is meant to only be used with the SQL interface. However, disabling checksums for the SQL interface also disables checksums for the non-SQL interfaces.

To disable log checksums, before you build the library edit the `build_windows/db_config.h` file, and delete the following line:

```
#define HAVE_LOG_CHECKSUM 1
```

Building the JDBC Driver

This section describes the steps to build the JDBC driver.

1. Configure your build environment. For information on how to configure to build Java applications, see [Building the Java API](#).
2. Build the SQL project in Debug mode.
3. Open Berkeley_DB.sln or Berkeley_DB_vs2010.sln in Visual Studio. If you are using Java 1.6, do the following:
 - In the Solution Explorer, right-click the `db_sql_jdbc` project and select properties.
 - In the Configuration Properties -> Build Events -> Pre-Build Event section, alter the command to refer to `JDBC2z` instead of `JDBC2x`.

If you are using Java 1.7, do the following:

- In the Solution Explorer, right-click the `db_sql_jdbc` project and select properties.
 - In the Configuration Properties -> Build Events -> Pre-Build Event section, alter the command to refer to `JDBC2z1` instead of `JDBC2x`. Also, remove the option of `"-target 1.5"`.
 - Go to `db\lang\sql\jdbc\SQLite`, and replace `JDBCDriver.java` with `JDBCDriver.java17`, and replace `JBCDataSource.java` with `JBCDataSource.java17`.
4. Build the `db_sql_jdbc` project in Visual Studio.

You can test the build by entering the following commands from the `db\build_windows\Win32\Debug` directory:

```
javac -cp ".;jdbc.jar" -d .\..\..\sql\jdbc\test3.java
java -cp ".;jdbc.jar" test3
```

Using the JDBC Driver

This section describes the steps to download, build, and run sample programs using the built JDBC driver.

Downloading JDBC Sample Code

The download link for JDBC sample code is available on the [Oracle Technology Network \(OTN\)](#). You can identify the link by the "JDBC programming examples from all three editions (ZIP format)" text beside it.

Modifying Sample Code

Before running the sample code, do the following:

1. Unzip the file containing the sample code to a new directory (for example, jdbc_ex).
2. Substitute `jdbc:sqlite:<db-file-name>` for the generic JDBC URL that appears in the code. That is, put `jdbc:sqlite:<db-file-name>` between the quotation marks in the line:

```
String url = "jdbc:mySubprotocol:myDataSource";
```

Note: The `<db-file-name>` can either be an absolute path name like `"D:\jdbc_ex_db\myDataSource"`, or a relative path-file-name like `"..\jdbc_ex_db\myDataSource"`, or a file name, like `"myDataSource"`, in which the database file will be stored at the current directory.

3. Substitute `SQLite.JDBC.Driver` for `myDriver.ClassName` in the line:

```
Class.forName("myDriver.ClassName");
```
4. Substitute the username and password you use for your database in the following:

```
"myLogin", "myPassword".
```

This is optional.

5. If your JDK version is above 1.5, change the variable name `enum` in `OutputApplet.java` to some other variable name because, as of JDK release 5 `enum` is a keyword and can not be used as an identifier.

Building and Running the JDBC Sample code

See [Building the JDBC Driver](#) for instructions about building JDBC driver.

To build and run the JDBC examples do the following:

1. In the `db\build_windows\Win32\Debug` directory, run following commands:

```
$ javac -classpath ".;jdbc.jar" -d . \path\to\jdbc_ex\*.java
$ java -classpath ".;jdbc.jar" <ClassName, eg. CreateCoffees>
```

2. After you run the `CreateCoffees` example, use the `dbsql` executable to open the `myDataSource` database file and check if the table `COFFEES` has been successfully created in the database.

```
$ dbsql myDataSourcedbsql> .tables
COFFEES
```

```

dbsql> .dump
PRAGMA foreign_keys=OFF;
BEGIN TRANSACTION;
CREATE TABLE COFFEES (COF_NAME varchar(32),/
SUP_ID int, PRICE float, SALES int, TOTAL int);
COMMIT;
dbsql>

```

3. Repeat step 2 to run other examples.

Note: Some examples like AutoGenKeys are not yet supported by BDB JDBC driver. The `SQLException` is displayed for those unsupported examples.

Building the ODBC Driver

This section describes the steps required to build the ODBC driver.

Configuring Your System

To configure your system prior to building the ODBC driver, do the following:

1. Download and install the latest SQLite ODBC driver Windows installer package for [32 bit Windows](#) or [64 bit Windows](#).
2. Download and install the latest [Microsoft Data Access Components \(MDAC\) SDK](#) . The MDAC SDK is only required for testing the installation.

Building the Library

1. Build the SQL project in Release mode. See [Building the SQL API](#).
2. Open Visual Studio.
3. Load the Berkeley_DB solution file into Visual Studio.
4. Set the build target to Release
5. Build the solution.
6. Select File -> Add -> Existing Project.
7. Select `build_windows` .
8. Select the appropriate directory for your compiler: `vs8` or `vs10` .
9. Select `db_sql_odbc.vcproj` and add it to the Berkeley_DB solution. This adds the `db_sql_odbc` Visual Studio project to the Berkeley_DB solution file.
10. Build the `db_sql_odbc` project. This can be done by right-clicking the `db_sql_odbc` project in the project explorer panel, and selecting `build` .

The `sqlite3odbc.dll` , `libdb_sql53.dll` and `libdb53.dll` files are now built.

Installing the Library

Copy the dll files built in the Building the Library section to the Windows system folder.

The Windows system folder is different on different systems, but is often `C:\WINDOWS\System32`.

Testing the ODBC Install

The steps to verify that the installed driver works are as follows:

1. Open the Unicode ODBCTest application. On Windows XP: Windows start -> Microsoft Data Access SDK 2.8 -> ODBCTest (Unicode, x86).
2. Select the Conn -> Full Connect... menu item.
3. Select SQLite3 Datasource and click OK.
4. Select the Stmt -> SQLExecDirect... menu item.
5. Enter `CREATE TABLE t1(x);` in the Statement text box and click OK.
6. Verify that no error messages were output to the error window.

Using the ADO.NET Driver

Go to the [Oracle Berkeley DB download page](#), and download the ADO.NET package. Build the package by following the [instructions](#) included in the package.

Building the Tcl API

Tcl support is not built automatically. See [Loading Berkeley DB with Tcl](#) for information on sites from which you can download Tcl and which Tcl versions are compatible with Berkeley DB. These notes assume that Tcl is installed as `d:\tcl`, but you can change that if you want.

The Tcl library must be built as the same build type as the Berkeley DB library (both Release or both Debug). We found that the binary release of Tcl can be used with the Release configuration of Berkeley DB, but you will need to build Tcl from sources for the Debug configuration. Before building Tcl, you will need to modify its makefile to make sure that you are building a debug version, including thread support. This is because the set of DLLs linked into the Tcl executable must match the corresponding set of DLLs used by Berkeley DB.

1. Set the include directories.
 - In Visual Studio 2005/Visual Studio 2008 - Choose Tools -> Options -> Projects -> VC++ Directories. Under the "Show directories for" pull-down, select "Include files". Add the full pathname for `d:\tcl\include`, then click OK.

- In Visual Studio 2010 - Right-click db_tcl project, choose Properties->Configuration Properties-> VC++ Directories->Include Directories. Add the full pathnames for d:\tcl\include, then click OK.

This is the directory that contains tcl.h.

2. Set the library files directory.

- In Visual Studio 2005/Visual Studio 2008 - Choose Tools -> Options -> Projects -> VC++ Directories. Under the "Show directories for" pull-down, select "Library files". Add the full pathname for the d:\tcl\lib directory, then click OK.
- In Visual Studio 2010 - Right-click db_tcl project, choose Properties->Configuration Properties-> VC++ Directories->Library Directories. Add the full pathname for the d:\tcl\lib directory, then click OK.

This is the directory needed to find tcl85g.lib (or whatever the library is named in your distribution).

3. Set the build type to Release or Debug in the drop-down on the tool bar.
4. To build, right-click on db_tcl and select Build. This builds the Tcl support library for Berkeley DB, placing the result into one of the following Berkeley DB subdirectories, depending upon the configuration that you chose:

build_windows\Win32\Debug\libdb_tcl53d.dll
--

build_windows\Win32\Release\libdb_tcl53.dll

If you use a version different from Tcl 8.5.x you will need to change the name of the Tcl library used in the build (for example, tcl85g.lib) to the appropriate name. To do this, right click on db_tcl, go to Properties -> Linker -> Input -> Additional dependencies and change tcl85g.lib to match the Tcl version you are using.

Distributing DLLs

When distributing applications linked against the DLL (not static) version of the library, the DLL files you need will be found in one of the following Berkeley DB subdirectories, depending upon the configuration that you chose:

build_windows\Win32\Debug

build_windows\Win32\Release

build_windows\Win32\Debug_static

build_windows\Win32\Release_static

build_windows\x64\Debug

build_windows\x64\Release
build_windows\x64\Debug_static
build_windows\x64\Release_static

You may also need to redistribute DLL files needed for the compiler's runtime. Generally, these runtime DLL files can be installed in the same directory that will contain your installed Berkeley DB DLLs. This directory may need to be added to your System PATH environment variable. Check your compiler's license and documentation for specifics on redistributing runtime DLLs.

Additional build options

There are several build options that you can configure when building Berkeley DB on Windows. To specify these, select `Project Properties -> C/C++ -> Command Line` and add the property.

These are some of the additional properties that you can specify when you are building Berkeley DB on Windows:

- `/D HAVE_LOCALIZATION`

Enable localized error message text, if available. This option should not be used when enabling stripped messages.

- `/D HAVE_MIXED_SIZE_ADDRESSING`

Allows for the sharing of the BDB database environment between 32-bit and 64-bit applications. Note that if you use this macro to rebuild your Berkeley DB library, then you need to also rebuild both your 32-bit and 64-bit applications using `/D HAVE_MIXED_SIZE_ADDRESSING`.

Note that use of this macro means that private environments are disabled for the library.

- `/D HAVE_STRIPPED_MESSAGES`

Causes all error messages to be stripped of their textual information. This option should not be used when enabling localization support. Use of this property can reduce your library footprint by up to 42KB (for DLLs) or 98KB (for a .lib).

Note that this option is automatically enabled if you build using the `db_small` project name. For more information on building a small library, see [Building a small memory footprint library](#).

If you use this build option, you can get an idea of what text should be issued for a given error message by using the [Message Reference for Stripped Libraries](#) guide.

Building a small memory footprint library

For applications that don't require all of the functionality of the full Berkeley DB library, an option is provided to build a static library with certain functionality disabled. In particular, cryptography, hash and queue access methods, replication and verification are all turned off. In addition, all message text is stripped from the library. This can reduce the memory footprint of Berkeley DB significantly.

Note

If your library has stripped messages, you can get an idea of what text should be issued for a given error message by using the [Message Reference for Stripped Libraries](#) guide.

In general on Windows systems, you will want to evaluate the size of the final application, not the library build. The Microsoft LIB file format (like UNIX archives) includes copies of all of the object files and additional information. The linker rearranges symbols and strips out the overhead, and the resulting application is much smaller than the library. There is also a Visual C++ optimization to "Minimize size" that will reduce the library size by a few percent.

A Visual C++ project file called `db_small` is provided for this small memory configuration. During a build, static libraries are created in `Release` or `Debug`, respectively. The library name is `libdb_small48sd.lib` for the debug build, or `libdb_small48s.lib` for the release build.

For assistance in further reducing the size of the Berkeley DB library, or in building small memory footprint libraries on other systems, please contact Berkeley DB support.

Running the test suite under Windows

[Building the software needed by the tests](#)

[Running the test suite under Windows](#)

[Building the software needed by the SQL tests](#)

To build the test suite on Windows platforms, you will need to configure Tcl support. You will also need sufficient main memory (at least 64MB), and disk (around 250MB of disk will be sufficient).

Building the software needed by the tests

The test suite must be run against a Debug version of Berkeley DB, so you will need a Debug version of the Tcl libraries. This involves building Tcl from its source. See the Tcl sources for more information. Then build the Tcl API - see [Building the Tcl API](#) for details.

Visual Studio 2005 or newer

To build for testing, perform the following steps:

1. Open the Berkeley DB solution.
2. Ensure that the target configuration is Debug
3. Right click the db_tcl project in the Solution Explorer, and select Build.
4. Right click the db_test project in the Solution Explorer, and select Build.

Running the test suite under Windows

Before running the tests for the first time, you must edit the file `include.tcl` in your build directory and change the line that reads:

```
set tclsh_path SET_YOUR_TCLSH_PATH
```

You will want to use the location of the `tclsh` program (be sure to include the name of the executable). For example, if Tcl is installed in `d:\tcl`, this line should be the following:

```
set tclsh_path d:\tcl\bin\tclsh85g.exe
```

If your path includes spaces be sure to enclose it in quotes:

```
set tclsh_path "c:\Program Files\tcl\bin\tclsh85g.exe"
```

Make sure that the path to Berkeley DB's tcl library is in your current path. On Windows NT/2000/XP, edit your PATH using the My Computer -> Properties -> Advanced -> Environment Variables dialog. On earlier versions of Windows, you may find it convenient to add a line to `c:\AUTOEXEC.BAT`:

```
SET PATH=%PATH%;c:\db\build_windows
```

Then, in a shell of your choice enter the following commands:

1. `cd build_windows`
2. `run d:\tcl\bin\tclsh85g.exe`, or the equivalent name of the Tcl shell for your system.

You should get a "%" prompt.

3. % source ../test/tcl/test.tcl

If no errors occur, you should get a "%" prompt.

You are now ready to run tests in the test suite; see [Running the test suite](#) for more information.

Building the software needed by the SQL tests

The SQL test suite must be run against a Debug version of Berkeley DB, so you need a Debug version of the Tcl libraries. This involves building Tcl from its source. See the [Tcl sources](#) for more information. Then build the Tcl API - see [Building the Tcl API](#) for details.

Before building for SQL tests, build the db_tcl and db_sql_testfixture projects. This requires Tcl 8.5 or above. If you are using a later version of Tcl, edit the Tcl library that db_tcl and db_sql_testfixture link to.

To do this right click the db_tcl/db_sql_testfixture project, select Properties->Configuration Properties->Linker->Input->Additional Dependencies and edit the Tcl library, tcl85g.lib, to match the version you are using.

Building the db_sql_testfixture project builds the testfixture.exe program in ../build_windows/Win32/Debug. It also builds the projects db and db_sql, on which it depends.

Visual Studio 2005 or newer

To build for testing, perform the following steps:

1. Open the Berkeley DB solution.
2. Ensure that the target configuration is Debug.
3. Right click the db_tcl project in the Solution Explorer, and select Build.
4. Right click the db_sql_testfixture project in the Solution Explorer, and select Build.

To test extensions, specify the following in the Preprocessor Definitions of the db_sql_testfixture project:

- SQLITE_ENABLE_FTS3 to enable the full text search layer
- SQLITE_ENABLE_RTREE to enable the R-Tree layer

Windows notes

If a system memory environment is closed by all processes, subsequent attempts to open it will return an error. To successfully open a transactional environment in this state,

recovery must be run by the next process to open the environment. For non-transactional environments, applications should remove the existing environment and then create a new database environment.

1. Berkeley DB does not support the Windows/95, Windows/98 or Windows/ME platforms.
2. On Windows, system paging file memory is freed on last close. For this reason, multiple processes sharing a database environment created using the `DB_SYSTEM_MEM` flag must arrange for at least one process to always have the environment open, or alternatively that any process joining the environment be prepared to re-create it.
3. When using the `DB_SYSTEM_MEM` flag, Berkeley DB shared regions are created without ACLs, which means that the regions are only accessible to a single user. If wider sharing is appropriate (for example, both user applications and Windows/NT service applications need to access the Berkeley DB regions), the Berkeley DB code will need to be modified to create the shared regions with the correct ACLs. Alternatively, by not specifying the `DB_SYSTEM_MEM` flag, filesystem-backed regions will be created instead, and the permissions on those files may be directly specified through the `DB_ENV->open()` method.
4. Applications that operate on wide character strings can use the Windows function `WideCharToMultiByte` with the code page `CP_UTF8` to convert paths to the form expected by Berkeley DB. Internally, Berkeley DB calls `MultiByteToWideChar` on paths before calling Windows functions.
5. Various Berkeley DB methods take a mode argument, which is intended to specify the underlying file permissions for created files. Berkeley DB currently ignores this argument on Windows systems.

It would be possible to construct a set of security attributes to pass to `CreateFile` that accurately represents the mode. In the worst case, this would involve looking up user and all group names, and creating an entry for each. Alternatively, we could call the `_chmod` (partial emulation) function after file creation, although this leaves us with an obvious race.

Practically speaking, however, these efforts would be largely meaningless on a FAT file system, which only has a "readable" and "writable" flag, applying to all users.

Windows FAQ

1. My Win* C/C++ application crashes in the Berkeley DB library when Berkeley DB calls `fprintf` (or some other standard C library function).

You should be using the "Debug Multithreaded DLL" compiler option in your application when you link with the `build_windows\Debug\libdb48d.lib` library (this .lib file is actually a stub for `libdb48d.DLL`). To check this setting in Visual C++, choose the Project/Settings menu item and select Code Generation under the tab marked C/C++; and see the box marked Use runtime library. This should be set to Debug Multithreaded DLL. If your application is linked against the static library, `build_windows\Debug\libdb48sd.lib`; then, you will want to set Use runtime library to Debug Multithreaded.

Setting this option incorrectly can cause multiple versions of the standard libraries to be linked into your application (one on behalf of your application, and one on behalf of the Berkeley DB library). That violates assumptions made by these libraries, and traps can result.

Also, using different Visual Studio compilers in the application and libraries can lead to a crash. So rebuild the application with the same Visual C++ version as that of the library.

2. Why are the build options for DB_DLL marked as "Use MFC in a Shared DLL"? Does Berkeley DB use MFC?

Berkeley DB does not use MFC at all. It does however, call `malloc` and `free` and other facilities provided by the Microsoft C runtime library. We found in our work that many applications and libraries are built assuming MFC, and specifying this for Berkeley DB solves various interoperation issues, and guarantees that the right runtime libraries are selected. Note that because we do not use MFC facilities, the MFC library DLL is not marked as a dependency for `libdb.dll`, but the appropriate Microsoft C runtime is.

3. How can I build Berkeley DB for [MinGW](#)?

Follow the instructions in [Building for UNIX/POSIX](#), and specify the `--enable-mingw` option to the configuration script. This configuration option currently only builds static versions of the library, it does not yet build a DLL version of the library, and file sizes are limited to 2GB (2^{32} bytes.)

4. How can I build a Berkeley DB for Windows 98/ME?

Windows 98/ME is no longer supported by Berkeley DB. The following is therefore only of interest to historical users of Berkeley DB.

By default on Windows, Berkeley DB supports internationalized filenames by treating all directory paths and filenames passed to Berkeley DB methods as UTF-8

encoded strings. All paths are internally converted to wide character strings and passed to the wide character variants of Windows system calls.

This allows applications to create and open databases with names that cannot be represented with ASCII names while maintaining compatibility with applications that work purely with ASCII paths.

Windows 98 and ME do not support Unicode paths directly. To build for those versions of Windows, either:

- Follow the instructions at [Microsoft's web site](#).
- Open the workspace or solution file with Visual Studio. Then open the Project properties/settings section for the project you need to build (at least db_dll). In the C/C++->Preprocessor->Preprocessor Definitions section, remove _UNICODE and UNICODE entries. Add in an entry of _MBCS. Build the project as normal.

The ASCII builds will also work on Windows NT/2K/XP/2003 and Windows7, but will not translate paths to wide character strings.

Chapter 7. Building Berkeley DB for UNIX/POSIX

Table of Contents

[Building for UNIX/POSIX](#)

[Building the Berkeley DB SQL Interface](#)

[Configuring Berkeley DB](#)

[Configuring the SQL Interface](#)

[Changing Compile Options](#)

[Enabling Extensions](#)

[Building the JDBC Driver](#)

[Using the JDBC Driver](#)

[Building the ODBC Driver](#)

[Building the BFILE extension](#)

[Building a small memory footprint library](#)

[Changing compile or load options](#)

[Cross-Compiling on Unix](#)

[Installing Berkeley DB](#)

[Dynamic shared libraries](#)

[Running the test suite under UNIX](#)

[Building SQL Test Suite on Unix](#)

[Architecture independent FAQ](#)

[AIX](#)

[FreeBSD](#)

[Apple iOS \(iPhone OS\)](#)

[IRIX](#)

[Linux](#)

[Mac OS X](#)

[QNX](#)

[SCO](#)

[Solaris](#)

[SunOS](#)

Building for UNIX/POSIX

[Building the Berkeley DB SQL Interface](#)

The Berkeley DB distribution builds up to four separate libraries: the base C API Berkeley DB library and the optional C++, Java, and Tcl API libraries. For portability reasons, each library is standalone and contains the full Berkeley DB support necessary to build applications; that is, the C++ API Berkeley DB library does not require any other Berkeley DB libraries to build and run C++ applications.

Building for Linux, Apple iOS (known as iPhone OS previously), Mac OS X or the QNX Neutrino release is the same as building for a conventional UNIX platform.

The Berkeley DB distribution uses the Free Software Foundation's [autoconf](#) and [libtool](#) tools to build on UNIX platforms. In general, the standard configuration and installation options for these tools apply to the Berkeley DB distribution.

To perform a standard UNIX build of Berkeley DB, change to the `build_unix` directory and then enter the following two commands:

```
../dist/configure  
make
```

This will build the Berkeley DB library.

To install the Berkeley DB library, enter the following command:

```
make install
```

To rebuild Berkeley DB, enter:

```
make clean  
make
```

If you change your mind about how Berkeley DB is to be configured, you must start from scratch by entering the following command:

```
make realclean  
../dist/configure  
make
```

To uninstall Berkeley DB, enter:

```
make uninstall
```

To build multiple UNIX versions of Berkeley DB in the same source tree, create a new directory at the same level as the `build_unix` directory, and then configure and build in that directory as described previously.

Building the Berkeley DB SQL Interface

To perform a standard UNIX build of the Berkeley DB SQL interface, go to the `build_unix` directory and then enter the following two commands:

```
../dist/configure --enable-sql  
make
```

This creates a library, `libdb_sql`, and a command line tool, `dbsql`. You can create and manipulate SQL databases using the `dbsql` shell.

You can optionally provide the `--enable-sql_compat` argument to the `configure` script. In addition to creating `libdb_sql` and `dbsql` this causes a thin wrapper library called `libsqlite3` and a command line tool called `sqlite3` to be built. This library can be used as a drop-in replacement for SQLite. The `sqlite3` command line tool is identical to the `dbsql` executable but is named so that existing scripts for SQLite can easily work with Berkeley DB.

```
../dist/configure --enable-sql_compat
make
```

There are several arguments you can specify when configuring the Berkeley DB SQL Interface. See [Configuring the SQL Interface](#) for more information.

Configuring Berkeley DB

There are several arguments you can specify when configuring Berkeley DB. Although only the Berkeley DB-specific ones are described here, most of the standard GNU autoconf arguments are available and supported. To see a complete list of possible arguments, specify the `--help` flag to the `configure` program.

The Berkeley DB specific arguments are as follows:

- `--disable-largefile`

Some systems, notably versions of Solaris, require special compile-time options in order to create files larger than 2^{32} bytes. These options are automatically enabled when Berkeley DB is compiled. For this reason, binaries built on current versions of these systems may not run on earlier versions of the system because the library and system calls necessary for large files are not available. To disable building with these compile-time options, enter `--disable-largefile` as an argument to `configure`.

- `--disable-shared`, `--disable-static`

On systems supporting shared libraries, Berkeley DB builds both static and shared libraries by default. (Shared libraries are built using [the GNU Project's Libtool](#) distribution, which supports shared library builds on many (although not all) systems.) To not build shared libraries, configure using the `--disable-shared` argument. To not build static libraries, configure using the `--disable-static` argument.

- `--disable-heap`

Disables the Heap access method so that it cannot be used by Berkeley DB applications.

- `--enable-compat185`

To compile or load Berkeley DB 1.85 applications against this release of the Berkeley DB library, enter `--enable-compat185` as an argument to configure. This will include Berkeley DB 1.85 API compatibility code in the library.

- `--enable-cxx`

To build the Berkeley DB C++ API, enter `--enable-cxx` as an argument to configure.

- `--enable-debug`

To build Berkeley DB with `-g` as a compiler flag and with `DEBUG` #defined during compilation, enter `--enable-debug` as an argument to configure. This will create a Berkeley DB library and utilities with debugging symbols, as well as load various routines that can be called from a debugger to display pages, cursor queues, and so forth. If installed, the utilities will not be stripped. This argument should not be specified when configuring to build production binaries.

- `--enable-debug_rop`

To build Berkeley DB to output log records for read operations, enter `--enable-debug_rop` as an argument to configure. This argument should not be specified when configuring to build production binaries.

- `--enable-debug_wop`

To build Berkeley DB to output log records for write operations, enter `--enable-debug_wop` as an argument to configure. This argument should not be specified when configuring to build production binaries.

- `--enable-diagnostic`

To build Berkeley DB with run-time debugging checks, enter `--enable-diagnostic` as an argument to configure. This causes a number of additional checks to be performed when Berkeley DB is running, and also causes some failures to trigger process abort rather than returning errors to the application. Applications built using this argument should not share database environments with applications built without this argument. This argument should not be specified when configuring to build production binaries.

- `--enable-dump185`

To convert Berkeley DB 1.85 (or earlier) databases to this release of Berkeley DB, enter `--enable-dump185` as an argument to configure. This will build the [db_dump185](#) utility, which can dump Berkeley DB 1.85 and 1.86 databases in a format readable by the Berkeley DB [db_load](#) utility.

The system libraries with which you are loading the [db_dump185](#) utility must already contain the Berkeley DB 1.85 library routines for this to work because the Berkeley DB distribution does not include them. If you are using a non-standard library for the Berkeley DB 1.85 library routines, you will have to change the Makefile that the configuration step creates to load the [db_dump185](#) utility with that library.

- `--enable-java`

To build the Berkeley DB Java API, enter `--enable-java` as an argument to configure. To build Java, you must also build with shared libraries. Before configuring, you must set your `PATH` environment variable to include `javac`. Note that it is not sufficient to include a symbolic link to `javac` in your `PATH` because the configuration process uses the location of `javac` to determine the location of the Java include files (for example, `jni.h`). On some systems, additional include directories may be needed to process `jni.h`; see [Changing compile or load options](#) for more information.

- `--enable-posixmutexes`

To force Berkeley DB to use the POSIX pthread mutex interfaces for underlying mutex support, enter `--enable-posixmutexes` as an argument to configure. This is rarely necessary: POSIX mutexes will be selected automatically on systems where they are the preferred implementation.

The `--enable-posixmutexes` configuration argument is normally used in two ways: First, when there are multiple mutex implementations available and the POSIX mutex implementation is not the preferred one (for example, on Solaris where the LWP mutexes are used by default). Second, by default the Berkeley DB library will only select the POSIX mutex implementation if it supports mutexes shared between multiple processes, as described for the `pthread_condattr_setpshared` and `pthread_mutexattr_setpshared` interfaces. The `--enable-posixmutexes` configuration argument can be used to force the selection of POSIX mutexes in this case, which can improve application performance significantly when the alternative mutex implementation is a non-blocking one (for example test-and-set assembly instructions). However, configuring to use POSIX mutexes when the implementation does not have inter-process support will only allow the creation of private database environments, that is, environments where the [DB_PRIVATE](#) flag is specified to the [DB_ENV->open\(\)](#) method.

Specifying the `--enable-posixmutexes` configuration argument may require that applications and Berkeley DB be linked with the `-lpthread` library.

- `--enable-pthread_api`

To configure Berkeley DB for a POSIX pthreads application (with the exception that POSIX pthread mutexes may not be selected as the underlying mutex implementation for the build), enter `--enable-pthread_api` as an argument to configure. The build will include the Berkeley DB replication manager interfaces and will use the POSIX standard `pthread_self` and `pthread_yield` functions to identify threads of control and yield the processor. The `--enable-pthread_api` argument requires POSIX pthread support already be installed on your system.

Specifying the `--enable-pthread_api` configuration argument may require that applications and Berkeley DB be linked with the `-lpthread` library.

- `--enable-sql`

To build the command tool `dbsql`, enter `--enable-sql` as an argument to configure. The `dbsql` utility provides access to the Berkeley DB SQL interface. See [Configuring the SQL Interface](#) for more information.

- `--enable-sql_compat`

To build the command tool `sqlite3`, enter `--enable-sql_compat` as an argument to configure. `Sqlite3` is a command line tool that enables you to manually enter and execute SQL commands. It is identical to the `dbsql` executable but named so that existing scripts for SQLite can easily work with Berkeley DB. See [Configuring the SQL Interface](#) for more information.

- `--enable-sql_codegen`

To build the command line tool `db_sql_codegen`, enter `--enable-sql_codegen` as an argument to configure. The `db_sql_codegen` utility translates a schema description written in a SQL Data Definition Language dialect into C code that implements the schema using Berkeley DB.

- `--enable-smallbuild`

To build a small memory footprint version of the Berkeley DB library, enter `--enable-smallbuild` as an argument to configure. The `--enable-smallbuild` argument is equivalent to individually specifying `--with-cryptography=no`, `--disable-hash`, `--disable-queue`, `--disable-replication`, `--disable-statistics` and `--disable-verify`, turning off cryptography support, the Hash and Queue access methods, database

environment replication support and database and log verification support. See [Building a small memory footprint library](#) for more information.

- `--enable-stl`

To build the Berkeley DB C++ STL API, enter `--enable-stl` as an argument to `configure`. Setting this argument implies that `--enable-cxx` is set, and the Berkeley DB C++ API will be built too.

There will be a `libdb_stl-X.X.a` and `libdb_stl-X.X.so` built, which are the static and shared library you should link your application with in order to make use of Berkeley DB via its STL API.

If your compiler is not ISO C++ compliant, the `configure` may fail with this argument specified because the STL API requires standard C++ template features. In this case, you will need a standard C++ compiler. So far gcc is the best choice, we have tested and found that gcc-3.4.4 and all its newer versions can build the Berkeley DB C++ STL API successfully.

For information on `db_stl` supported compilers, see the [Portability section](#) in the Programmer's Reference Guide.

And you need to include the STL API header files in your application code. If you are using the Berkeley DB source tree, the header files are in `<Berkeley DB Source Root>/stl` directory; If you are using the installed version, these header files are in `<Berkeley DB Installed Directory>/include`, as well as the `db.h` and `db_cxx.h` header files.

- `--enable-tcl`

To build the Berkeley DB Tcl API, enter `--enable-tcl` as an argument to `configure`. This configuration argument expects to find Tcl's `tclConfig.sh` file in the `/usr/local/lib` directory. See the `--with-tcl` argument for instructions on specifying a non-standard location for the Tcl installation. See [Loading Berkeley DB with Tcl](#) for information on sites from which you can download Tcl and which Tcl versions are compatible with Berkeley DB. To build Tcl, you must also build with shared libraries.

- `--enable-test`

To build the Berkeley DB test suite, enter `--enable-test` as an argument to `configure`. To run the Berkeley DB test suite, you must also build the Tcl API. This argument should not be specified when configuring to build production binaries.

- `--enable-uimutexes`

To force Berkeley DB to use the UNIX International (UI) mutex interfaces for underlying mutex support, enter `--enable-uimutexes` as an argument to configure. This is rarely necessary: UI mutexes will be selected automatically on systems where they are the preferred implementation.

The `--enable-uimutexes` configuration argument is normally used when there are multiple mutex implementations available and the UI mutex implementation is not the preferred one (for example, on Solaris where the LWP mutexes are used by default).

Specifying the `--enable-uimutexes` configuration argument may require that applications and Berkeley DB be linked with the `-lthread` library.

- `--enable-umrw`

Rational Software's Purify product and other run-time tools complain about uninitialized reads/writes of structure fields whose only purpose is padding, as well as when heap memory that was never initialized is written to disk. Specify the `--enable-umrw` argument during configuration to mask these errors. This argument should not be specified when configuring to build production binaries.

- `--enable-dtrace` [`--enable-perfmon-statistics`]

To build Berkeley DB with performance event monitoring probes add `--enable-dtrace` to the configuration options. Both native DTrace (on Solaris and Mac OS X) and the Statically Defined Tracing compatibility layer in Linux SystemTap version 1.1 or better are supported. That compatibility package may be called `systemtap-sdt-devel`; it includes `sys/sdt.h`.

If `--enable-perfmon-statistics` is combined with `--enable-dtrace` then additional probes are defined for the tracking variables from which DB's statistics are obtained. They allow DTrace and SystemTap access to these values when they are updated, are the basis of the statistics as displayed `db_stat` and the API functions that return statistics.

The `--enable-dtrace` option may not be specified at the same time as `--disable-statistics`.

For information on Berkeley DB Performance Event Monitoring, see the [Performance Event Monitoring](#) section in the Programmer's Reference Guide.

- `--enable-localization`

Enable localized error message text, if available. This option should not be used when `--enable-stripped_messages` is in use.

- `--enable-stripped_messages`

Causes all error messages to be stripped of their textual information. Instead, only error return codes are used. This option should not be used when `--enable-localization` is in use. Use of this build option can reduce your library foot print by up to 44KB (.so) or 50KB (.a).

If you use this configuration option, you can get an idea of what text should be issued for a given error message by using the [Message Reference for Stripped Libraries](#) guide.

- `--with-cryptography`

To build Berkeley DB with support for cryptography, enter `--with-cryptography=yes` as an argument to configure.

To build Berkeley DB without support for cryptography, enter `--with-cryptography=no` as an argument to configure.

To build Berkeley DB with support for cryptography using Intel's Performance Primitive (IPP) library, enter `--with-cryptography=ipp` as an argument to configure. Additionally, set the following arguments:

`-L/path/to/ipp/sharedlib` to `LDFLAGS`

`-I/path/to/ipp/include` to `CPPFLAGS`

`-lippcpem64t -lpthread` to `LIBS`

An example configuration command for IPP encryption is as follows:

```
../dist/configure --with-cryptography=ipp
CPPFLAGS="-I/opt/intel/ipp/6.1.3.055/em64t/include"
LDFLAGS="-L/opt/intel/ipp/6.1.3.055/em64t/sharedlib"
LIBS="-lippcpem64t -lpthread"
```

See the [Intel Documentation](#) for specific instructions on configuring environment variables.

Note: The `--with-cryptography=ipp` argument works only on Linux.

- `--with-mutex=MUTEX`

To force Berkeley DB to use a specific mutex implementation, configure with `--with-mutex=MUTEX`, where `MUTEX` is the mutex implementation you want. For example, `--with-mutex=x86/gcc-assembly` will configure Berkeley DB to use the x86 GNU gcc

compiler based test-and-set assembly mutexes. This is rarely necessary and should be done only when the default configuration selects the wrong mutex implementation. A list of available mutex implementations can be found in the distribution file `dist/aclocal/mutex.m4`.

- `--with-tcl=DIR`

To build the Berkeley DB Tcl API, enter `--with-tcl=DIR`, replacing `DIR` with the directory in which the Tcl `tclConfig.sh` file may be found. See [Loading Berkeley DB with Tcl](#) for information on sites from which you can download Tcl and which Tcl versions are compatible with Berkeley DB. To build Tcl, you must also build with shared libraries.

- `--with-uniquename=NAME`

To build Berkeley DB with unique symbol names (in order to avoid conflicts with other application modules or libraries), enter `--with-uniquename=NAME`, replacing `NAME` with a string that to be appended to every Berkeley DB symbol. If `"=NAME"` is not specified, a default value of `"_MAJORMINOR"` is used, where `MAJORMINOR` is the major and minor release numbers of the Berkeley DB release. See [Building with multiple versions of Berkeley DB](#) for more information.

Configuring the SQL Interface

[Changing Compile Options](#)

[Enabling Extensions](#)

[Building the JDBC Driver](#)

[Using the JDBC Driver](#)

[Building the ODBC Driver](#)

[Building the BFILE extension](#)

There are a set of options you can provide to configure in order to control how the Berkeley DB SQL interface is built. These configuration options include:

`--disable-log-checksum`

Disables checksums in log records. This provides a boost to performance at the risk of log files having undetectable corruption that could prevent proper data recovery in case of database corruption.

Note that while this option is meant for use with the SQL interface, it will also disable checksum for the non-SQL interfaces.

`--enable-sql`

Causes the `dbsql` command line interpreter to be built. Along with `dbsql`, this argument also builds the `libdb_sqlXX.{so|la}` library, a C API library that mirrors the SQLite C API.

`--enable-sql_compat`

Causes the `sqlite3` command line tool to be built. This tool is identical to the `dbsql` command line tool, except that it has the same name as the command line tool that comes with standard SQLite.

In addition, the `libsqlite3.{so|la}` C API library is built if this option is specified. This library is identical to the `libdb_sqlXX.{so|la}` library that is normally built for Berkeley DB's `sql` interface, except that it has the same name as the library which is built for standard SQLite.

Warning

Use this compatibility option with extreme care. Standard SQLite is used by many programs and utilities on many different platforms. Some platforms, such as Mac OS X, come with standard SQLite built in because default applications for the platform use that library.

Use of this option on platforms where standard SQLite is in production use can cause unexpected runtime errors either for your own application, or for applications and utilities commonly found on the platform, depending on which library is found first in the platform's library search path.

Use this option only if you know exactly what you are doing.

This option is provided so that there is an easy upgrade path for legacy SQLite tools and scripts that want to use BDB SQL without rewriting the tool or script. However, data contained in standard SQLite databases must be manually migrated from the old database to your BDB SQL database even if you use this option. See the Berkeley DB Getting Started with the SQL APIs guide for information on migrating data from standard SQLite to BDB SQL databases.

Note that in addition to the renamed command line tool and library, this option also causes versions of the command line tool and library to be built that use the normal BDB SQLite names (`dbsql` and `libdb_sqlXX.{so|la}`).

`--enable-test`

Cause the Berkeley DB SQL interface test suite to be built. This argument can also be used with either `--enable-sql` or `--enable-sql_compat` to build the SQLite Tcl test runner.

--enable-jdbc

Causes the JDBC driver to be built. Setting this option implies that `--enable-sql` is set, which means that the Berkeley DB SQL API will be built too.

The following configuration options are useful when debugging applications:

--enable-debug

Builds the Berkeley DB SQL interface with debug symbols.

--enable-diagnostic

Builds the Berkeley DB SQL interface with run-time debugging checks.

Any arguments that you can provide to the standard SQLite configure script can also be supplied when configuring Berkeley DB SQL interface.

Changing Compile Options

There are several configuration options you can specify as an argument to the configure script using the standard environment variable, CFLAGS.

BDBSQL_DEFAULT_PAGE_SIZE

To set the default page size when you create a database, specify the `BDBSQL_DEFAULT_PAGE_SIZE` flag. The value assigned must be a 0, 512, 1024, 2048, 4096, 8192, 16384, 32768, or 65536. The default value is 4096. If the value is set to zero, Berkeley DB queries the file system to determine the best page size, and the value of `SQLITE_DEFAULT_PAGE_SIZE` is used to calculate the cache size, as the cache size is specified as a number of pages.

BDBSQL_FILE_PER_TABLE

To generate each table in a separate file, rather than as subdatabases in a single file, specify the `BDBSQL_FILE_PER_TABLE` flag. When this option is enabled, the SQL database name is used as a directory name. This directory contains one file for the metadata and one file each for every table created by the SQL API. Note that adding or deleting files from the database directory may corrupt your database. To backup the metadata (schema), make a copy of the `metadata` and `table00001` files from the database directory. Make a new copy whenever the schema is changed.

BDBSQL_LOG_REGIONMAX

To configure the log region size for the underlying storage engine, specify the `BDBSQL_LOG_REGIONMAX` flag. For more information, see [DB_ENV->get_lg_regionmax\(\)](#).

BDBSQL_OMIT_LEAKCHECK

For Berkeley DB to use the default system allocation routines rather than the SQLite allocation routines, specify the `BDBSQL_OMIT_LEAKCHECK` flag.

BDBSQL_OMIT_LOG_REMOVE

Berkeley DB automatically removes log files that are not required any more, that is, files that are older than the most recent checkpoint. To disable this functionality, specify the BDBSQL_OMIT_LOG_REMOVE flag. It is necessary to provide this flag if you are using replication with Berkeley DB SQL.

BDBSQL_OMIT_SHARING

To create a private environment rather than a shared environment, specify the BDBSQL_OMIT_SHARING flag. That is, the cache and other region files will be created in memory rather than using file backed shared memory. For more information, see the DB_PRIVATE flag of [DB_ENV->open\(\)](#).

BDBSQL_SINGLE_THREAD

To disable locking and thread safe connections, specify the BDBSQL_SINGLE_THREAD flag. If an application is going to use Berkeley DB from a single thread and a single process, enabling this flag can deliver significant performance advantages.

SQLITE_DEFAULT_CACHE_SIZE

SQLite provides an in-memory cache which you size according to the maximum number of database pages that you want to hold in memory at any given time. Berkeley DB's in-memory cache feature performs the same function as SQLite. To specify the suggested maximum number of pages of disk cache that will be allocated per open database file specify the SQLITE_DEFAULT_CACHE_SIZE flag. Default value is 2000 pages. For more information, see the SQLite documentation on [PRAGMA default_cache_size](#).

SQLITE_DEFAULT_JOURNAL_SIZE_LIMIT

For SQLite, this pragma identifies the maximum size that the journal file is allowed to be. Berkeley DB does not have a journal file, but it writes and uses log files. A new log file is created when the current log file has reached the defined maximum size. To define this maximum size for a log file, specify the SQLITE_DEFAULT_JOURNAL_SIZE_LIMIT flag. Default value is 10 MB for the Berkeley DB SQL interface.

Enabling Extensions

The Berkeley DB SQL API provides extensions such as full text search and R-Tree index. By default, these extensions are disabled. To enable an extension in the Berkeley DB SQL interface, specify the related option as an argument to the configure script using the standard environment variable, CPPFLAGS.

SQLITE_ENABLE_FTS3

Enable building the Berkeley DB full text search layer

SQLITE_ENABLE_RTREE

Enables the Berkeley DB R-Tree layer.

See the SQLite Documentation for more information on [full text search](#) and [R-Tree](#).

Building the JDBC Driver

This section describes how to build the JDBC driver code using `autoconf`, which is the only method supported and tested by the Berkeley DB team.

To build the JDBC driver, you must have Sun Java Development Kit 1.1 or above installed.

```
cd build_unix
../dist/configure --enable-jdbc --prefix=<install path>
make install
```

You can test the build by entering the following commands from the `build_unix/jdbc` directory:

```
javac -classpath ./sqlite.jar test3.java
java -Djava.library.path=./libs -classpath ./sqlite.jar:. test3
```

Using the JDBC Driver

This section describes how to download, build, and run sample programs using the built JDBC driver.

Downloading JDBC Sample Code

The download link for JDBC sample code is available on the [Oracle Technology Network \(OTN\)](#) page. You can identify the link by the "JDBC programming examples from all three editions (ZIP format)" text beside it.

Modifying Sample Code

Before running the example code, do the following:

1. Unzip the file containing the sample code to a new directory (for example, `jdbc_ex`).
2. Substitute `jdbc:sqlite:<db-file-name>` for the generic JDBC URL that appears in the code. That is, put `jdbc:sqlite:<db-file-name>` between the quotation marks in the line:

```
String url = "jdbc:mySubprotocol:myDataSource";
```

Note: The `<db-file-name>` can either be an absolute path name like `"/jdbc_ex_db/myDataSource"`, or a relative path-file-name like `"/jdbc_ex_db/myDataSource"`, or a file name, like `"myDataSource"`, in which case the database file will be stored at the current directory.

3. Substitute `SQLite.JDBC.Driver` for `myDriver.ClassName` in the line:
`Class.forName("myDriver.ClassName");`
4. Optionally substitute the username and password you use for your database in the following: `"myLogin"`, `"myPassword"`.
5. If your JDK version is above 1.5, change the variable name `enum` in `OutputApplet.java` to some other variable name because, as of JDK release 5 `enum` is a keyword and can not be used as an identifier.

Building and Running the JDBC Sample code

See [Building the JDBC Driver](#) for instructions on building the JDBC driver.

To build and run the JDBC examples do the following:

1. Copy `build_unix/jdbc/sqlite.jar` and `build_unix/jdbc/.libs/libsqlite_jni.so` to the `jdbc_ex` directory.
2. In the `jdbc_ex` directory, run the following commands:

```
$ javac -classpath ./sqlite.jar *.java
$ java -classpath ./sqlite.jar -Djava.library.path=. \
  <ClassName, eg. CreateCoffees>
```

3. After you run the `CreateCoffees` example, use the `dbsql` executable to open the `myDataSource` database file and check if the table `COFFEES` has been successfully created in the database.

```
$ dbsql myDataSourcedbsql> .tables
COFFEES
dbsql> .dump
PRAGMA foreign_keys=OFF;
BEGIN TRANSACTION;
CREATE TABLE COFFEES (COF_NAME varchar(32),\
SUP_ID int, PRICE float, SALES int, TOTAL int);
COMMIT;
dbsql>
```

4. Repeat step 3 to run other examples.

Note: Some examples, such as `AutoGenKeys`, are not yet supported by BDB JDBC driver. The `SQLException` is displayed for those unsupported examples.

Building the ODBC Driver

This section describes how to build the ODBC driver.

Configuring Your System

To configure your system prior to building the ODBC driver, do the following:

1. Download and install the latest [unixODBC](#) if ODBC is not already installed on your system.
2. Configure the ODBC server to work with SQLite databases. Follow [these instructions](#) from Christian Werner.

Building the Library

To build the library, do the following:

```
$ cd db-5.3.XX/build_unix
$ CFLAGS="-fPIC" ../dist/configure --enable-sql_compat --disable-shared
$ make
$ cd ../lang/sql/odbc
$ CFLAGS="-DHAVE_ERRNO_H -I../././build_unix -I../././src/dbinc \
-I../sqlite/src" LDFLAGS="../././build_unix/libdb-5.3.a" \
./configure --with-sqlite3=../generated
$ make
```

The `libsqlite3odbc.so` library containing a statically linked version of Berkeley DB SQL is now built.

NOTE: The final `make` command above is known to generate a warning when using GCC. The warning states: Warning: Linking the shared library `libsqlite3odbc.la` against the static library `.././build_unix/libdb-5.3.a` is not portable! . It is generally safe to ignore the warning when using the generated library.

Testing the ODBC Driver

The steps to verify that the installed driver works are as follows:

1. Alter the `/etc/odbcinst.ini` and `~/odbc.ini` configuration files to refer to the `libsqlite3odbc.so` file built above.
2. Create a data source, and launch a data source viewer application by doing the following:

```
$ mkdir ~/databases
$ cd ~/databases
$ /path/to/Berkeley DB/build_unix/sqlite3 mytest.db
dbsql> CREATE TABLE t1(x);
dbsql> .quit;
$ DataManager
```

The final step opens a GUI application that displays ODBC data sources on a system. You should be able to find the `mytest.db` data source just created.

Building the BFILE extension

The BFILE extension allows you to store binary files outside of the database, but still operate upon them as if they were stored within the database. To enable this extension, use the `--enable-load-extension` configuration flag. For example:

```
$ cd <db>/build_unix
$ export DBSQL_DIR=$PWD/../install
$ ../dist/configure --enable-sql --enable-load-extension \
  --prefix=$DBSQL_DIR && make && make install
$ cd ../lang/sql/sqlite/ext/bfile/build
$ make && make install
```

BFILE extensions are only supported for Unix platforms.

Note that the extension support has two interfaces: SQL expressions and a C-functions API. By default, the SQL expressions are built when you use `--enable-load_extension`. To use the C-functions API, edit `<db>/lang/sql/ext/bfile/build/Makefile` and set `ENABLE_BFILE_CAPI` to `1`.

Once you have enabled the extension and built the library, you can run the included example:

```
$ cd lang/sql/sqlite/ext/bfile/build
$ export LD_LIBRARY_PATH=$PWD:$DBSQL_DIR/lib
$ ./bfile_example_sql    # for SQL expressions interface
$ ./bfile_example_capi   # for C-functions API
```

For more information on using the BFILE extension, see the Berkeley DB Getting Started with the SQL APIs guide.

Building a small memory footprint library

There are a set of configuration options to assist you in building a small memory footprint library. These configuration options turn off specific functionality in the Berkeley DB library, reducing the code size. These configuration options include:

`--enable-smallbuild`

Equivalent to individually specifying all of the following configuration options. In addition, when compiling building with the GNU gcc compiler, this option uses the `-Os` compiler build flag instead of the default `-O3`.

`--with-cryptography=no`

Builds Berkeley DB without support for cryptography.

`--disable-hash`

Builds Berkeley DB without support for the Hash access method.

`--disable-heap`

Builds Berkeley DB without support for the Heap access method.

`--disable-queue`

Builds Berkeley DB without support for the Queue access method.

`--disable-replication`

Builds Berkeley DB without support for the database environment replication.

`--disable-statistics`

Builds Berkeley DB without support for the statistics interfaces.

`--disable-verify`

Builds Berkeley DB without support for database verification.

`--enable-stripped_messages`

Strips message text from the error messages issued by Berkeley DB. This can reduce the size of the library by roughly another 22KB.

If your library has stripped messages, you can get an idea of what text should be issued for a given error message by using the [Message Reference for Stripped Libraries](#) guide.

Note

`--disable-cryptography` and `--enable-cryptography` are deprecated in the Berkeley DB 11gR2 release. Use `--with-cryptography=no` and `--with-cryptography=yes` instead.

The following configuration options will increase the size of the Berkeley DB library dramatically and are only useful when debugging applications:

`--enable-debug`

Build Berkeley DB with symbols for debugging.

`--enable-debug_rop`

Build Berkeley DB with read-operation logging.

--enable-debug_wop

Build Berkeley DB with write-operation logging.

--enable-diagnostic

Build Berkeley DB with run-time debugging checks.

In addition, static libraries are usually smaller than shared libraries. By default Berkeley DB will build both shared and static libraries. To build only a static library, configure Berkeley DB with the [Configuring Berkeley DB](#) option.

The size of the Berkeley DB library varies depending on the compiler, machine architecture, and configuration options. As an estimate, production Berkeley DB libraries built with GNU gcc version 4.X compilers have footprints in the range of 600KB to 1.4MB on 32-bit x86 architectures, and in the range of 700KB to 1.6MB on 64-bit x86 architectures.

For assistance in further reducing the size of the Berkeley DB library, or in building small memory footprint libraries on other systems, please contact Berkeley DB support.

Changing compile or load options

You can specify compiler and/or compile and load time flags by using environment variables during Berkeley DB configuration. For example, if you want to use a specific compiler, specify the CC environment variable before running configure:

```
prompt: env CC=gcc ../dist/configure
```

Using anything other than the native compiler will almost certainly mean that you'll want to check the flags specified to the compiler and loader, too.

To specify debugging and optimization options for the C compiler, use the CFLAGS environment variable:

```
prompt: env CFLAGS=-O2 ../dist/configure
```

To specify header file search directories and other miscellaneous options for the C preprocessor and compiler, use the CPPFLAGS environment variable:

```
prompt: env CPPFLAGS=-I/usr/contrib/include ../dist/configure
```

To specify debugging and optimization options for the C++ compiler, use the CXXFLAGS environment variable:

```
prompt: env CXXFLAGS=-Woverloaded-virtual ../dist/configure
```

To specify miscellaneous options or additional library directories for the linker, use the `LD_FLAGS` environment variable:

```
prompt: env LD_FLAGS="-N32 -L/usr/local/lib" ../dist/configure
```

If you want to specify additional libraries, set the `LIBS` environment variable before running `configure`. For example, the following would specify two additional libraries to load, `"posix"` and `"socket"`:

```
prompt: env LIBS="-lposix -lsocket" ../dist/configure
```

Make sure that you prepend `-L` to any library directory names and that you prepend `-l` to any include file directory names! Also, if the arguments you specify contain blank or tab characters, be sure to quote them as shown previously; that is with single or double quotes around the values you are specifying for `LIBS`.

The `env` command, which is available on most systems, simply sets one or more environment variables before running a command. If the `env` command is not available to you, you can set the environment variables in your shell before running `configure`. For example, in `sh` or `ksh`, you could do the following:

```
prompt: LIBS="-lposix -lsocket" ../dist/configure
```

In `csh` or `tcsh`, you could do the following:

```
prompt: setenv LIBS "-lposix -lsocket"
prompt: ../dist/configure
```

See your command shell's manual page for further information.

Cross-Compiling on Unix

The purpose of cross-compiling is to build a Berkeley DB library on one platform that is to be used on a different platform. This section demonstrates how to build a library compatible with the ARM platform from a 32-bit x86 Linux platform. You will need to adjust the build tools and options to match your particular platforms.

1. Download, install, and test a toolchain that can build the desired target binaries. In this example, we use the `arm-linux-gnueabi-gcc` package.

2. Decide on your configuration options, and set up your environment to use the appropriate compiler and compiler tools. It is often easiest to set this up as a small shell script. For example:

```
#!/bin/sh -f
env \
CC=/usr/bin/arm-linux-gnueabi-gcc \
STRIP=/usr/bin/arm-linux-gnueabi-strip \
./dist/configure \
--build=i686-pc-linux-gnu \
--host=arm-linux-gnueabi \
--enable-java \
--enable-sql \
--enable-jdbc \
--enable-smallbuild \
--enable-shared \
--enable-stripped_messages \
--prefix=$HOME/ARM-linux/install \
$*
```

The `--build` flag indicates the system on which you are compiling and the `--host` flag indicates the target platform. Adjust or omit the `--enable` and `--prefix` flag settings as dictated by your own needs.

3. Unpack your Berkeley DB distribution and go to the `build_unix` directory. Alternatively, you may create a more appropriately-named directory at the same level and build in it. For example, `build_arm`.
4. In your build directory, configure by executing your script, then `make` and `make install`. Review any compiler warnings and fix if necessary, repeating the `configure` and `make` steps until you are satisfied. The resulting libraries and executables can then be transferred to your target system.

Installing Berkeley DB

Berkeley DB installs the following files into the following locations, with the following default values:

Configuration Variables	Default value
<code>--prefix</code>	<code>/usr/local/BerkeleyDB.Major.Minor</code>
<code>--exec_prefix</code>	<code>\$(prefix)</code>
<code>--bindir</code>	<code>\$(exec_prefix)/bin</code>
<code>--includedir</code>	<code>\$(prefix)/include</code>
<code>--libdir</code>	<code>\$(exec_prefix)/lib</code>

Configuration Variables	Default value
docdir	\$(prefix)/docs
Files	Default location
include files	\$(includedir)
libraries	\$(libdir)
utilities	\$(bindir)
documentation	\$(docdir)

With one exception, this follows the GNU Autoconf and GNU Coding Standards installation guidelines; please see that documentation for more information and rationale.

The single exception is the Berkeley DB documentation. The Berkeley DB documentation is provided in HTML format, not in UNIX-style man or GNU info format. For this reason, Berkeley DB configuration does not support `--infodir` or `--mandir`. To change the default installation location for the Berkeley DB documentation, modify the Makefile variable, `docdir`.

When installing Berkeley DB on filesystems shared by machines of different architectures, please note that although Berkeley DB include files are installed based on the value of `$ (prefix)`, rather than `$(exec_prefix)`, the Berkeley DB include files are not always architecture independent.

To move the entire installation tree to somewhere besides `/usr/local`, change the value of `prefix`.

To move the binaries and libraries to a different location, change the value of `exec_prefix`. The values of `includedir` and `libdir` may be similarly changed.

Any of these values except for `docdir` may be set as part of the configuration:

```
prompt: ../dist/configure --bindir=/usr/local/bin
```

Any of these values, including `docdir`, may be changed when doing the install itself:

```
prompt: make prefix=/usr/contrib/bdb install
```

The Berkeley DB installation process will attempt to create any directories that do not already exist on the system.

Dynamic shared libraries

Warning: the following information is intended to be generic and is likely to be correct for most UNIX systems. Unfortunately, dynamic shared libraries are not standard between UNIX systems, so there may be information here that is not correct for your system. If you have problems, consult your compiler and linker manual pages, or your system administrator.

The Berkeley DB dynamic shared libraries are created with the name `libdb-major.minor.so`, where `major` is the major version number and `minor` is the minor version number. Other shared libraries are created if Java and Tcl support are enabled: specifically, `libdb_java-major.minor.so` and `libdb_tcl-major.minor.so`.

On most UNIX systems, when any shared library is created, the linker stamps it with a "SONAME". In the case of Berkeley DB, the SONAME is `libdb-major.minor.so`. It is important to realize that applications linked against a shared library remember the SONAMEs of the libraries they use and not the underlying names in the filesystem.

When the Berkeley DB shared library is installed, links are created in the install `lib` directory so that `libdb-major.minor.so`, `libdb-major.so`, and `libdb.so` all refer to the same library. This library will have an SONAME of `libdb-major.minor.so`.

Any previous versions of the Berkeley DB libraries that are present in the install directory (such as `libdb-2.7.so` or `libdb-2.so`) are left unchanged. (Removing or moving old shared libraries is one drastic way to identify applications that have been linked against those vintage releases.)

Once you have installed the Berkeley DB libraries, unless they are installed in a directory where the linker normally looks for shared libraries, you will need to specify the installation directory as part of compiling and linking against Berkeley DB. Consult your system manuals or system administrator for ways to specify a shared library directory when compiling and linking applications with the Berkeley DB libraries. Many systems support environment variables (for example, `LD_LIBRARY_PATH` or `LD_RUN_PATH`), or system configuration files (for example, `/etc/ld.so.conf`) for this purpose.

Warning: some UNIX installations may have an already existing `/usr/lib/libdb.so`, and this library may be an incompatible version of Berkeley DB.

We recommend that applications link against `libdb.so` (for example, using `-ldb`). Even though the linker uses the file named `libdb.so`, the executable file for the application remembers the library's SONAME (`libdb-major.minor.so`). This has the effect of marking the applications with the versions they need at link time. Because applications locate their needed SONAMEs when they are executed, all previously linked applications will

continue to run using the library they were linked with, even when a new version of Berkeley DB is installed and the file `libdb.so` is replaced with a new version.

Applications that know they are using features specific to a particular Berkeley DB release can be linked to that release. For example, an application wanting to link to Berkeley DB major release "3" can link using `-ldb-3`, and applications that know about a particular minor release number can specify both major and minor release numbers; for example, `-ldb-3.5`.

If you want to link with Berkeley DB before performing library installation, the "make" command will have created a shared library object in the `.libs` subdirectory of the build directory, such as `build_unix/.libs/libdb-major.minor.so`. If you want to link a file against this library, with, for example, a major number of "3" and a minor number of "5", you should be able to do something like the following:

```
cc -L BUILD_DIRECTORY/.libs -o testprog testprog.o -ldb-3.5
env LD_LIBRARY_PATH="BUILD_DIRECTORY/.libs:$LD_LIBRARY_PATH" ./testprog
```

where `BUILD_DIRECTORY` is the full directory path to the directory where you built Berkeley DB.

The `libtool` program (which is configured in the build directory) can be used to set the shared library path and run a program. For example, the following runs the `gdb` debugger on the `db_dump` utility after setting the appropriate paths:

```
libtool gdb db_dump
```

Libtool may not know what to do with arbitrary commands (it is hardwired to recognize "gdb" and some other commands). If it complains the mode argument will usually resolve the problem:

```
libtool --mode=execute my_debugger db_dump
```

On most systems, using `libtool` in this way is exactly equivalent to setting the `LD_LIBRARY_PATH` environment variable and then executing the program. On other systems, using `libtool` has the virtue of knowing about any other details on systems that don't behave in this typical way.

Running the test suite under UNIX

[Building SQL Test Suite on Unix](#)

The Berkeley DB test suite is built if you specify `--enable-test` as an argument when configuring Berkeley DB. The test suite also requires that you configure and build the Tcl interface to the library.

Before running the tests for the first time, you may need to edit the `include.tcl` file in your build directory. The Berkeley DB configuration assumes that you intend to use the version of the `tclsh` utility included in the Tcl installation with which Berkeley DB was configured to run the test suite, and further assumes that the test suite will be run with the libraries prebuilt in the Berkeley DB build directory. If either of these assumptions are incorrect, you will need to edit the `include.tcl` file and change the following line to correctly specify the full path to the version of `tclsh` with which you are going to run the test suite:

```
set tclsh_path ...
```

You may also need to change the following line to correctly specify the path from the directory where you are running the test suite to the location of the Berkeley DB Tcl library you built:

```
set test_path ...
```

It may not be necessary that this be a full path if you have configured your system's shared library mechanisms to search the directory where you built or installed the Tcl library.

All Berkeley DB tests are run from within `tclsh`. After starting `tclsh`, you must source the file `test.tcl` in the test directory. For example, if you built in the `build_unix` directory of the distribution, this would be done using the following command:

```
% source ../test/tcl/test.tcl
```

If no errors occur, you should get a `"%"` prompt.

You are now ready to run tests in the test suite; see [Running the test suite](#) for more information.

Building SQL Test Suite on Unix

The Berkeley DB SQL interface test suite is built if you specify `--enable-test` and `--enable-sql` as arguments, when configuring Berkeley DB. The test suite also requires that you build the Berkeley DB Tcl API.

```
../dist/configure --enable-sql --enable-test --with-tcl=/usr/lib
```

This builds the testfixture project in `../build_unix/sql`.

To enable extensions like full text search layer and R-Tree layer in the SQL test suite, configure with `--enable-amalgamation`.

Architecture independent FAQ

1. I have gcc installed, but configure fails to find it.

Berkeley DB defaults to using the native C compiler if none is specified. That is usually "cc", but some platforms require a different compiler to build multithreaded code. To configure Berkeley DB to build with gcc, run configure as follows:

```
env CC=gcc ../dist/configure ...
```

2. When compiling with gcc, I get unreferenced symbols; for example the following:

```
symbol __muldi3: referenced symbol not found
symbol __cmpdi2: referenced symbol not found
```

Berkeley DB often uses 64-bit integral types on systems supporting large files, and gcc performs operations on those types by calling library functions. These unreferenced symbol errors are usually caused by linking an application by calling "ld" rather than by calling "gcc": gcc will link in libgcc.a and will resolve the symbols. If that does not help, another possible workaround is to reconfigure Berkeley DB using the `--disable-largefile` configuration option and then rebuild.

3. My C++ program traps during a failure in a DB call on my gcc-based system.

We believe there are some severe bugs in the implementation of exceptions for some gcc compilers. Exceptions require some interaction between compiler, assembler, and runtime libraries. We're not sure exactly what is at fault, but one failing combination is gcc 2.7.2.3 running on SuSE Linux 6.0. The problem on this system can be seen with a rather simple test case of an exception thrown from a shared library and caught in the main program.

A variation of this problem seems to occur on AIX, although we believe it does not necessarily involve shared libraries on that platform.

If you see a trap that occurs when an exception might be thrown by the Berkeley DB runtime, we suggest that you use static libraries instead of shared libraries. See the documentation for configuration. If this doesn't work and you have a choice of compilers, try using a more recent gcc- or a non-gcc based compiler to build Berkeley DB.

Finally, you can disable the use of exceptions in the C++ runtime for Berkeley DB by using the [DB_CXX_NO_EXCEPTIONS](#) flag with the [DbEnv](#) or [Db](#) constructors. When this flag is on, all C++ methods fail by returning an error code rather than throwing an exception.

4. I get unexpected results and database corruption when running threaded programs.

I get error messages that mutex (for example, `pthread_mutex_XXX` or `mutex_XXX`) functions are undefined when linking applications with Berkeley DB.

On some architectures, the Berkeley DB library uses the ISO POSIX standard pthreads and UNIX International (UI) threads interfaces for underlying mutex support; Solaris is an example. You can specify compilers or compiler flags, or link with the appropriate thread library when loading your application to resolve the undefined references:

```
cc ... -lpthread ...
cc ... -lthread ...
xlc_r ...
cc ... -mt ...
```

See the appropriate architecture-specific Reference Guide pages for more information.

On systems where more than one type of mutex is available, it may be necessary for applications to use the same threads package from which Berkeley DB draws its mutexes. For example, if Berkeley DB was built to use the POSIX pthreads mutex calls for mutex support, the application may need to be written to use the POSIX pthreads interfaces for its threading model. This is only conjecture at this time, and although we know of no systems that actually have this requirement, it's not unlikely that some exist.

In a few cases, Berkeley DB can be configured to use specific underlying mutex interfaces. You can use the [--enable-posixmutexes](#) and [--enable-uimutexes](#) configuration options to specify the POSIX and Unix International (UI) threads packages. This should not, however, be necessary in most cases.

In some cases, it is vitally important to make sure that you load the correct library. For example, on Solaris systems, there are POSIX pthread interfaces in the C library, so applications can link Berkeley DB using only C library and not see any undefined symbols. However, the C library POSIX pthread mutex support is insufficient for Berkeley DB, and Berkeley DB cannot detect that fact. Similar errors can arise when applications (for example, `tcsh`) use `dlopen` to dynamically load Berkeley DB as a library.

If you are seeing problems in this area after you confirm that you're linking with the correct libraries, there are two other things you can try. First, if your platform supports interlibrary dependencies, we recommend that you change the Berkeley DB Makefile to specify the appropriate threads library when creating the Berkeley DB shared library, as an interlibrary dependency. Second, if your application is using dlopen to dynamically load Berkeley DB, specify the appropriate thread library on the link line when you load the application itself.

5. I get core dumps when running programs that fork children.

Berkeley DB handles should not be shared across process forks, each forked child should acquire its own Berkeley DB handles.

6. I get reports of uninitialized memory reads and writes when running software analysis tools (for example, Rational Software Corp.'s Purify tool).

For performance reasons, Berkeley DB does not write the unused portions of database pages or fill in unused structure fields. To turn off these errors when running software analysis tools, build with the `--enable-umrw` configuration option.

7. Berkeley DB programs or the test suite fail unexpectedly.

The Berkeley DB architecture does not support placing the shared memory regions on remote filesystems -- for example, the Network File System (NFS) or the Andrew File System (AFS). For this reason, the shared memory regions (normally located in the database home directory) must reside on a local filesystem. See [Shared memory region](#) for more information.

With respect to running the test suite, always check to make sure that TESTDIR is not on a remote mounted filesystem.

8. The `db_dump` utility fails to build.

The `db_dump185` utility is the utility that supports the conversion of Berkeley DB 1.85 and earlier databases to current database formats. If the build errors look something like the following, it means the `db.h` include file being loaded is not a Berkeley DB 1.85 version include file:

```
db_dump185.c: In function `main':
db_dump185.c:210: warning: assignment makes pointer from integer
without a cast
db_dump185.c:212: warning: assignment makes pointer from integer
without a cast
db_dump185.c:227: structure has no member named `seq'
db_dump185.c:227: `R_NEXT' undeclared (first use in this function)
```

If the build errors look something like the following, it means that the Berkeley DB 1.85 code was not found in the standard libraries:

```
cc -o db_dump185 db_dump185.o
ld:
Unresolved:
dbopen
```

To build the [db_dump185](#) utility, the Berkeley DB version 1.85 code must already been built and available on the system. If the Berkeley DB 1.85 header file is not found in a standard place, or if the library is not part of the standard libraries used for loading, you will need to edit your Makefile, and change the following lines:

```
DB185INC=
DB185LIB=
```

So that the system Berkeley DB 1.85 header file and library are found; for example:

```
DB185INC=/usr/local/include
DB185LIB=-ldb185
```

AIX

1. I can't compile and run multithreaded applications.

Special compile-time flags are required when compiling threaded applications on AIX. If you are compiling a threaded application, you must compile with the `_THREAD_SAFE` flag and load with specific libraries; for example, `-lc_r`. Specifying the compiler name with a trailing `_r` usually performs the right actions for the system.

```
xlc_r ...
cc -D_THREAD_SAFE -lc_r ...
```

The Berkeley DB library will automatically build with the correct options.

2. I can't run using the [DB_SYSTEM_MEM](#) option to [DB_ENV->open\(\)](#).

AIX 4.1 allows applications to map only 10 system shared memory segments. In AIX 4.3, this has been raised to 256K segments, but only if you set the environment variable `"export EXTSHM=ON"`.

3. On AIX 4.3.2 (or before) I see duplicate symbol warnings when building the C++ shared library and when linking applications.

We are aware of some duplicate symbol warnings with this platform, but they do not appear to affect the correct operation of applications.

4. On AIX 4.3.3 I see undefined symbols for `DbEnv::set_error_stream`, `Db::set_error_stream` or `DbEnv::verify` when linking C++ applications. (These undefined symbols also appear when building the Berkeley DB C++ example applications).

By default, Berkeley DB is built with `_LARGE_FILES` set to 1 to support the creation of "large" database files. However, this also affects how standard classes, like `iostream`, are named internally. When building your application, use a `"-D_LARGE_FILES=1"` compilation option, or insert `"#define _LARGE_FILES 1"` before any `#include` statements.

5. I can't create database files larger than 1GB on AIX.

If you're running on AIX 4.1 or earlier, try changing the source code for `os/os_open.c` to always specify the `O_LARGEFILE` flag to the `open(2)` system call, and recompile Berkeley DB from scratch.

Also, the documentation for the IBM Visual Age compiler states that it does not support the 64-bit filesystem APIs necessary for creating large files; the `ibmcxx` product must be used instead. We have not heard whether the GNU `gcc` compiler supports the 64-bit APIs or not.

Finally, to create large files under AIX, the filesystem has to be configured to support large files and the system wide user hard-limit for file sizes has to be greater than 1GB.

6. I see errors about "open64" when building Berkeley DB applications.

System include files (most commonly `fcntl.h`) in some releases of AIX and Solaris redefine "open" when large-file support is enabled for applications. This causes problems when compiling applications because "open" is a method in the Berkeley DB APIs. To work around this problem:

1. Avoid including the problematical system include files in source code files which also include Berkeley DB include files and call into the Berkeley DB API.
2. Before building Berkeley DB, modify the generated include file `db.h` to itself include the problematical system include files.
3. Turn off Berkeley DB large-file support by specifying the [--disable-largefile](#) configuration option and rebuilding.

7. I see the error "Redeclaration of `lseek64`" when building Berkeley DB with the `--enable-sql` and `--enable-test` options.

In some releases of AIX, the system include files (most commonly `unistd.h`) redefine `lseek` to `lseek64` when large-file support is enabled even though `lseek` may have already been defined when the `_LARGE_FILE_API` macro is on. To work around this problem, do either one of the following:

1. Disable large-file support in Berkeley DB by specifying the `--disable-largefile` configuration option and rebuilding.
2. Edit `db.h` manually after running the configure command, and remove the line that includes `unistd.h`.

FreeBSD

1. I can't compile and run multithreaded applications.

Special compile-time flags are required when compiling threaded applications on FreeBSD. If you are compiling a threaded application, you must compile with the `_THREAD_SAFE` and `-pthread` flags:

```
cc -D_THREAD_SAFE -pthread ...
```

The Berkeley DB library will automatically build with the correct options.

2. I see `fsync` and `close` system call failures when accessing databases or log files on NFS-mounted filesystems.

Some FreeBSD releases are known to return `ENOLCK` from `fsync` and `close` calls on NFS-mounted filesystems, even though the call has succeeded. The Berkeley DB code should be modified to ignore `ENOLCK` errors, or no Berkeley DB files should be placed on NFS-mounted filesystems on these systems.

Apple iOS (iPhone OS)

Building Berkeley DB in Apple iOS (known as iPhone OS previously) is the same as building for a conventional UNIX platform. This section lists the commands for building Berkeley DB in both the iPhone simulator (a software simulator included in the iPhone SDK that you can use to test your application without using the iPhone/iPod Touch) and the iPhone device.

Prior to building BDB in an iPhone simulator/iPhone device, set the required environment variables for iOS (iPhone OS). For a simulator/iPhone version 4.2 or older, set the `LDFLAGS` variable as follows:

```
export LDFLAGS="-L$SDKROOT/usr/lib/"
```

Otherwise, set LDFLAGS as follows:

```
export LDFLAGS="-L$SDKROOT/usr/lib/system/"
```

The steps to build BDB in an iPhone simulator are as follows:

```
export CFLAGS="-arch i386 -pipe -no-cpp-precomp --sysroot=$SDKROOT"
export CXXFLAGS="-arch i386 -pipe -no-cpp-precomp --sysroot=$SDKROOT"
cd $BDB_HOME/build_unix
../dist/configure --host=i386-apple-darwin\
    --prefix=$SDKROOT ...
make
```

The steps to build BDB in an iPhone device are as follows:

```
export CFLAGS="-arch armv6 -pipe -Os -gdwarf-2\
    -no-cpp-precomp -mthumb -isysroot $SDKROOT "
export CXXFLAGS="-arch armv6 -pipe -Os -gdwarf-2\
    -no-cpp-precomp -mthumb -isysroot $SDKROOT "
cd $BDB_HOME/build_unix
../dist/configure --host=arm-apple-darwin9\
    --prefix=$SDKROOT ...
make
```

Both sets of commands create the BDB dynamic library - libdb-5.3.dylib. To build the static library, libdb-5.3.a, add the `--enable-shared=no` option while configuring.

IRIX

1. I can't compile and run multithreaded applications.

Special compile-time flags are required when compiling threaded applications on IRIX. If you are compiling a threaded application, you must compile with the `_SGI_MP_SOURCE` flag:

```
cc -D_SGI_MP_SOURCE ...
```

The Berkeley DB library will automatically build with the correct options.

Linux

1. I can't compile and run multithreaded applications.

Special compile-time flags are required when compiling threaded applications on Linux. If you are compiling a threaded application, you must compile with the `_REENTRANT` flag:

```
cc -D_REENTRANT ...
```

The Berkeley DB library will automatically build with the correct options.

2. I see database corruption when accessing databases.

Some Linux filesystems do not support POSIX filesystem semantics. Specifically, ext2 and early releases of ReiserFS, and ext3 in some configurations, do not support "ordered data mode" and may insert random data into database or log files when systems crash. Berkeley DB files should not be placed on a filesystem that does not support, or is not configured to support, POSIX semantics.

3. What scheduler should I use?

In some Linux kernels you can select schedulers, and the default is the "anticipatory" scheduler. We recommend not using the "anticipatory" scheduler for transaction processing workloads.

Mac OS X

1. When trying to link multiple Berkeley DB language interfaces (for example, Tcl, C++, Java, Python) into a single process, I get "multiple definitions" errors from dyld.

To fix this problem, set the environment variable `MACOSX_DEPLOYMENT_TARGET` to 10.3 (or your current version of OS X), and reconfigure and rebuild Berkeley DB from scratch. See the OS X `ld(1)` and `dyld(1)` man pages for information about how OS X handles symbol namespaces, as well as undefined and multiply-defined symbols.

2. When trying to use system-backed shared memory on OS X I see failures about "too many open files".

The default number of shared memory segments on OS X is too low. To fix this problem, edit the file `/etc/rc`, changing the `kern.sysv.shmmax` and `kern.sysv.shmseg` values as follows:

```
*** /etc/rc.orig      Fri Dec 19 09:34:09 2003
--- /etc/rc          Fri Dec 19 09:33:53 2003
*****
*** 84,93 ****
    # System tuning
    sysctl -w kern.maxvnodes=$(echo $(sysctl -n hw.physmem) '33554432 /
```

```

512 * 1024 +p'|dc)
! sysctl -w kern.sysv.shmmax=4194304
sysctl -w kern.sysv.shmmin=1
sysctl -w kern.sysv.shmmni=32
! sysctl -w kern.sysv.shmseg=8
sysctl -w kern.sysv.shmall=1024
if [ -f /etc/sysctl-macosxserver.conf ]; then
    awk '{ if (!-1 && -1) print $1 }' <
/etc/sysctl-macosxserver.conf | while read
--- 84,93 ---
# System tuning
sysctl -w kern.maxvnodes=$(echo $(sysctl -n hw.physmem) '33554432 /
512 * 1024 +p'|dc)
! sysctl -w kern.sysv.shmmax=134217728
sysctl -w kern.sysv.shmmin=1
sysctl -w kern.sysv.shmmni=32
! sysctl -w kern.sysv.shmseg=32
sysctl -w kern.sysv.shmall=1024
if [ -f /etc/sysctl-macosxserver.conf ]; then
    awk '{ if (!-1 && -1) print $1 }' <
/etc/sysctl-macosxserver.conf | while read

```

and then reboot the system.

QNX

1. To what versions of QNX has DB been ported?

Berkeley DB has been ported to the QNX Neutrino technology which is commonly referred to as QNX RTP (Real-Time Platform). Berkeley DB has not been ported to earlier versions of QNX, such as QNX 4.25.

2. Building Berkeley DB shared libraries fails.

The `/bin/sh` utility distributed with some QNX releases drops core when running the GNU libtool script (which is used to build Berkeley DB shared libraries). There are two workarounds for this problem: First, only build static libraries. You can disable building shared libraries by specifying the configuration flag when configuring Berkeley DB.

Second, build Berkeley DB using an alternate shell. QNX distributions include an accessories disk with additional tools. One of the included tools is the GNU bash shell, which is able to run the libtool script. To build Berkeley DB using an alternate shell, move `/bin/sh` aside, link or copy the alternate shell into that location, configure, build and install Berkeley DB, and then replace the original shell utility.

3. Are there any QNX filesystem issues?

Berkeley DB generates temporary files for use in transactionally protected file system operations. Due to the filename length limit of 48 characters in the QNX filesystem, applications that are using transactions should specify a database name that is at most 43 characters.

4. What are the implications of QNX's requirement to use `shm_open(2)` in order to use `mmap(2)`?

QNX requires that files mapped with `mmap(2)` be opened using `shm_open(2)`. There are other places in addition to the environment shared memory regions, where Berkeley DB tries to memory map files if it can.

The memory pool subsystem normally attempts to use `mmap(2)` even when using private memory, as indicated by the `DB_PRIVATE` flag to `DB_ENV->open()`. In the case of QNX, if an application is using private memory, Berkeley DB will not attempt to map the memory and will instead use the local cache.

5. What are the implications of QNX's mutex implementation using microkernel resources?

On QNX, the primitives implementing mutexes consume system resources. Therefore, if an application unexpectedly fails, those resources could leak. Berkeley DB solves this problem by always allocating mutexes in the persistent shared memory regions. Then, if an application fails, running recovery or explicitly removing the database environment by calling the `DB_ENV->remove()` method will allow Berkeley DB to release those previously held mutex resources. If an application specifies the `DB_PRIVATE` flag (choosing not to use persistent shared memory), and then fails, mutexes allocated in that private memory may leak their underlying system resources. Therefore, the `DB_PRIVATE` flag should be used with caution on QNX.

6. The make clean command fails to execute when building the Berkeley DB SQL interface.

Remove the build directory manually to clean up and proceed.

SCO

1. If I build with gcc, programs such as `db_dump` and `db_stat` core dump immediately when invoked.

We suspect gcc or the runtime loader may have a bug, but we haven't tracked it down. If you want to use gcc, we suggest building static libraries.

Solaris

1. I can't compile and run multithreaded applications.

Special compile-time flags and additional libraries are required when compiling threaded applications on Solaris. If you are compiling a threaded application, you must compile with the `D_REENTRANT` flag and link with the `libpthread.a` or `libthread.a` libraries:

```
cc -mt ...
cc -D_REENTRANT ... -lthread
cc -D_REENTRANT ... -lpthread
```

The Berkeley DB library will automatically build with the correct options.

2. I've installed gcc on my Solaris system, but configuration fails because the compiler doesn't work.

On some versions of Solaris, there is a `cc` executable in the user's path, but all it does is display an error message and fail:

```
% which cc
/usr/ucb/cc
% cc
/usr/ucb/cc: language optional software package not installed
```

Because Berkeley DB always uses the native compiler in preference to `gcc`, this is a fatal error. If the error message you are seeing is the following, then this may be the problem:

```
checking whether the C compiler (cc -O) works... no
configure: error: installation or configuration problem: C compiler
cannot create executables.
```

The simplest workaround is to set your `CC` environment variable to the system compiler and reconfigure; for example:

```
env CC=gcc ../dist/configure
```

If you are using the `--configure-cxx` option, you may also want to specify a C++ compiler, for example the following:

```
env CC=gcc CCC=g++ ../dist/configure
```

3. I see the error "libc internal error: _rmutex_unlock: rmutex not held", followed by a core dump when running threaded or JAVA programs.

This is a known bug in Solaris 2.5 and it is fixed by Sun patch 103187-25.

4. I see error reports of nonexistent files, corrupted metadata pages and core dumps.

Solaris 7 contains a bug in the threading libraries (-lpthread, -lthread), which causes the wrong version of the pwrite routine to be linked into the application if the thread library is linked in after the C library. The result will be that the pwrite function is called rather than the pwrite64. To work around the problem, use an explicit link order when creating your application.

Sun Microsystems is tracking this problem with Bug Id's 4291109 and 4267207, and patch 106980-09 to Solaris 7 fixes the problem:

```
Bug Id: 4291109
Duplicate of: 4267207
Category: library
Subcategory: libthread
State: closed
Synopsis: pwrite64 mapped to pwrite
Description:
When libthread is linked after libc, there is a table of functions in
libthread that gets "wired into" libc via _libc_threads_interface().
The table in libthread is wrong in both Solaris 7 and on28_35 for the
TI_PWRITE64 row (see near the end).
```

5. I see corrupted databases when doing hot backups or creating a hot failover archive.

The Solaris cp utility is implemented using the mmap system call, and so writes are not blocked when it reads database pages. See [Berkeley DB recoverability](#) for more information.

6. Performance is slow and the application is doing a lot of I/O to the disk on which the database environment's files are stored.

By default, Solaris periodically flushes dirty blocks from memory-mapped files to the backing filesystem. This includes the Berkeley DB database environment's shared memory regions and can affect Berkeley DB performance. Workarounds include creating the shared regions in system shared memory ([DB_SYSTEM_MEM](#)) or application private memory ([DB_PRIVATE](#)), or configuring Solaris to not flush memory-mapped pages. For more information, see the "Solaris Tunable Parameters Reference Manual: fsflush and Related Tunables".

7. I see errors about "open64" when building Berkeley DB applications.

System include files (most commonly `fcntl.h`) in some releases of AIX and Solaris redefine "open" when large-file support is enabled for applications. This causes problems when compiling applications because "open" is a method in the Berkeley DB APIs. To work around this problem:

1. Avoid including the problematical system include files in source code files which also include Berkeley DB include files and call into the Berkeley DB API.
2. Before building Berkeley DB, modify the generated include file `db.h` to itself include the problematical system include files.
3. Turn off Berkeley DB large-file support by specifying the `--disable-largefile` configuration option and rebuilding.

SunOS

1. I can't specify the `DB_SYSTEM_MEM` flag to `DB_ENV->open()`.

The `shmget` (2) interfaces are not used on SunOS releases prior to 5.0, even though they apparently exist, because the distributed include files did not allow them to be compiled. For this reason, it will not be possible to specify the `DB_SYSTEM_MEM` flag to those versions of SunOS.

Chapter 9. Upgrading Berkeley DB 11.2.5.2 applications to Berkeley DB 11.2.5.3

Table of Contents

[Introduction](#)

[Changes to the `build\_windows` Folder](#)

[Replication Connection Status in the Java API](#)

[New Function](#)

[New Class](#)

[Deprecated Function](#)

[Exclusive Database Handles](#)

[New Functions](#)

[Configure the Region Size of Heap Databases](#)

[New Functions](#)

New Hotbackup Interface

New Functions

Flags Accepted by DB_ENV->backup()

Flags Accepted by DB_ENV->dbbackup()

Enumerations Accepted by DB_ENV->set_backup_config()

Updated JDBC Version

Configure Directory to Store Metadata Files

New Functions

Changes in the SQL API Build

New Berkeley DB SQL API PRAGMAs

New PRAGMAs

Replication for Existing Databases in the SQL API

PRAGMAs With Permanent Effects

PRAGMAs That Can Now Operate on Existing Databases

Berkeley DB X/Open Compliant XA Resource Manager and Transaction Snapshots

Berkeley DB Library Version 11.2.5.3 Change Log

Changes between 11.2.5.3.21 and 11.2.5.3.28

Changes between 11.2.5.3.15 and 11.2.5.3.21

Database or Log File On-Disk Format Changes

New Features

Database Environment Changes

Access Method Changes

SQL API Changes

Java-specific API changes

Replication Changes

[Locking Subsystem Changes](#)[Logging Subsystem Changes](#)[Memory Pool Subsystem Changes](#)[Mutex Subsystem Changes](#)[Transaction Subsystem Changes](#)[Utility Changes](#)[Configuration, Documentation, Sample Apps, Portability and Build Changes](#)[Known Bugs](#)

Introduction

The following pages describe how to upgrade applications coded against the Berkeley DB 11.2.5.2 library to the Berkeley DB 11.2.5.3 library (both of which belong to Berkeley DB 11g Release 2). This information does not describe how to upgrade Berkeley DB 1.85 release applications.

For information on the general process of upgrading Berkeley DB installations and upgrade instructions related to historical releases, see the [Berkeley DB Upgrade Guide](#).

Changes to the build_windows Folder

Visual Studios projects in the `build_windows` folder have been moved into two new sub folders `build_windows/VS8` and `build_windows/VS10`. Visual Studios 2010 projects have been moved into the folder `VS10`, and Visual Studios 2008 projects have been moved into the folder `VS8`.

Replication Connection Status in the Java API

[New Function](#)[New Class](#)[Deprecated Function](#)

The Java function `ReplicationManagerSiteInfo.isConnected()` is now deprecated. To get the replication connection status, use `ReplicationManagerSiteInfo.getConnectionStatus()`, which

returns the new class `ReplicationManagerConnectionStatus`, which has the values `CONNECTED`, `DISCONNECTED`, and `UNKNOWN`.

New Function

- `ReplicationManagerSiteInfo.getConnectionStatus()`

New Class

- `ReplicationManagerConnectionStatus`

Deprecated Function

- `ReplicationManagerSiteInfo.isConnected()`

Exclusive Database Handles

New Functions

Database handles can now be configured to allow exclusive access to the database. To enable exclusive access, call `DB->set_lk_exclusive()` before calling `DB->open()`. Set `nowait_onoff` to non-zero to have `DB->open()` return immediately, with the error `DB_LOCK_NOTGRANTED` if it cannot immediately get exclusive access to the database, and to 0 to have `DB->open()` block until it can gain exclusive access.

New Functions

- `DB->set_lk_exclusive()`
- `DB->get_lk_exclusive()`

Configure the Region Size of Heap Databases

New Functions

The region size of heap databases is now configurable. Configuring the region size is useful in controlling the growth of a heap database. To set the region size, call `DB->set_heap_regionsize()` with the number of pages that the region should have, before the database is created. The function is ignored if it is called after the database is created.

New Functions

- `DB->set_heap_regionsize()`
- `DB->get_heap_regionsize()`

New Hotbackup Interface

New Functions

Flags Accepted by `DB_ENV->backup()`

Flags Accepted by `DB_ENV->dbbackup()`

Enumerations Accepted by `DB_ENV->set_backup_config()`

Two new functions have been added to the API that perform hotbackups, `DB_ENV->backup()` and `DB_ENV->dbbackup()`. `DB_ENV->backup()` creates a hotbackup of all databases in the specified environment, and `DB_ENV->dbbackup()` creates a hotbackup of the specified database. The functions `DB_ENV->set_backup_callbacks()` and `DB_ENV->set_backup_config()` can be called to customize the behavior of hotbackup. Note that this interface must be used to create a hotbackup on all non-BSD or Unix based systems.

New Functions

- `DB_ENV->backup()`
- `DB_ENV->dbbackup()`
- `DB_ENV->set_backup_callbacks()`
- `DB_ENV->set_backup_config()`

Flags Accepted by `DB_ENV->backup()`

- `DB_BACKUP_CLEAN`
- `DB_BACKUP_FILES`
- `DB_BACKUP_NO_LOGS`
- `DB_BACKUP_SINGLE_DIR`
- `DB_BACKUP_UPDATE`
- `DB_CREATE`
- `DB_EXCL`
- `DB_VERB_BACKUP`

Flags Accepted by `DB_ENV->dbbackup()`

- `DB_CREATE`
- `DB_EXCL`

Enumerations Accepted by `DB_ENV->set_backup_config()`

- `DB_BACKUP_WRITE_DIRECT`
- `DB_BACKUP_READ_COUNT`
- `DB_BACKUP_READ_SLEEP`

- `DB_BACKUP_SIZE`

Updated JDBC Version

The JDBC library version included with Berkeley DB has been updated to version 20110827. The new version supports the Embedded Java JSR 169 standard, which is the official specification for JDBC with embedded Java platforms.

Configure Directory to Store Metadata Files

New Functions

The directory in which persistent metadata files are stored can now be configured. By default persistent metadata files are stored in the environment home directory. The files that will be stored in the metadata directory are `__db.rep.system`, `__db.rep.gen`, `__db.rep.egen` and `__db.rep.init`.

To set the metadata file directory, call `DB_ENV->set_metadata_dir()` with the path to the directory in which to store metadata files. The metadata directory can also be set in the `DB_CONFIG` file using `set_metadata_dir`.

New Functions

- `DB_ENV->set_metadata_dir()`
- `DB_ENV->get_metadata_dir()`

Changes in the SQL API Build

Several changes have been made to the SQL API build. Encryption is now disabled by default on Windows, and the compile time flag `BDBSQL_OMIT_SHARING` has been changed to `BDBSQL_SINGLE_PROCESS`.

Encryption is now disabled by default in the Windows SQL API build. This makes it consistent with builds of the SQL API on other systems. To enable encryption in Visual Studios right click the `db_sql` project, and select `Properties->Configuration Properties->C/C++->Preprocessor` and add `SQLITE_HAS_CODEC` to `Preprocessor Definitions`.

New Berkeley DB SQL API PRAGMAs

New PRAGMAs

Two new Berkeley DB SQL API specific pragmas have been added, `bdbsql_shared_resources` and `bdbsql_set_lock_tablesize`. `bdbsql_shared_resources` is used to set the maximum amount of memory, in bytes, to be used by shared structures in the main environment region, which is useful in applications with a large number of tables, transactions, or threads.

`bdbsql_set_lock_tablesize` is used to set the number of buckets in the lock object hash table in the Berkeley DB environment, which is useful if an application has many concurrent long running transactions.

New PRAGMAs

- `PRAGMA bdbsql_shared_resources[=N]`
- `PRAGMA bdbsql_set_lock_tablesize[=N]`

Replication for Existing Databases in the SQL API

PRAGMAs With Permanent Effects

PRAGMAs That Can Now Operate on Existing Databases

Replication can now be enabled on existing SQL databases, and replication is now disabled permanently instead of temporarily. Replication is enabled on an existing database the same way it is enabled on a new database, with one restriction. The existing database must configure itself as the initial master of a new replication group. To disable replication on a database permanently, use `pragma replication=OFF;`.

PRAGMAs With Permanent Effects

- `pragma replication=OFF;`

PRAGMAs That Can Now Operate on Existing Databases

- `pragma replication_local_site="host:port";`
- `pragma replication_initial_master=ON;`
- `pragma replication=ON;`

Berkeley DB X/Open Compliant XA Resource Manager and Transaction Snapshots

The transactions managed by the Berkeley DB X/open compliant XA resource manager can now be enabled for transaction snapshots. To enable snapshots open an XA managed database with the flag, `DB_MULTIVERSION`, and all XA managed transactions that operate on that database will use transaction snapshots.

Berkeley DB Library Version 11.2.5.3 Change Log

[Changes between 11.2.5.3.21 and 11.2.5.3.28](#)

[Changes between 11.2.5.3.15 and 11.2.5.3.21](#)

[Database or Log File On-Disk Format Changes](#)

[New Features](#)

[Database Environment Changes](#)

[Access Method Changes](#)

[SQL API Changes](#)

[Java-specific API changes](#)

[Replication Changes](#)

[Locking Subsystem Changes](#)

[Logging Subsystem Changes](#)

[Memory Pool Subsystem Changes](#)

[Mutex Subsystem Changes](#)

[Transaction Subsystem Changes](#)

[Utility Changes](#)

[Configuration, Documentation, Sample Apps, Portability and Build Changes](#)

[Known Bugs](#)

This is the changelog for Berkeley DB 11g Release 2 (library version 11.2.5.3).

Changes between 11.2.5.3.21 and 11.2.5.3.28

1. Fixed tcl library linking for AIX 7. [#17109]
2. Fixed a bug that could cause a trap if a lock timeout occurred while opening a database. [#21098]
3. Fixed missing encryption support for the Android JDBC driver. [#21129]

4. Fixed an incorrect message being displayed when the -l option was specified to the db_hotbackup utility. [#21313]
5. Fixed a bug that DB_ENV->log_get_config did not work correctly before DB_ENV->open. [#21359]
6. Fixed a bug that could cause a SQL build failure when FTS3 is enabled. [#21382]
7. Fixed a bug that prevented in-memory SQL database from being created properly. They can now be created without the use of the SQLITE_OPEN_CREATE flag. [#21456]
8. Fixed a memory leak in SQL online backup. [#21460]
9. Added additional examples for C++ [#21477]
10. Fixed a bug that could make odbc fail to build with sqlite source code. [#21490]
11. Fixed bugs in compaction of large keys in the upper levels of btrees. [#21569]
12. The db utilities (db_xxxx) no longer operate on replication clients that are being automatically initialized. The DB_REP_LOCKOUT error is now returned. [#21593]
13. Using DB_TXN_SNAPSHOT on an HA client will now result in an error. [#21601]
14. Fixed a bug that prevented a sub-database from being created under the directory identified in DB->set_create_dir. [#21603]
15. Fixed a race condition on a cursor when using a multi-threaded application with the SQL API. [#21714]
16. Fixed a race condition in the failchk code when cleaning up mutexes. [#21796]
17. Fixed a resource leak in the db->verify() function for btrees. The bug would slow down verification and possibly cause it to run out of memory. [#21917]
18. Removed a potential hang when compacting databases with many duplicates. [#21975]
19. Fixed a bug that caused a crash when reentering dbsql and specifying replication=on for a SQL database where replication was already enabled. [#22116]
20. Fixed an incorrect recursive call dealing with joins. [#22398]
21. Fix build failures in ado.net [#22405]

- 22. Fixed a bug that could cause a JDBC build failure on recent versions of Visual Studio. [#22497]
- 23. Correct build problems when building dbstl with gcc-4.7.3 [#22615]

Changes between 11.2.5.3.15 and 11.2.5.3.21

- 1. Fixed incompatibility problems of Java DPL with JDK7, so DPL will now work with JDK7. [#20586]
- 2. Added a flag to allow database locking to be disabled from the SQL API. [#20928]
- 3. Fixed a bug that could allocate a heap data page in a region after the region creation has been undone. [#20939]
- 4. Redundant whitespaces are now ignored in DB_CONFIG lines pertaining to directories, e.g. set_data_dir. [#20158]
- 5. Fixed a rare race condition that could cause a crash if two processes opened the same database at the same time. [#21041]
- 6. Fixed a bug that caused DB_ENV->backup to stop early if DB_BACKUP_FILES was not set and a non-DB file was in the data directory. [#21076]
- 7. Fixed missing cross compiling capability for the JDBC driver. [#21101]
- 8. Allow the same system/machine to host both a master and a replica database through the use of relative pathnames. [#21105]
- 9. Fixed a bug in the Java API where EnvironmentConfig.setCreateDir would fail to configure the environment. [#21127]
- 10. Fixed an assert failure in btreeCompare when allocating memory in the wrong thread was causing a memory leak. [#21232]
- 11. Fixed a bug in the Java API where concurrent operations that change the database schema could lead to a hang. [#21265]
- 12. Added JDBC code to the code base and updated the windows build files to include the JDBC solution. [#21294]
- 13. Fixed a bug where the heap's region size was not getting swapped correctly in mixed-endian environments. [#21295]
- 14. Fixed a bug in the db_sql_jdbc project file for vs2010 that was preventing it from building correctly. [#21332]

Database or Log File On-Disk Format Changes

1. Existing database file formats were unchanged in library version 11.2.5.3.
2. The log file format changed in library version 11.2.5.3.

New Features

1. Added support for verifying named in-memory dbs. [#16941]
2. Added an integer key comparison function to improve performance through the SQL API. [#19609]
3. Support build on the platforms where pthread_t is a struct. [#19876]
4. Added an API call so the user can specify the size of the region in a heap db. [#19914]
5. Improved Replication Manager's ability to recover from the (perhaps rare) phenomenon of two sites trying to connect to each other simultaneously, which used to result in loss of both connections, requiring a retry after the CONNECTION_RETRY timeout period. [#19980]
6. Enhanced the interface for copying databases for a hot backup. Added configure support for --enable-atomicfileread. [#20129]
7. Enhanced the log reading routine to detect that a log file is missing rather than returning that a zero length record was found. [#20130]
8. Added pragma bdbsql_shared_resources to set or report the maximum amount of memory to be used by shared structures in the main environment region and bdbsql_lock_tablesize to set or report the number of buckets in the lock object hash table. These are advanced tuning features for applications with large number of tables or needs to reduce locking on concurrent long running transactions. [#20156]
9. Added set_metadata_dir() and get_metadata_dir() to enable storage of persistent metadata files in a location other than the environment home directory. [#20174]
10. Improved the error handling through the SQL API. Errors can be sent to a file with the use of the BDBSQL_ERROR_FILE pragma. [#20213]
11. Database handles can now be configured to give exclusive access to the database. [#20331]
12. XA transactions will now use transaction snapshots if the XA databases they operate on were configured with DB_MULTIVERSION. [#20332]

13. Added additional stats fields into the C# API [#20693]
14. Added pragma bdbsql_single_process to keep the Berkeley DB environment information on the heap instead of in shared memory. This option cannot be used if the database is accessed from multiple processes. [#20789]
15. Improved the ability of DB->compact to move DB_HASH database pages to the beginning of the file. [#20815]

Database Environment Changes

1. Fixed a bug that could cause a segmentation violation when closing an environment handle which has open database handles on partition databases. [#20281]
2. Fixed a bug that could cause a segmentation violation if a region was extended leaving a very small fragment at the end. [#20414]
3. Changed the behavior of the DB_REGISTER | DB_RECOVER flag combination, so that recovery is always run if the environment panic bit is set. [#20767]

Access Method Changes

1. Fixed a bug where database configuration settings could be lost when the database was opened if the open operation was blocked for any amount of time. [#20860]
2. Fixed a bug that bulk update operations did not work correctly on compressed databases. [#19017]
3. Improved the log flushing performance when ftruncate() is not available on a system. [#19725]
4. When performing partial puts in a heap database, empty pieces will no longer be left in a split record chain. [#20052]
5. Fixed a bug where, on systems without FTRUNCATE, db_verify will return an error for truncated heap databases. [#20195]
6. Fixed a bug where BDB could run out of mutexes when many databases are renamed. [#20228]
7. Fixed a bug where the metadata page in hash databases would not be flushed to disk. [#20265]
8. Fixed a bug that could leave deleted pages from a HEAP database in the buffer cache. [#20309]

9. Fixed a bug where the library would fail to put records with overflow keys into hash duplicate database. [#20329]
10. Fixed a bug in DB->compact of btrees that could cause a bad pointer reference. [#20360]
11. Fixed a bug that could cause the last page number stored on the metadata page to be wrong after rolling forward a db->compact operation that freed more pages than will fit in a single log record. [#20646]
12. Fixed a bug that could cause DB->stat to block on a mutex while holding a lock on the metadata page. [#20770]
13. Fixed a bug that could cause DB->compact of a DB_HASH database to fail to mark a page it updated as dirty. [#20778]
14. Fixed a bug where internal HEAP structures were not rebuilt during database handle refresh. [#20821]
15. Fixed a bug with secondary indices, off-page duplicates and DB_READ_COMMITTED which could erroneously release the lock on the page holding a returned record. [#20853]
16. Fixed a bug that could cause a hang or improperly report an empty queue when the queue record numbers wrapped around at 2^{32} . [#20956]
17. Fixed a bug on Linux or Windows that could generate a checksum error if a database file was being opened while the meta data page happened to be flushed from the cache. [#20996]

SQL API Changes

1. Fixed several memory leaks in the Online Backup API. [#19850]
2. Fixed a bug in the SQL API when using large blob items and multiple concurrent connections. [#19945]
3. To avoid a race condition that could cause a snapshot reader to see a wrong version it is now not permitted to open a DB handle specifying DB_MULTIVERSION if that database is currently opened by a handle which did not specify DB_MULTIVERSION. [#19940]
4. Pragma replication=on can now enable replication on an existing database. Turning replication off is now permanent. [#20180]

5. Fixed a bug in the SQL API where it was possible for a schema update to be ignored when accessing a database from multiple processes. [#20319]
6. Fixed a bug where aborting an exclusive transaction followed by an auto-commit read operation causes an assert failure. [#20567]
7. Fixed a bug in the SQL API where using the journal_mode pragma could cause a crash when used as the first operation in a connection on an existing database. [#20620]
8. Turn off the DBSQL encryption option on Windows/WinCE by default to match the behavior on the other platforms. [#20671]
9. Renamed the BDBSQL_OMIT_SHARING preprocessor flag to BDBSQL_SINGLE_PROCESS. [#20789]
10. Fixed a bug dealing with handle lock modes not reflecting the correct state which was causing a deadlock in the SQL API. [#20862]

Java-specific API changes

1. Added ReplicationManagerConnectionStatus class and ReplicationManagerSiteInfo.getConnectionStatus(). Deprecated ReplicationManagerSiteInfo.isConnected(). [#18068]
2. Updated EID_MASTER to be "public static final" so that it would be exposed in Java docs. [#20184]
3. Fixed a bug where calls that return Stat objects could cause a segfault. [#20377]

Replication Changes

1. Fixed quorum computation when most sites are unelectable. [#15251]
2. Made Replication more resilient to random input on its port. [#15712]
3. Fixed a bug where the datadir structure was not maintained during an internal init. [#19041]
4. Fixed a repmgr memory leak when using DB_PRIVATE. [#19363]
5. Fixed a minor bug to handle ENOMEM when using an in-memory temp database. [#20197]
6. Fixed a bug where multiple long running transactions across checkpoints could cause Log Sequence errors on client systems. [#20421]

7. Fixed a bug where multiple Replication Manager processes would sometimes not all conform to replication-group-aware log archiving. [#20342]
8. Fixed a bug where a Replication Manager master could stop functioning after accepting an obsolete group membership site list from another site. [#21804]

Locking Subsystem Changes

1. Fixed a bug that could cause an early lock timeout if a previous error left a lock timeout value set. [#19973]

Logging Subsystem Changes

1. Fixed a bug which could cause an incompletely written log record to be recognized as valid, resulting in recovery failing with the message "Illegal record type <integer, usually 0> in log". [#17851]
2. Fixed a bug where printlog would fail on in-memory heap databases. [#20269]

Memory Pool Subsystem Changes

1. Fixed a bug which overstated the number of clean and dirty pages evicted from the cache. [#20410]
2. Fixed a bug that left a small fragment at the end of a region when extending. [#20414]
3. Fixed a bug where the file bucket was always zero when creating a mpoolfile using the mpool API. [#20468]
4. Fixed a bug with multiversion concurrency control which could cause versions of pages to remain in the cache even though they are no longer needed. [#20570]
5. The memory pool allocator will now start freezing MVCC versions of buffers if it sees more than 1/4 of the available buffers are taken up by versions. [#20836]

Mutex Subsystem Changes

1. Fixed a bug in which DB_ENV->mutex_set_align() could cause DB_ENV->mutex_stat_print(dbenv, DB_STAT_ALL) to display only the first mutex. [#20522]
2. Fixed a bug with DB_ENV->mutex_stat_print() in which the information on some mutexes would not be displayed, if any mutex had been freed and not yet reallocated. [#20533]

3. Fix a race condition between DB_ENV->failchk() and the allocation of a mutex. [#21796]

Transaction Subsystem Changes

1. Fixed a bug where a malloc failure could result in a segfault when doing a put on a database with secondaries. [#20641]

Utility Changes

1. Fixed a bug that would cause verify to call the wrong compare function if there are user defined compare functions used and the database has multilevel off page sorted duplicate trees. [#20284]
2. Fixed a bug that could cause recovery to fail if DB->compact moved the meta data page of a HASH subdatabase. [#20708]
3. Fixed three problems with db_hotbackup's backup of transaction logs. A hot backup did not use any configured log directory, but would try to open the logs in the environment home. The second fix corrected an error path, in which the memory was freed by the wrong function, possibly causing a guard byte error. The third fix fixed the issue that a wrong message would be displayed when only "-l" was specified. [#21313]

Configuration, Documentation, Sample Apps, Portability and Build Changes

1. The DB_CONFIG configuration commands which specify directory pathnames ("set_data_dir", "set_lg_dir", and "set_tmp_dir") now accept names containing whitespace characters. [#20158]

Known Bugs

1. If two SQL processes are concurrently altering the schema of the same tables in a database, there is a race condition that can cause the application to hang. [#20513]
2. Replication groups including machines of different endianness do not support the heap access method. [#21016]
3. If a txn that is attempting to remove a region page from a heap database is aborted and another txn is trying to update that same page then it can cause the original txn to abort. This is timing dependant. [#20939]
4. Utilities can operate on a replication client which is being automatically initialized and may therefore be in an inconsistent state. This can cause the utility to fail or to return invalid results. You can use replication statistics to check the site's role

(st_status) and master generation (st_gen) before and after a utility runs; if neither has changed, the utility's results are valid. [#21593]

Chapter 10. Upgrading Berkeley DB 11.2.5.1 applications to Berkeley DB 11.2.5.2

Table of Contents

[Introduction](#)

[SQLite Interface Upgrade](#)

[32bit/64bit Compatibility on Windows](#)

[Read Only flag for DBT](#)

[New Flag](#)

[Dynamic Environment Configuration](#)

[New Functions](#)

[Deprecated Functions](#)

[Exclusive Transactions in the SQL Layer](#)

[Group Membership in Repmgr](#)

[Upgrading](#)

[New Functions](#)

[Modified Functions](#)

[New Events](#)

[Removed Functions](#)

[New Parameters](#)

[New Structure](#)

[Heap Access Method](#)

[New Functions](#)

[Modified Functions](#)

[New Definition](#)

[Enabling Transaction Snapshots in the SQL Layer](#)

[New Pragmas](#)

[2SITE_STRICT Enabled by Default in Replication](#)

[Enabling Replication in the SQL Layer](#)

[New Pragmas](#)

[Repmgr Message Channels](#)

[New Functions](#)

[Sequence Support in the SQL Layer](#)

[New Functions](#)

[Berkeley DB X/Open Compliant XA Resource Manager](#)

[Constraints](#)

[New Flag](#)

[Modified Function](#)

[Hot Backup Changes](#)

[Berkeley DB Library Version 11.2.5.2 Change Log](#)

[Database or Log File On-Disk Format Changes](#)

[New Features](#)

[Database Environment Changes](#)

[Concurrent Data Store Changes](#)

[Access Method Changes](#)

[SQL API Changes](#)

[C API Changes](#)

[Tcl-specific API Changes](#)

[C#-specific API Changes](#)

[Replication Changes](#)

[Locking Subsystem Changes](#)

[Logging Subsystem Changes](#)

[Memory Pool Subsystem Changes](#)

[Mutex Subsystem Changes](#)

[Transaction Subsystem Changes](#)

[Test Suite Changes](#)

[Utility Changes](#)

[Configuration, Documentation, Sample Apps, Portability and Build Changes](#)

[Example Changes](#)

[Miscellaneous Bug Fixes](#)

[Deprecated Features](#)

[Known Bugs](#)

Introduction

The following pages describe how to upgrade applications coded against the Berkeley DB 11.2.5.1 library to the Berkeley DB 11.2.5.2 library (both of which belong to Berkeley DB 11g Release 2). This information does not describe how to upgrade Berkeley DB 1.85 release applications.

For information on the general process of upgrading Berkeley DB installations and upgrade instructions related to historical releases, see the [Berkeley DB Upgrade Guide](#).

SQLite Interface Upgrade

Berkeley DB's SQL interface includes code from SQLite. The version of SQLite used has been upgraded, so DB SQL is compatible with SQLite version 3.7.6.2. Please see the release notes at <http://sqlite.org/changes.html> for further information.

32bit/64bit Compatibility on Windows

Berkeley DB can now be compiled on Windows so that 32 bit and 64 bit applications can concurrently access a BDB environment. To enable this feature, build both the 32 bit BDB library and application and the 64 bit library and application with the flag `/D HAVE_MIXED_SIZE_ADDRESSING`. Note that private environments are disabled under the compatibility mode.

Read Only flag for DBT

New Flag

A [DBT](#) can now be set as read-only, when passed to the [DB->get\(\)](#) method, using the flag `DB_DBT_READONLY`. This is useful when using a static string as a key value, because this flag will prevent Berkeley DB from updating the [DBT](#).

New Flag

- `DB_DBT_READONLY`

Dynamic Environment Configuration

New Functions

Deprecated Functions

Memory is now allocated incrementally as needed, instead of all at once during environment initialization, for structures that support locks, transactions, threads, and mutexes. With this change new functions have been added that configure how much memory is allocated initially, and how much that memory is allowed to grow. The old memory configuration functions have been deprecated.

New Functions

- [DB_ENV->set_memory_init\(\)](#)
- [DB_ENV->get_memory_init\(\)](#)
- [DB_ENV->set_memory_max\(\)](#)
- [DB_ENV->get_memory_max\(\)](#)
- [DB_ENV->set_lk_tablesize\(\)](#)
- [DB_ENV->get_lk_tablesize\(\)](#)
- [DB_ENV->mutex_set_init\(\)](#)
- [DB_ENV->mutex_get_init\(\)](#)

Deprecated Functions

- [DB_ENV->mutex_set_max\(\)](#)
- [DB_ENV->mutex_get_max\(\)](#)
- [DB_ENV->set_lk_max_lockers\(\)](#)
- [DB_ENV->get_lk_max_lockers\(\)](#)
- [DB_ENV->set_lk_max_locks\(\)](#)
- [DB_ENV->get_lk_max_locks\(\)](#)
- [DB_ENV->set_lk_max_objects\(\)](#)
- [DB_ENV->get_lk_max_objects\(\)](#)
- [DB_ENV->set_thread_count\(\)](#)
- [DB_ENV->get_thread_count\(\)](#)
- [DB_ENV->set_tx_max\(\)](#)
- [DB_ENV->get_tx_max\(\)](#)

Exclusive Transactions in the SQL Layer

Issuing the SQL command `BEGIN TRANSACTION EXCLUSIVE` will now cause any other transactions accessing the database to block, or return a `SQLITE_BUSY` or `SQLITE_LOCK` error, until the exclusive transaction is committed or aborted. Previously, non-exclusive transactions could execute concurrently with an exclusive transaction.

Group Membership in Repmgr

[Upgrading](#)

[New Functions](#)

[Modified Functions](#)

[New Events](#)

[Removed Functions](#)

[New Parameters](#)

[New Structure](#)

Replication Manager now manages group membership much more closely, making it much easier for applications to add and remove sites from a replication group without risk of transaction loss. In order to accomplish this, the API for configuring group membership has changed significantly. The `repmgr_set_local_site()` and `repmgr_add_remote_site()` methods no longer exist; they are replaced by a new handle type, `DB_SITE`. The `repmgr_get_local_site()`

method has been replaced by `DB_ENV->repmgr_site()`, which now returns a `DB_SITE` handle instead of a raw host/port network address.

Replication Manager applications may no longer call the `DB_ENV->rep_set_nsites()` method, because the Replication Manager now tracks the number of sites in the replication group for you. Replication Manager applications may still call `DB_ENV->rep_get_nsites()`, but only after a successful call to `DB_ENV->repmgr_start()`.

For applications using the replication Base API there is no change, except that they may now call `DB_ENV->rep_set_nsites()` to change the group size even when Master Leases are in use.

The new Replication Manager group membership functionality is described in the [Managing Replication Manager Group Membership](#) chapter in the Berkeley DB Programmer's Reference Guide.

Replication Manager no longer prints an error message on a connection failure. Instead it generates an event with the equivalent information (invoking the application's event-handling call-back function).

Upgrading

An existing application running a previous version of BDB can do a "live upgrade" so that only one site at a time has to be shut down. To do this, restart each site in the group, with the old master being shutdown last. When each site is restarted, use `DB_SITE` to configure the local site with the flag `DB_LEGACY`, and create a `DB_SITE` handle with a full specification of all the remote site addresses for all other sites currently in the group, and configure each handle with the `DB_LEGACY` flag. When the old master is restarted and a new master has been established, the new master is ready to manage membership changes, and new sites can be added as usual. But the application must not try to add new sites, or remove existing sites, during the mixed-version transitional phase.

To do a non-live upgrade shutdown the entire replication group. Then restart the group with each site configured with the `DB_LEGACY` flag, and in `DB_REP_ELECTION` mode.

New Functions

- `DB_ENV->repmgr_site()`
- `DB_ENV->repmgr_site_by_eid()`
- `DB_SITE->set_config()`
- `DB_SITE->get_config()`
- `DB_SITE->remove()`
- `DB_SITE->get_eid()`
- `DB_SITE->get_address()`

- [DB_SITE->close\(\)](#)

Modified Functions

- [DB_ENV->rep_set_nsites\(\)](#) is no longer used by the Replication Manager, but is still used by the Base API. It can now be used to change the number of sites dynamically, even when master leases are in use.

New Events

- DB_EVENT_REP_SITE_ADDED
- DB_EVENT_REP_SITE_REMOVED
- DB_EVENT_REP_LOCAL_SITE_REMOVED
- DB_EVENT_REP_CONNECT_BROKEN
- DB_EVENT_REP_CONNECT_ESTD
- DB_EVENT_REP_CONNECT_TRY_FAILED
- DB_EVENT_REP_INIT_DONE

Removed Functions

- DB_ENV->repmgr_set_local_site()
- DB_ENV->repmgr_add_local_site()
- DB_ENV->repmgr_add_remote_site()
- DB_ENV->repmgr_get_local_site()

New Parameters

The following new parameters are passed to [DB_SITE->set_config\(\)](#).

- DB_BOOTSTRAP_HELPER
- DB_GROUP_CREATOR
- DB_LEGACY
- DB_LOCAL_SITE
- DB_REPMGR_PEER

New Structure

- DB_REPMGR_CONN_ERR encapsulates an EID and an integer system error code.

Heap Access Method

New Functions

Modified Functions

New Definition

Databases can now be configured as heaps by passing the access type `DB_HEAP` to `DB->open()`. Heap size can be configured with `DB->set_heapsize()`, and `DB->stat()` now returns heap statistics in the structure `DB_HEAP_STAT` when applied to a heap database.

New Functions

- `DB->set_heapsize()`
- `DB->get_heapsize()`

Modified Functions

- `DB->open()` now accepts `DB_HEAP` as an access type.
- `DB->stat()` now returns heap statistics in the structure `DB_HEAP_STAT`.

New Definition

- `DB_HEAP_RID` is the defined heap key value.

Enabling Transaction Snapshots in the SQL Layer

New Pragmas

Read/write concurrency can now be enabled in the SQL API by using `PRAGMA multiversion=on` before accessing any tables in the database. After multiversion has been enabled, it can be temporarily disabled using the `PRAGMA transaction_snapshots=on/off`.

New Pragmas

For more details on pragmas concerning Transaction Snapshots read [Using Multiversion Concurrency Control](#) in the Berkeley DB Getting Started with the SQL APIs guide.

- `PRAGMA multiversion=ON|OFF;`
- `PRAGMA snapshot_isolation=ON|OFF`

2SITE_STRICT Enabled by Default in Replication

The `2SITE_STRICT` replication configuration parameter is now turned on by default. This configuration parameter is controlled using the `DB_REPMGR_CONF_2SITE_STRICT` flag on the `DB_ENV->rep_set_config()` method.

Enabling Replication in the SQL Layer

New Pragmas

Replication can now be enabled and configured in the SQL layer using pragmas. The pragmas `replication_local_site`, `replication_initial_master`, and `replication_remote_site` can be used to configure the replication group. Note that when the BDB SQL replicated application is initially started, a specific master site must be explicitly designated. After configuring the replication group, start replication using `PRAGMA replication=ON`.

To display replication statistics in the `dbsql` shell, use:

```
dbsql> .stat :rep
```

New Pragmas

For more details on the replication pragmas see [Replication PRAGMAs](#) in the Berkeley DB Getting Started with the SQL APIs guide.

- `PRAGMA replication=ON|OFF`
- `PRAGMA replication_initial_master=ON|OFF`
- `PRAGMA replication_local_site="hostname:port"`
- `PRAGMA replication_remote_site="hostname:port"`
- `PRAGMA replication_remove_site="host:port"`
- `PRAGMA replication_verbose_output=ON|OFF`
- `PRAGMA replication_verbose_file=filename`

Repmgr Message Channels

New Functions

Application components running at various sites within a replication group can now use the Replication Manager's existing TCP/IP communications infrastructure to send and process messages among themselves, using the `DB_CHANNEL` handle. `DB_ENV->repmgr_channel()` is used to create the `DB_CHANNEL` handle. `DB_CHANNEL->send_msg()` and `DB_CHANNEL->send_request()` are used to send synchronous and asynchronous messages that are handled by the function set by `DB_ENV->repmgr_msg_dispatch()`. `DB_CHANNEL->set_timeout()` is used to configure channel time out, and `DB_CHANNEL->close()` closes the channel and frees resources held by it.

New Functions

- `DB_ENV->repmgr_msg_dispatch()`

- `DB_ENV->repmgr_channel()`
- `DB_CHANNEL->send_msg()`
- `DB_CHANNEL->send_request()`
- `DB_CHANNEL->set_timeout()`
- `DB_CHANNEL->close()`

Sequence Support in the SQL Layer

New Functions

A partial implementation of the sequence API defined in the SQL 2003 specification has been added to the SQL layer. A sequence is created using the syntax:

```
SELECT create_sequence("sequence_name"...)

```

The sequence numbers are accessed using

```
SELECT nextval("sequence_name")

```

and

```
SELECT currval("sequence_name")

```

Finally, a sequence can be dropped using

```
SELECT drop_sequence("sequence_name")

```

New Functions

The four new functions, which have to be called as part of a `SELECT` statement, are describe in more detail in [Using Sequences](#) in the Berkeley DB Getting Started with the SQL APIs guide.

- `create_sequence`
- `seq_nextval`
- `seq_currval`
- `seq_drop_sequence`.

Berkeley DB X/Open Compliant XA Resource Manager

Constraints

New Flag

Modified Function

The Berkeley DB X/open compliant XA resource manager has been restored. (It was removed from the product after the 4.7 release.) The new implementation includes support for multi-threaded servers. Consult the documentation of your chosen transaction manager to learn how to implement a multi-threaded server.

Constraints

Applications that use a BDB XA resource manager must now take into account the following constraints.

- No in-memory logging.
- No application-level child transactions.
- All database-level operations (open, close, create and the like) must be performed outside of a global transactions (i.e., they can be performed in local BDB transactions, but not while a distributed XA transaction is active).
- Environment configuration must be done using a DB_CONFIG file.
- Cursors must be closed before a service invocation returns.

New Flag

- DB_XA_CREATE - This flag is passed to [db_create\(\)](#) to create a DB handle that supports XA transactions.

Modified Function

- [DB->stat\(\)](#) now returns the field DB_TXN_STAT->DB_TXN_ACTIVE->xa_status, which contains information on the XA transactions.

Hot Backup Changes

Because non-UNIX systems do not support atomic file system reads, the [db_hotbackup](#) utility has been modified to read data through the environment. If your application is running on a UNIX based system such as Solaris, HP-UX, BSD or Mac OS, you can specify the -F flag to read directly from the filesystem. Please refer to [Recovery procedures](#) in the Berkeley DB Programmer's Reference Guide for more information on safely backing up your databases.

Berkeley DB Library Version 11.2.5.2 Change Log

[Database or Log File On-Disk Format Changes](#)

[New Features](#)

[Database Environment Changes](#)

[Concurrent Data Store Changes](#)

[Access Method Changes](#)

[SQL API Changes](#)

[C API Changes](#)

[Tcl-specific API Changes](#)

[C#-specific API Changes](#)

[Replication Changes](#)

[Locking Subsystem Changes](#)

[Logging Subsystem Changes](#)

[Memory Pool Subsystem Changes](#)

[Mutex Subsystem Changes](#)

[Transaction Subsystem Changes](#)

[Test Suite Changes](#)

[Utility Changes](#)

[Configuration, Documentation, Sample Apps, Portability and Build Changes](#)

[Example Changes](#)

[Miscellaneous Bug Fixes](#)

[Deprecated Features](#)

[Known Bugs](#)

This is the changelog for Berkeley DB 11g Release 2 (library version 11.2.5.2).

Database or Log File On-Disk Format Changes

1. Existing database file formats were unchanged in library version 11.2.5.2. However, a new database file format, "heap", was introduced.
2. The log file format changed in library version 11.2.5.2.

New Features

1. Replication Manager now manages Group Membership. This allows sites to be added to and removed from the replication group dynamically. Replication Manager also now automatically keeps track of the group size (nsites). [#14778]
2. Initial allocations for various non-pagebuffer (mpool) system resources may now be specified, as well as a total maximum of memory to use, rather than specifying a maximum value for each resource. [#16334]
3. Implemented Berkeley DB globalization support architecture to enable localized and stripped error and output messages. [#16863]
4. Added a new access method, DB_HEAP. Heap aims for efficient use (and re-use) of disk space. Keys in a heap database are automatically generated by BDB, it is recommended that one or more secondary indexes be used with a heap database. For full details on DB_HEAP, see the Programmer's Reference Guide. [#17627]
5. Added a compatible mode for 32bit and 64bit Windows environment. [#18225]
6. For the SQL API, concurrency between read and write transactions can now be enabled using "PRAGMA multiversion". Added several pragmas that can be used to configure the Berkeley DB datastore. [#18521]
7. Add several new pragmas to provide in-process support for replication in the SQL API. [#18528]
8. The Berkeley DB X/open compliant XA resource manager has been restored, including support for multi-threaded servers. [#18701]
9. Improved the ability to recover from an application crash on connections through the SQL API. Berkeley DB will try to automatically clean up locks, mutexes and transactions from the failed process. [#18713]
10. Add support for sequence usage in the SQL API using SQLite custom functions. [#19007]
11. Add a pragma in the SQL API to allow execution of a cache trickle command. [#19202]

12. Add a pragma in the SQL API to allow configuration of DB_SYSTEM_MEM environments. [#19249]
13. The new db_env_set_win_security(SEcurity_ATTRIBUTES *) function allows an application to specify the particular Microsoft Windows security attributes to be used by Berkeley DB. This helps support applications which reduce their privileges after opening the environment. [#19529]

Database Environment Changes

1. None

Concurrent Data Store Changes

1. None

Access Method Changes

1. Modified the queue access method so that it only uses latches on the metadata page rather than a latch and a lock. This was done to improve performance. [#18749]
2. Fixed several bugs that could cause an update running with MVCC to get the wrong version of a page or improperly update the metadata last page number. [#19063]
3. The database open code will no longer log the open and close of the master database in a file when opening a sub database in that file. [#19071]
4. Fixed a bug where an error during an update to a hash database with DB_NOOVERWRITE set could return DB_KEYEXIST rather than the correct error. [#19077]
5. Fixed a bug that could cause the wrong page number to be on a root or metadata page if DB->compact moved the page and the operation was later rolled forward. [#19167]
6. Fixed a bug that could cause the close of a secondary index database to fail if the transaction doing the open aborted. [#19169]
7. Fixed a bug that could prevent an update to a primary recno or queue database with DB_NOOVERWRITE set. [#19230]
8. Fixed a bug when an update to a database with DB_NOOVERWRITE set could incorrectly return DB_KEYEXIST rather than the correct error (e.g., DB_LOCK_DEADLOCK). [#19345]

9. Fixed a bug preventing the use of the set_re_len and set_re_pad methods with a RECNO database when configuring with --disable-queue. [#19367]
10. Fixed a bug in DB->compact on BTREE databases that did not check if the last page in the database could be moved to a lower numbered page. [#19394]
11. Fixed a bug that could cause a Log Sequence Error when recovering the deallocation of a multiple page overflow chain. [#19474]
12. Fixed a bug that could cause a diagnostic assertion if MVCC was in use and multiple levels of a btree needed to be split. [#19481]
13. Fixed a few error paths that could cause a Panic with an "unpinned page returned" error. [#19493]
14. Fixed a bug that closed a race condition that under heavy mult-threaded appending to a queue database could cause some records to be lost. [#19498]
15. Fixed a bug that might cause DB->compact to mis-estimate the size of an overflow record when merging two pages. This may cause the page to have more data than desired. [#19562]
16. Fixed a bug in DB_ENV->fileid_reset that did not update the fileid's on the metadata pages of subdatabases if the database file was not in native byte order. [#19608]
17. Fixed a bug that caused the first directory specified in the create of a partitioned database to get too many partitions. [#20041]

SQL API Changes

1. Fixed a race condition that would cause a corruption error in one process when two processes created the same SQL database. [#18929]
2. Fixed a bug that would cause a constraint violation when updating the primary key with the same value. [#18976]
3. Overwriting an old backup with a new backup using the SQL online backup API will no longer double the size of the database. [#19021]
4. Implemented index optimizations for indexes on large values. [#19094]
5. Fixed a bug that could cause an undetected deadlock between a thread which moved a metadata or root page via a DB->compact operation and another thread trying to open the database if the old page was being removed from the file. [#19186]

6. Fix a bug in the BDBSQL_FILE_PER_TABLE option, to allow absolute path names. [#19190]
7. Add a pragma to allow configuration of DB_SYSTEM_MEM environments. [#19249]
8. Exclusive transactions will now block new transactions and will prevent existing transactions from making forward progress. [#19256]
9. Fixed a bug that would cause assert error when opening an in-memory hash database with thread count configured when compiled with --enable-diagnostic. [#19357]
10. Upgrade the bundled version of SQLite to 3.7.6.2 [#19376]
11. Fixed a performance bug with the cache victim selection algorithm when there were multiple cache regions. [#19385]
12. Fixed a bug which could cause two SQL threads to have an undetected deadlock when opening or closing tables. [#19386]
13. Fix a bug that could cause a hang when deleting a table if there are multiple connections to a database from different processes. [#19419]
14. Fixed a bug which could cause multiple threads performing DB->compact on the same database file to overrun the in-memory freelist, which could potentially lead to memory corruption. [#19571]
15. Fixed a bug in DB->compact that could cause a loop if an attempt to move a sub-database meta data page deadlocked. [#20028]

C API Changes

1. Fixed a bug where encryption could not be enabled for individual databases in an encrypted environment. [#18891]
2. Removed two unused error codes, DB_NOSERVER_HOME and DB_NOSERVER_ID. [#18978]
3. Added a DB_DBT_READONLY flag so that users can pass in a non-usermem key (DB_DBT_USERMEM) for get operations. [#19360]
4. Fixed a bug in DB/DBC->get/pget that the partial flags are silently ignored with positional flags and return inconsistent DBT. [#19540]
5. Fixed a bug which prevented items from being deleted on a secondary database. [#19573]

6. Fixed a bug to correctly handle the DB_BUFFER_SMALL case on delete operations when compression is enabled. [#19660]

Tcl-specific API Changes

1. None.

C#-specific API Changes

1. Added support for partial put/get in the C# API. [#18795]
2. Fixed a bug in compare delegate for secondary db. [#18935]

Replication Changes

1. Replication Manager now allows differing ack policies at different sites throughout the group, and supports dynamic changes to the ack policy. (The ack policy in force is determined by the current master.) [#14993]
2. Replication Manager "channels" feature allows applications to share repmgr's communication facilities. [#17228]
3. Add example program for RepMgr "channels" feature: ex_rep_chan. [#17387]
4. Replication Manager now allows dynamic changes to a site's "electability" (changes between zero and non-zero priority). This feature should be used with care, because electability changes can in boundary cases invalidate durability guarantees granted for previous transactions. [#17497]
5. Changed election criteria so that later group transactions won't get overwritten by earlier generations with more log. [#17815]
6. Added changes to master lease checks that result in improved performance when using master leases. [#18960]
7. A log write failure on a replication master will now cause a panic since the transaction may be committed on some clients. [#19054]
8. Fixed a few memory leak conditions on error paths. [#19131]
9. Change lease code so that zero priority sites do not count in lease guarantees since they cannot be elected. [#19154]
10. Repmgr rerequest processing is moved from a dedicated thread to heartbeat messages. Repmgr clients using heartbeats can now detect and rerequest missing final master log records without master activity. [#19197]

11. Repmgr statistics are now included in full statistics output for an environment. [#19198]
12. Fix an inefficiency in mixed version elections. We now check if an election is won via the EID instead of priority. [#19254]
13. Changed election LSNs to use the last txn commit LSN instead of the end of the log. [#19278]
14. Create replication internal database files in the environment home directory rather than the data directory so that they are in the same location as the other internal replication files. [#19403]
15. Fix a bug that was preventing repmgr from calling an election when starting a site with the DB_REP_ELECTION flag. [#19546]
16. Fixed a bug which could cause a segfault at a replication master if a named in-memory database was being created around the same time as a client site were synchronizing (in "internal init") with the master. [#19583]
17. Adjust lease code to consider timeout length when retrying. [#19705]
18. Fixed a bug that could cause a crash in replication groups of more than 10 sites, with multiple processes sharing each DB environment concurrently. [#19818]
19. Fix a bug where an assertion failure could happen if pages in a database were deallocated during a client internal initialization. [#19851]
20. Fix a bug where an internal initialization of a queue database with non-contiguous extent files could return an error. [#19925]
21. The 2SITE_STRICT replication configuration parameter is now turned on by default. It can be turned off via a call to DB_ENV->rep_set_config(). [#19937]
22. Repmgr heartbeats can now help detect a duplicate master without the need for application activity. [#19950]

Locking Subsystem Changes

1. Fixed a bug where an updater supporting DB_READ_UNCOMMITTED might downgrade its lock too soon if there was an error during the update. [#19155]
2. Fixed a bug where transaction timeouts could have been specified in a database environment where the locking subsystem was disabled. [#19582]

3. Fixed a bug in a diagnostic assertion that was improperly triggered by the removal of a sub-database. [#19683]
4. Fixed a bug that would cause DB_ENV->failcheck to free locks for a locker associated with a database handle after the thread that opened the handle exited. [#19881]

Logging Subsystem Changes

1. Enhanced recovery so that it will not output extra checkpoint or transaction id recycle log records if there was no activity since the last checkpoint. [#15330]
2. Log checksums can now be disabled using the compile argument --disable-log-checksum. This will give a performance increase at the risk of undetectable corruption in the log records, which would make recovery impossible. [#19143]
3. Fixed a bug that could cause a page that should have been removed from the end of a file still be in the copy of the file in a hot backup. [#19996]

Memory Pool Subsystem Changes

1. Fixed a bug in MPOOLFILE->get that did not permit the DB_MPOOL_DIRTY flag to be used with other flags. [#19421]

Mutex Subsystem Changes

1. Fixed a bug when the mutex region needs to be larger than 4GB, the region size was incorrectly adjusted to be slightly too small to fit the mutexes. [#18968]
2. Fixed a performance problem with hybrid shared latches in which a request for exclusive access would busy-wait (rather than put itself to sleep) if the latch were held by a shared reader. This also fixed the timeout handling of hybrid mutexes. In some cases the timeout would not be honored, resulting in delays for the replication "read your writes" feature which were longer than requested. [#18982]
3. Fixed the timeout handling of the pthreads mutexes used by the replication "read your writes" feature. When a timeout occurred there was a race condition which might result in a hang. [#19047]

Transaction Subsystem Changes

1. Fixed a leak of log file ids when a database is closed before the end of a transaction that references it. [#15957]

2. Fixed a bug that would cause a panic if a child transaction performed a database rename, then aborted, and then the parent transaction committed. [#18069]
3. Fixed a bug where we released the metadata page lock too early if a non-transactional update was being done. [#19036]
4. Removed the possibility that checkpoints will overlap in the log, decreasing the time to recover. [#19062]

Test Suite Changes

1. Require Tcl 8.5 or greater.

Utility Changes

1. Added a new utility, db_tuner, which analyzes the data in a btree database, and suggests a reasonable pagesize. [#18910]
2. Fixed some bugs in log_verify when there are in-memory database logs and subdb logs. [#19157]
3. Modified db_hotbackup to not read from the file system as required on non-UNIX systems. Also provided the db_copy function for this purpose. [#19863]
4. Fixed db_hotbackup so that when -d/-l or -D is not specified, DB_CONFIG is used to determine the locations of the databases and logs in the source environment. [#19994]

Configuration, Documentation, Sample Apps, Portability and Build Changes

1. Changed SQL API library built on *nix to link with libpthreads when necessary. [#19098]
2. Added CPPFLAGS into our --enable-jdbc configuration. [#19234]
3. Added encryption support into the Windows CE build project for SQL API. [#19632]
4. Fixed a bug in the STAT_INC_VERB() dtrace probe that was causing compiler warnings. [#19707]
5. Fixed a bug that could cause a trap in db_dump using salvage mode if a page was found that was not associated with any database in the file. [#19974]
6. On Cygwin, circumvented a bug in libtool that is exposed when building the BDB SQL API in a directory path containing whitespace characters. [#19812]

Example Changes

1. Update repmgr C, C#, C++, Java examples(ex_rep_mgr, ex_rep_gsg_repmgr, ex_rep_chan, excs_repquote, excxx_repquote, excxx_epquote_gsg, repquote, repquote_gsg) with their related API changes for group membership. [#19586] [#19622]
2. Port ex_rep_chan, ex_rep_gsg_repmgr,ex_rep_gsg_simple, excxx_repquote_gsg_repmgr, excxx_repquote_gsg_simple to Window.[#19890]

Miscellaneous Bug Fixes

1. Fixed a bug where memory copied from the Java API could leak if flags were not correctly configured. [#19152]

Deprecated Features

1. None

Known Bugs

1. The SQL API has a known issue when using a blob field with a lot of content and multiple concurrent connections to the database. [#19945]
2. Rollback of a dropped table in the SQL layer contains a mutex leak, which can consume all mutex resources if enough rollbacks of table drops are performed. [#20077]
3. The DB_CONFIG configuration parameters which specify path names currently do not support names containing any whitespace characters. [#20158]
4. The BFile module has a known crash issue when using BFile handle for SQL expressions interface on 64bit platforms. [#20193]
5. On systems without FTRUNCATE, db_verify will return an error for truncated heap databases. This is a bug in db_verify, the database has been truncated correctly and can be used in the future. [#20195]
6. An application using queue extents which is append mostly could see a decrease in the buffer pool hit rate due to the failure to remove pages from closed extents from the buffer pool. [#20217]

Chapter 11. Upgrading Berkeley DB 11.2.5.0 applications to Berkeley DB 11.2.5.1

Table of Contents

[Introduction](#)

[DPL Applications must be recompiled](#)

[Source Tree Rearranged](#)

[SQLite Interface Upgrade](#)

[Mod_db4 Support Discontinued](#)

[Berkeley DB Library Version 11.2.5.1 Change Log](#)

[Database or Log File On-Disk Format Changes](#)

[New Features](#)

[Database Environment Changes](#)

[Concurrent Data Store Changes](#)

[Access Method Changes](#)

[API Changes](#)

[SQL-Specific API Changes](#)

[Tcl-Specific API Changes](#)

[Java-Specific API Changes](#)

[C#-Specific API Changes](#)

[Direct Persistence Layer \(DPL\), Bindings and Collections API](#)

[Replication Changes](#)

[Locking Subsystem Changes](#)

[Logging Subsystem Changes](#)

[Memory Pool Subsystem Changes](#)

[Mutex Subsystem Changes](#)

[Transaction Subsystem Changes](#)[Test Suite Changes](#)[Utility Changes](#)[Configuration, Documentation, Sample Apps, Portability, and Build Changes](#)[Example Changes](#)[Miscellaneous Bug Fixes](#)[Deprecated Features](#)[Known Bugs](#)

Introduction

The following pages describe how to upgrade applications coded against the Berkeley DB 11.2.5.0 library to the Berkeley DB 11.2.5.1 library (both of which belong to Berkeley DB 11g Release 2). This information does not describe how to upgrade Berkeley DB 1.85 release applications.

For information on the general process of upgrading Berkeley DB installations and upgrade instructions related to historical releases, see the [Berkeley DB Upgrade Guide](#).

DPL Applications must be recompiled

Applications that use the Java interface's Direct Persistence Layer must be recompiled, due to a change in the return type of the setter methods in `StoreConfig` and `EvolveConfig` classes. The setter methods now return `this` instead of `void`.

Source Tree Rearranged

The source code hierarchy has been reorganized. Source files that belong to Berkeley DB core are now in a top-level directory named `src`. Files related to language interfaces are in `lang`, and all examples are collected under `examples`.

SQLite Interface Upgrade

Berkeley DB's SQL interface includes code from SQLite. The version of SQLite used has been upgraded, so DB SQL is compatible with SQLite version 3.7.0.1. Please see the release notes at <http://sqlite.org/changes.html> for further information.

Mod_db4 Support Discontinued

The mod_db4 apache module is no longer included in the release.

Berkeley DB Library Version 11.2.5.1 Change Log

[Database or Log File On-Disk Format Changes](#)

[New Features](#)

[Database Environment Changes](#)

[Concurrent Data Store Changes](#)

[Access Method Changes](#)

[API Changes](#)

[SQL-Specific API Changes](#)

[Tcl-Specific API Changes](#)

[Java-Specific API Changes](#)

[C#-Specific API Changes](#)

[Direct Persistence Layer \(DPL\), Bindings and Collections API](#)

[Replication Changes](#)

[Locking Subsystem Changes](#)

[Logging Subsystem Changes](#)

[Memory Pool Subsystem Changes](#)

[Mutex Subsystem Changes](#)

[Transaction Subsystem Changes](#)

[Test Suite Changes](#)

[Utility Changes](#)

[Configuration, Documentation, Sample Apps, Portability, and Build Changes](#)

[Example Changes](#)

Miscellaneous Bug Fixes

Deprecated Features

Known Bugs

This is the changelog for Berkeley DB 11g Release 2 (library version 11.2.5.1).

Database or Log File On-Disk Format Changes

1. The database file format was unchanged in 11gR2 library version 11.2.5.1.
2. The log file format was unchanged in 11gR2 library version 11.2.5.1.

New Features

1. Added Performance event monitoring support for DTrace and SystemTap which can be enabled during configuration. Static probes have been defined where statistics values are updated, where mutex or transactional consistency lock waits occur, and where some other potentially lengthy operations may be initiated. [#15605]
2. Added a new acknowledge policy - DB_REPMGR_ACKS_ALL_AVAILABLE. [#16762]
3. Added transactional bulk loading optimization for non-nested transactions. [#17669]
4. Added exclusive transaction support for the SQL API. [#17822]
5. Added support for bulk update and delete in C# API. [#18011]
6. Added a db_replicate utility. [#18326]
7. Added an implementation of the Online Backup API. [#18500]
8. Added support in Berkeley DB SQL for the vacuum and incremental vacuum pragmas [#18545]
9. Added an option to automatically convert SQLite databases to Berkeley DB on opening. [#18531]
10. Added BDBSQL_SHARE_PRIVATE, an option to enable inter-process sharing of DB_PRIVATE environments using multiple-reader. [#18533]
11. Added database-level locking to optimize single-threaded operations and remove locking limitations for database load operations. [#18549]

12. Added support for DB_INIT_REP, DB_PRIVATE and DB_THREAD in DB_CONFIG file. [#18555]
13. Added support for the BDBSQL_DEFAULT_PAGE_SIZE pragma to override Berkeley DB's choice of page size depending on the filesystem. Use SQLITE_DEFAULT_PAGE_SIZE rather than a hard-coded default. [#18577]
14. Added an extension that allows access to binary files stored outside of the database. What is stored in the database is a pointer to the binary file. [#18635]
15. Added .stat command to dbsql shell to print environment, table, and index statistics. [#18640]
16. Added enhancements to reduce the size of indexes in the SQL API by allowing duplicates in the index database and moving the rowid from the index key into the index data. [#18653]
17. Added a compile time flag BDBSQL_FILE_PER_TABLE that causes each table to be created in a separate file. This flag replaces the BDBSQL_SPLIT_META_TABLE flag. [#18664]
18. Added the handling of read only and read write open of the same database in BDB SQL [#18672]
19. Added an encryption implementation to the SQL API [#18683]

Database Environment Changes

1. Fixed failchk behavior on QNX. [#17403]
2. Fixed a bug that prevented the same process from connecting to the database after recovery is performed. [#18535]
3. Fixed a bug which would occur when recovery checkpoint was not written because the cache ran out of space attempting to flush the memory pool cache. The environment would be recovered and all database where made available, but some databases would incorrectly closed. This would cause a subsequent recovery to fail on its backward pass with the error "PANIC: No such file or directory". [#18590]
4. Fixed a bug that could cause recovery to fail with the error "DB_LOGC->get: log record LSN %u/%u: checksum mismatch" if the last log file was nearly full and ended with a partially written log record which was smaller than a checkpoint record. It now erases the invalid partial record before switching to the new log file. [#18651]

Concurrent Data Store Changes

1. None

Access Method Changes

1. Fixed a bug such that segmentation fault does not occur if DB->set_partition_dirs is called before DB->set_partition. [#18591]
2. Fixed a bug such that the error "unknown path" does not occur if you put duplicate records into a duplicated sorted HASH database with DB_OVERWRITE_DUP option. [#18607]
3. Added the ability to specify that data should not be logged when removing pages from a database. This can be used if the ability to recover the data is not required after the database has been removed. [#18666]
4. Fixed a bug that caused an aborting transaction to fail if it aborted while a DB->compact of the same HASH database was compacting the dynamic hash table [#18695]
5. Fixed a bug that could cause DB->compact to loop on a DB_RECNO database or a database with an multilevel unsorted off page duplicate tree. [#18722]
6. Fixed a bug that could cause an illegal page type error when using a HASH database with MVCC and the HASH table was contracted and then extended. [#18785]
7. Fixed locking bugs: [#18789] The Db->compact method of BTREE with MVCC would return an unpinned page. The RECNO option would fail to lock the next page when splitting a leaf page.
8. Fixed a bug that could cause data to not be returned in a HASH database that was one of multiple databases in a file and was opened prior to running DB->compact method on that database in another thread of control [#18824]
9. Fixed a bug where doing a bulk insert with secondaries could return an error incorrectly. [#18878]
10. Fixed a bug that would return DB_NOTFOUND instead of DB_BUFFER_SMALL when the first item in a HASH database is larger than the user supplied buffer. [#18829]

API Changes

1. Fixed various items uncovered by extending DB_CONFIG support: [#18720] - Added missing set_cache_max method, and fixed name of log_set_config (was

set_log_config). - Added new DB_ENV->repmgr_get_local_site method. - Fixed a bug which could fail to allocate enough mutexes when specifying a maximum cache size. - Fixed a bug that could allocate multiple caches when a small cache size was specified.

SQL-Specific API Changes

1. Allowed SQL applications to attach to the same database multiple times unless shared cache mode is explicitly requested. [#18340]
2. Fixed a bug where auto-removal of log files after writing a checkpoint was not functioning correctly. [#18413]
3. Fixed a race between opening and closing SQL databases from multiple threads that could lead to the error "DB_REGISTER limits processes to one open DB_ENV handle per environment". [#18538]
4. Optimized the SQL adapter for joins. Reduce the number of Berkeley DB operations in a join by caching the maximum key in the primary. [#18566]
5. A SQLITE_LOCKED or SQLITE_BUSY error returned by a statement in an explicit transaction will no longer invalidate the entire transaction, but just the statement that returned the error. [#18582]
6. Changed how multiple connections to the same database are detected. Used a fileid so that different paths can be used without error. [#18646]
7. Fixed a bug where the journal (environment) directory was being created prior to the actual environment. [#18656]
8. Added a new PRAGMA to allow tuning of when checkpoints are run. [#18657]
9. Fixed spurious "column <x> not unique" error messages.[#18667]
10. Fixed a segmentation fault that could happen when memory could not be allocated for the index key in the SQL API.[#18783]
11. Fixed a bug causing a segfault when releasing a savepoint that was already released. [#18784]

Tcl-Specific API Changes

1. Changed to link tcl8.5 by default on Windows[#18244]

Java-Specific API Changes

1. Fixed a bug where `getAllowPopulate` and `getImmutableSecondaryKey` method always returned false for `SecondaryConfig` objects returned by `SecondaryDatabase.getSecondaryConfig` method. [#16018]
2. Fixed a bug which made it impossible to (re)set `VerboseConfig.REPLICATION_SYSTEM` on the Java API. [#17561]
3. Fixed a bug where populating a `SecondaryDatabase` on open could lead to an `OutOfMemoryException`. [#18529]
4. Fixed a bug such that segmentation fault does not occur when putting records into callback-partitioned database. [#18596]
5. Fixed a bug where `DatabaseConfig.getUnsortedDuplicates` method returned true when the database had been configured for sorted duplicates. [#18612]
6. Initialized `DatabaseConfig.pageSize` so that it can be queried. [#18691]
7. Fixed a bug by opening a write cursor for Direct Persistent Layer(DPL) entity's put operation in the Concurrent Data Store product. [#18692]
8. Synchronized Java persistence code and tests from Java Edition to Berkeley DB. [#18711]
9. Introduced the `EnvironmentConfig.setReplicationInMemory` method as a way to configure in-memory internal replication files before opening the `Environment` handle on the Java API. [#18719]
10. Fixed a bug in the bulk `DatabaseEntry` class, where it was possible to overflow the buffer. [#18850]
11. Added `LEASE_TIMEOUT` field to the `ReplicationTimeoutType` class that enables configuring the amount of time a client grants its Master Lease to a master. [#18867]

C#-Specific API Changes

1. Fixed a bug in `BTree` prefix comparison method such that there is no problem when the application needs to save a number larger than or equal to 2^{31} . The `BTree` prefix comparison function now returns an unsigned int instead of a signed int. [#18481]
2. Fixed a bug which caused the `HasMultiple` method to throw an exception when there were multiple databases in a single database file. [#18483]

3. Fixed a bug to ensure the CachePriority is set for Database and Cursor objects. [#18716]
4. Fixed a bug that use leading to the error: "Transaction that opened the DB handle is still active" when applications used different transactional handles in the associate and open methods in a secondary database. [#18873]

Direct Persistence Layer (DPL), Bindings and Collections API

1. All setter methods in the DPL |StoreConfig| and |EvolveConfig| now return |this| rather than having a |void| return type. This change requires that applications using the DPL be recompiled. [#17021]
2. Improve performance of |StoredCollection.removeAll|. This method no longer iterates over all records in the stored collection. [#17727]
3. Several new tuple formats and binding classes have been added in the |com.sleepycat.bind.tuple| package:
 - Packed integer formats have been added that support default natural sorting. These are intended to replace the old unsorted packed integer formats.
 - Two new |BigDecimal| formats have been added. One format supports default natural sorting. The other format is unsorted, but has other advantages: trailing zeros after the decimal place are preserved, and a more compact, faster serialization format is used. See the |com.sleepycat.bind.tuple| package description for an overview of the new bindings and a comparative description of all tuple bindings. [#18379]
4. The following classes are now certified to be serializable. [#18738]
 - com.sleepycat.persist.IndexNotAvailableException
 - com.sleepycat.persist.StoreExistsException
 - com.sleepycat.persist.StoreNotFoundException
 - com.sleepycat.persist.evolve.DeletedClassException
 - com.sleepycat.persist.evolve.IncompatibleClassException

Replication Changes

1. Replication Manager now uses the standard system implementation of getaddrinfo() when running on Windows, which means that it can support IPv6 addresses if support is present and configured in the operating system. [#18263]
2. Fixed a bug which caused a "full election" to fail if a majority of sites were not ready when the election started. [#18456]

3. Fixed a bug which could occur when using bulk transfer with Replication Manager. When closing a DB_ENV handle, any remaining bulk buffer contents are flushed, and Replication Manager could have tried to send the resulting messages even though its connections had already been closed, leading in rare circumstances to spurious EBADF error reports, or possibly even arbitrary memory corruption. [#18469]
4. Fixed a bug which caused Replication Manager to wait for acknowledgement from client, even if it had failed to send a log record, due to "queue limit exceeded". Replication Manager now returns immediately, with a PERM_FAILED indication, to avoid a pointless delay to the commit() operation. [#18682]
5. Fixed a bug where changes made in one process to Replication Manager configuration values (such as ack policy or ack timeout) were not observed in other processes sharing the same database environment. [#18839]
6. Fixed bugs that could prevent client synchronization from completing due to a failure to request missing log records. [#18849]
7. Fixed a bug where a client that had rolled back transactions to synchronize with a new master, failed to invalidate existing database handles later used for cursor operations based on an explicitly provided transaction. [#18862]
8. Fixed a bug where Replication Manager called for an election after a DUPMASTER event, even when using Master Leases. In such a case it now simply accepts the new (remote) master. [#18864]
9. Fixed a bug which would cause failure if client env attempted to perform sync-up recovery to a point in the log that happened to fall exactly on a log file boundary. [#18907]

Locking Subsystem Changes

1. Moved the wait mutex from the lock structure to the locker structure, reducing the number of mutexes required in the system. [#18685]

Logging Subsystem Changes

1. None.

Memory Pool Subsystem Changes

1. Fixed a race condition that was causing the error: "Unable to allocate space from the buffer cache". The error can only be triggered when multiple memory pool regions are used and there is a periodic gathering and clearing of statistics. This also fixes a second bug where if you compile without statistics and explicitly set the memoru

pool default pagesize, other environment handles to that environment would not see the correct memory pool default pagesize. [#18386]

2. Fixed a bug where the `get_cachesize` method and the `mpool_stat` method returned the initial cache size, even if the cache size had been changed. [#18706]
3. Changed memory pool allocation so that the EIO error is returned rather than the ENOMEM error when the memory cannot be allocated because dirty pages cannot be written. [#18740]

Mutex Subsystem Changes

1. Fixed problems with the printed statistics for DB_MUTEX_SHARED latches. The DB_STAT_CLEAR flag (as specified by `db_stat -Z`) did not clear the counts of the number of times a shared latch either had to wait to get access or was able to get the latch without waiting. Also, the ownership state of a test-and-set latch (not a hybrid one) was always displayed as not owned, even when it was held. [#17585] [#18743]

Transaction Subsystem Changes

1. Fixed bugs that could caused PANIC or DB_RUNRECOVERY errors when the synchronization of the transaction log failed. [#18588]
2. Fix javadoc to note the exception to the rule that a transaction handle may not be accessed after `commit()` operaton (`getCommitToken()` is allowed). [#18730]

Test Suite Changes

1. None. [#18831]

Utility Changes

1. Modified the `db_printlog` and `db_dump -da` so that they use the same formatting. The `db_logprint` utility now uses the message stream. Both `db_dump -da` and `db_printlog` accept a `-D` flag to indicate the numer of bytes of data items to display. You can set this value when calling the `DB_ENV->set_data_len` method or in the DB_CONFIG. [#18365]
2. Fixed a bug that caused a segmentation violation when using the `db_printlog` utility. [#18694]
3. Fixed a bug in `db_hotbackup` that would cause a trap if the `-D` flag is used and the DB_CONFIG file does not specify a log directory. [#18841]

Configuration, Documentation, Sample Apps, Portability, and Build Changes

1. Added support and documentation for iPhone OS. [#18223]
2. Fix a bug to make the configuration option `--enable-debug` work when `CFLAGS` is set. [#18432]
3. Updated Visual Studio project files to enable ADO.NET support. [#18448]
4. Enhanced the source tree layout making it easier to navigate. [#18492]
5. Fixed the `--enable-dbm` argument to configure. [#18497]
6. Fixed Visual Studio project files so that they can load into Visual Studio 2010. [#18505]
7. Updated Windows CE build files to be consistent with desktop Windows build files. Added Windows Mobile 6.5.3 Professional as a target platform. [#18516]
8. Added a fix so that an error message is displayed when the 'ar' utility is missing when configured. [#18619]
9. Added tighter integration of JDBC on POSIX/autoconf by including an argument `--enable-jdbc` to configure. [#18621]
10. Fix build conflicts in log verify with other configurations. [#18658]
11. Upgraded Berkeley DB SQL to SQLite version 3.7.0 [#18857]

Example Changes

1. Renamed `examples/c/bench_001` to `examples/c/ex_bulk`. [#18537]

Miscellaneous Bug Fixes

1. Provided a functionality on the Windows platform to choose a default page size based on the underlying file system sector size. [#16538]
2. Changed `DB_NOSYNC` from an "operation" constant to a flag value. [#17775]
3. Changed the default permissions for files in the release tree to allow write access. [#17974]
4. Fixed a bug which caused database verification to hang when verifying a database in a Concurrent Data Store environment that performs locking on an environment-wide basis (`DB_CDB_ALLDB.`) [#18571]

Deprecated Features

1. [#18871] Removed the mod_db4 PHP/Apache wrapper. It only supported Apache 1.3 and has not been actively supported. Use php_db4 instead.

Known Bugs

1. None

Chapter 12. Upgrading Berkeley DB 4.8 applications to Berkeley DB 11.2.5.0

Table of Contents

[Introduction](#)

[db_sql Renamed to db_sql_codegen](#)

[DB_REP_CONF_NOAUTOINIT Replaced](#)

[Support for Multiple Client-to-Client Peers](#)

[Cryptography Support](#)

[DB_NOSYNC Flag to Flush Files](#)

[Dropped Support](#)

[Changing Stack Size](#)

[Berkeley DB 11g Release 2 Change Log](#)

[Changes between 11.2.5.0.26 and 11.2.5.0.32](#)

[Changes between 11.2.5.0.21 and 11.2.5.0.26](#)

[Changes between 4.8 and 11.2.5.0.21](#)

[Known Bugs](#)

Introduction

The following pages describe how to upgrade applications coded against the Berkeley DB 4.8 release interfaces to the Berkeley DB 11g Release 2 interfaces. (Library version 11.2.5.0). This information does not describe how to upgrade Berkeley DB 1.85 release applications.

db_sql Renamed to db_sql_codegen

The `db_sql` utility is now called `db_sql_codegen`. This command line utility is not built by default. To build `db_sql_codegen`, specify `--enable-sql_codegen` when configuring Berkeley DB.

DB_REP_CONF_NOAUTOINIT Replaced

In this release, the `DB_REP_CONF_NOAUTOINIT` flag is replaced by the `DB_REP_CONF_AUTOINIT` flag. This option is ON by default. To turn off automatic internal initialization, call the `DB_ENV->rep_set_config` method with the `which` parameter set to `DB_REP_CONF_AUTOINIT` and the `onoff` parameter set to zero.

Support for Multiple Client-to-Client Peers

A Berkeley DB Replication Manager application can now designate one or more remote sites (called its "peers") to receive client-to-client requests.

In previous releases, there could be only one peer at a time. If you called the `DB_ENV->repmgr_add_remote_site` method specifying site "A" as a peer and you made another call specifying site "B" as a peer, site "B" would become the only peer, and site "A" would no longer be a peer.

Starting with Berkeley DB 11gR2, the same sequence of calls results in both site "A" and site "B" being possible peers. Replication Manager may select any of a site's possible peers to use for client-to-client requests. If the first peer that the Replication Manager selects cannot be used (for example, it is unavailable or it is the current master), Replication Manager attempts to use a different peer if there is more than one peer.

To get the pre-11gR2 peer behavior in this example, you must make an additional call to the `DB_ENV->repmgr_add_remote_site` method, specifying site "A" and a flag value that excludes the `DB_REPMGR_PEER` bit value to remove site "A" as a possible peer.

Cryptography Support

In this release, the configuration options, `--disable-cryptography` and `--enable-cryptography` are deprecated. `--disable-cryptography` is replaced by `--with-cryptography=no` and `--enable-cryptography` is replaced by `--with-cryptography=yes`.

To build Berkeley DB with support for cryptography, enter `--with-cryptography=yes` as an argument to configure instead of `--enable-cryptography`.

To build Berkeley DB without support for cryptography, enter `--with-cryptography=no` as an argument to configure instead of `--disable-cryptography`.

Berkeley DB now supports encryption using Intel's Performance Primitive (IPP) on Linux. To build Berkeley DB with support for cryptography using Intel's Performance Primitive (IPP) library, enter `--with-cryptography=ipp` as an argument to configure.

Note: The `--with-cryptography=ipp` argument works only on Linux.

DB_NOSYNC Flag to Flush Files

Applications must now pass the `DB_NOSYNC` flag to the methods - `DB->remove`, `DB->rename`, `DB_ENV->dbremove`, and `DB_ENV->dbrename`, to avoid a multi-database file to be flushed from cache. This flag is applicable if you have created the database handle in a non-transactional environment.

By default, all non-transactional database remove/rename operations cause data to be synced to disk. This can now be overridden using the `DB_NOSYNC` flag so that files can be accessed outside the environment after the database handles are closed.

Dropped Support

Berkeley DB no longer supports Visual Studio 6.0. The earliest version supported is Visual Studio 2005. The build files for Windows Visual Studio 6.0 are removed.

Berkeley DB no longer supports Win9X, Windows Me (Millenium edition), and Windows NT 4.0. The minimum supported windows platform is Windows 2000.

Changing Stack Size

Prior to the 11gR2 release, Berkeley DB limited the stack size for threads it created using the POSIX thread API to 128 KB for 32-bit platforms and 256 KB for 64-bit platforms. In this release, the system default stack size is used unless you run the Berkeley DB configure script with the `--with-stacksize=SIZE` argument to override the default.

Berkeley DB 11g Release 2 Change Log

[Changes between 11.2.5.0.26 and 11.2.5.0.32](#)

[Changes between 11.2.5.0.21 and 11.2.5.0.26](#)

[Changes between 4.8 and 11.2.5.0.21](#)

Known Bugs

Changes between 11.2.5.0.26 and 11.2.5.0.32

1. Added Visual Studio 2010 support. Users can find Visual Studio 2010 solutions and projects on build_windows. [#18889]
2. Fixed a leak of log file ids when a database is closed before the end of a transaction that references it. [#15957]
3. Fixed a race condition that was causing an "unable to allocate space from the buffer cache" error. The error can only be triggered when multiple mpool regions are used and there is a periodic gathering and clearing of statistics. This also fixes a second bug where if you compile without statistics and explicitly set the mpool default pagesize, other environment handles to that environment would not see the correct mpool default pagesize. [#18386]
4. Fix failure to flush pages to disk. [#18760]
5. Fix a general I/O problem on Windows where system doesn't always return ENOENT when file is missing. [#18762]
6. Fixed locking bugs: [#18789]
 - Db->compact of BTREE with MVCC could return an unpinned page.
 - RECNO would fail to lock the next page when splitting a leaf page.
7. Don't await ack if message not sent due to queue limit exceeded. [#18682]
8. Fixed a bug that could cause data to not be returned in a HASH database that was one of multiple databases in a file and it was opened prior to running DB->compact on that database in another thread of control [#18824]
9. Return HANDLE_DEAD on cursor creation that names a specific txn after client callback. [#18862]
10. Remove parting_shot rep_start(CLIENT) in election thread because it can occasionally conflict with rep_start(MASTER) in another thread. [#18946]
11. Fixed a bug that would cause handle locks to be left referencing the wrong metadata page if DB->compact moved the metadata page of a sub-database. [#18944]
12. Fixed a bug that might cause an update to a HASH database to fail with an "unpinned page returned" error if it first gets an I/O error while logging. [#18985]

13. Fixed a bug that failed to dirty a page when DB->compact moved records within a hash bucket [#18994]
14. Fixed a bug in page allocation where if a non-transactional update was being done, then we release the metadata page lock too early possibly leading to the corruption of the in memory page list used by DB->compact. [#19036]
15. A log write failure on a replication master will now cause a panic since the transaction may be committed on some clients. [#19054]
16. Removed the possibility that checkpoints will overlap in the log, decreasing the time to recover [#19062]
17. Fixed a bug that could leave a hash bucket overflow page not linked to the bucket if the unlink of that page aborted. [#19001]
18. Fixed a bug that would leave the next page pointer of a hash bucket that was removed pointing to an invalid page. [#19004]
19. Fixed several bugs that could cause an update running with MVCC to get the wrong version of a page or improperly update the metadata last page number. [#19063]
20. Fixed a bug where an error during an update to a hash database with DB_NOOVERWRITE set could return DB_KEYEXIST rather than the correct error. [#19077]
21. Fixed a bug where an updater supporting DB_READ_UNCOMMITTED might downgrade its lock too soon if there was an error during the update [#19155]
22. Fixed a bug that could cause the wrong page number to be on a root or metadata page if DB->compact moved the page and the operation was later rolled forward [#19167]
23. Fixed a bug that could cause the close of a secondary index database to fail if the transaction doing the open aborted [#19169]
24. The database open code will no longer log the open and close of the master database in a file when opening a sub database in that file [#19071]

Changes between 11.2.5.0.21 and 11.2.5.0.26

1. Fixed a bug that might cause recovery to fail if processed part of the log that had previously been recovered and a database which was not present was opened in the log and not closed. [#18459]

2. Fixed a bug which could occur when using bulk transfer with Replication Manager. When closing a DB_ENV handle, any remaining bulk buffer contents are flushed, and Replication Manager could have tried to send the resulting messages even though its connections had already been closed, leading in rare circumstances to spurious EBADF error reports, or possibly even arbitrary memory corruption. [#18469]
3. Fixed a bug in C# HasMultiple() that this function always throws exceptions when there are multiple databases in a single db file. [#18483]
4. Fixed the '--enable-dbm' argument to configure. [#18497]
5. Fixed a bug in the Java API where populating a SecondaryDatabase on open could lead to an OutOfMemoryException. [#18529]
6. Fixed a bug where DB SQL reports "The database disk image is malformed" in "group by" operations. [#18531]
7. Fixed a bug that prevented the same process from reconnecting to the database when DB_REGISTER is being used. [#18535]
8. Fix a race between opening and closing SQL databases from multiple threads that could lead to the error "DB_REGISTER limits processes to one open DB_ENV handle per environment". [#18538]
9. Fixed some bugs that could cause a panic or a DB_RUN_RECOVERY error if the sync of the transaction log failed. [#18588]
10. Fixed a bug which would occur when recovery checkpoint was not written because the cache ran out of space attempting to flush the mpool cache. The environment was recovered and all database were made available, but some databases were incorrectly closed. This would cause a subsequent recovery to fail on its backward pass with the error "PANIC: No such file or directory". [#18590]
11. Fixed a bug that segmentation fault would occur if DB->set_partition_dirs was called before DB->set_partition. [#18591]
12. Fixed a bug that the error of "unknown path" would occur if putting duplicate records to duplicated sorted hash database with DB_OVERWRITE_DUP.[#18607]
13. Fixed a bug where DatabaseConfig.getUnsortedDuplicates() returned true when the database had been configured for sorted duplicates. [#18612]
14. Fixed a bug that could cause recovery to fail with the error "DB_LOGC->get: log record LSN %u/%u: checksum mismatch" if the last log file was nearly full and ended with a partially written log record which was smaller than a checkpoint

record. It now erases the invalid partial record before switching to the new log file. [#18651]

15. Initialize DatabaseConfig.pageSize so that it can be queried from Java. [#18691]

16. Fixed a bug that might cause an aborting transaction to fail if it aborted while a DB->compact of the same HASH database was compacting the dynamic hash table [#18695]

Changes between 4.8 and 11.2.5.0.21

Database or Log File On-Disk Format Changes

1. The log file format changed in 11.2.5.0.21

New Features

1. Replication Manager sites can specify one or more possible client-to-client peers. [#14776]
2. Added resource management feature in all Berkeley DB APIs to automatically manage cursor and database handles by closing them when they are not required, if they are not yet closed. [#16188]
3. Added a SQL interface to the Berkeley DB library. The interface is based on - and a drop-in-replacement for - the SQLite API. It can be accessed via a command line utility, a C API, or existing APIs built for SQLite. [#16809]
4. Added hash databases support to the DB->compact interface. [#16936]
5. Renamed the "db_sql" utility to "db_sql_codegen". This utility is not built by default. To build this utility, enter --enable-sql_codegen as an argument to configure. [#18265]
6. Added transactional support in db_sql_codegen utility. Specify TRANSACTIONAL or NONTRANSACTIONAL in hint comments in SQL statement, db_sql_codegen enable/disable transaction in generated code accordingly. [#17237]
7. Added the feature read-your-writes consistency that allows client application to check, or wait for a specific transaction to be replicated from the master before reading database. [#17323]
8. Added DB log verification feature, accessible via the API and a new utility. This feature can help debugging and analysis. [#17420]

9. Added support for applications to assign master/client role explicitly at any time. Replication Manager can now be configured not to initiate elections. [#17484]
10. Enhanced the DB->compact method so that it can reassign metadata and root pages from subdatabases to lower numbered pages while compacting a database file that contains multiple databases. This feature helps to free the higher numbered pages and truncate the file. [#17554]
11. Added system diagnostic messages that are ON by default. [#17561]
12. Added the feature to assign a priority level to transactions. When resolving a deadlock:
 - if the transactions have differing priority, the lowest priority transaction is aborted
 - if all transactions have the same priority, the same policy that existed before priorities were introduced is used [#17604]
13. Added a feature in which log_archive uses group-wide information for archiving purposes if Replication Manager is in use. [#17664]
14. Added a feature by which the Replication Manager application clients now automatically request any missing information, even when there is no master transaction activity. [#17665]
15. Added support for sharing logs across mixed-endian systems. [#18032]
16. Added an option to specify the first and last pages to the db_dump utility. You can do this by providing -F and -L flags to the db_dump -d option. [#18072]
17. Added Intel Performance Primitive (IPP) AES encryption support. [#18110]
18. Removed support for the configuration option --with-mutex=UNIX/fcntl as of version 4.8. If Berkeley DB was configured to use this type of mutex in an earlier release, switch to a different mutex type or contact Oracle for support. [#18361]

Database Environment Changes

1. Fixed a bug to reflect the correct configuration of the logging subsystem when the DB_ENV->log_set_config method is called with the DB_LOG_ZERO flag in a situation where a DB_ENV handle is open and an environment exists. [#17532]
2. Fixed a bug to prevent memory leak caused when the environment is closed by the named in-memory database in a private database environment which has open named in-memory databases. [#17816]

3. Fixed a race condition in an internal directory-scanning function that returns the ENOENT ("No such file or directory") error, if a file is removed just before a call to stat() or its equivalent. [#17850]

Access Method Changes

1. Fixed a bug to prevent a page in the hash database from carrying the wrong header information when a group allocation is rolled forward by recovery. [#15414]
2. Improved the sort function for bulk put operations. [#17440]
3. Fixed a bug in the DB->compact method to ensure locking of leaf pages when merging higher level interior nodes or when freeing interior nodes when descending to find a non-zero length key. [#17485][#16466]
4. Fixed a bug to prevent a trap if a cursor is opened or closed when another thread is adjusting cursors due to an update in the same database. [#17602]
5. Fixed a bug that incorrectly lead to the error message "library build did not include support for the Hash access method" [#17672]
6. Fixed a bug to ensure that the DB->exists method accepts the DB_AUTO_COMMIT flag. [#17687]
7. In the past, removing a database from a multi-database file that was opened in an environment always caused dirty pages in the file to be flushed from the cache. In this release, there is no implicit flush as part of a DB->remove for handles opened in an environment. Applications that expect the database file to be flushed will need to add an explicit flush. [#17775]
8. Fixed a bug so that the code does not loop if a DB->compact operation processed a 3 or more level non-sorted off page duplicate tree. [#17831]
9. Fixed a bug that could leave pages pinned in the cache if an allocation failed during a DB->compact operation. [#17845]
10. Fixed a bug to ensure sequences are closed when an EntityStore is closed. [#17951]
11. Fixed a bug that prevented retrieval of a non-duplicate record with DB_GET_BOTH_RANGE in hash sorted duplicate db. In a database configured with sorted duplicate support, when the DBcursor->get method is passed the DB_GET_BOTH_RANGE flag, the data item should be retrieved that is the smallest value greater than or equal to the value provided by the data parameter (as determined by the comparison function). [#17997]

12. Fixed a bug that causes the wrong file to be removed if multiple cascading renames are done in the same transaction. [#18069]
13. Fixed a bug to prevent the DB->compact method specified with the DB_AUTO_COMMIT flag from acquiring too many locks. [#18072]
14. Fixed a bug that might cause DB->compact on a DB_BTREE database to get a spurious read lock on the metadata page. If the database was opened non-transactionally the lock would get left behind. [#18257]
15. Fixed a bug that could lead to btree structure corruption if the DB->compact method ran out of locks [#18361]
16. Fixed a bug that would generate an error if a non-BDB file was used to create a database and the DB_TRUNCATE flag was specified. [#18373]
17. Fixed a bug that might cause a trap reading uninitialized memory when backing out a merge of a duplicate tree leaf page during DB->compact. [#18461]

Locking Subsystem Changes

1. Fixed a bug to ensure deadlock detection works even when there are blocked transactions, configured with and without timeouts. [#17555]
2. Fixed a bug to ensure a call to the DB->key_range method from outside a transaction does not lock pages. [#17930]
3. Fixed a bug that could cause a segmentation fault if the lock manager ran out of mutexes [#18428]

Logging Subsystem Changes

1. Limited the size of a log record generated by freeing pages from a database, so that it fits in the log file size. [#17313]

Memory Pool Subsystem Changes

1. Fixed a bug to ensure multiple versions of a buffer are not created when MVCC is not set. [#17495]
2. Fixed a bug to detect if cache size is being set when the cache is not configured. [#17556]

3. Fixed a bug to ensure the error message "unable to allocate space from the buffer cache" generated when there is still some space available, can be cleared by running recovery. [#17630]
4. Fixed a race condition that causes an operation to return EPERM when the buffer cache is nearly filled with pages belonging to recently closed queue extents. [#17840]
5. Fixed a bug that could cause a page needed by a snapshot reader to be overwritten rather than copied when it was freed. [#17973]
6. Enabled set_mp_pagesize to be specified in the DB_CONFIG file. [#18015]
7. Fixed a bug to ensure single-version or obsolete buffers were selected over any intermediate version. [#18114]

Mutex Subsystem Changes

1. Fixed a bug on HP-UX when specifying --with-mutex=HP/msem_init during configure. It would generate the error "TAS: mutex not appropriately aligned" at runtime, when initializing the first mutex. [#17489]
2. Fixed a race condition which could cause unnecessary retrying of btree searches when several threads simulatenously attempted to get a shared latch. [#18078]
3. Exclusive Transactions have been implemented for the SQL API. See the documentation for details on the behavior of this feature. [#17822]

Tcl-specific API Changes

1. Fixed a bug in Tcl API to prevent a segmentation fault from occurring when the get_dbname method is called to get the db name and the db handle is opened without providing either the filename or dbname. [#18037]

C#-specific API Changes

1. Fixed a bug in C# to prevent a System.AccessViolationException from occurring on Windows7 when trying to open new database. [#18422]
2. Fixed a bug in the C# API to make DB_COMPACT consistent with __db_compact in teh C API. [#18246]

API Changes

1. Added the `dbstl_thread_exit` method to release thread specific resources on thread exit. [#17595]
2. Fixed the parser to allow configuration API flags set in the `DB_CONFIG` file to accept an optional ON/OFF string. The `DB_REP_CONF_NOAUTOINIT` flag has been removed. It is replaced by `DB_REP_CONF_AUTOINIT`. However, replication's default behavior remains the same. [#17795]

Replication Changes

1. Fixed bug where a not-in-sync client could service a peer request. [#18279]
2. Fixed bug where page gaps, once filled, would not immediately request the next page gap it finds. This was already fixed for logs. [#18219]
3. Fixed a bug so that only one thread waits for the meta-page lock during internal initialization and broadcasts out the information rather than all threads waiting. Removed the former retry code. [#17871]
4. Added a feature by which the `DB->open` method now allows the `DB_CREATE` flag on a replication client. It is ignored, but this allows a replication application to make one call that can work on either master or client. It fixes a possible race that could develop in a Replication Manager application if a call to `DB->open` is made around the same time as a master/client role change. [#15167]
5. The `DB_ENV->repmgr_site_list` method now returns an indication on whether the site is a client-to-client peer. [#16113]
6. Fixed a bug that could occasionally lead to elections failing to complete. [#17105]
7. Fixed a bug that could cause `DB_ENV->txn_stat` to trap. [#17198]
8. Added a new `JOIN_FAILURE` event to notify Replication Manager applications which refuse auto-initialization. [#17319]
9. Fixed a bug where a failed master lease check at a client site causes an `ASSERT` when processing a master lease grant at the master site. [#17869]
10. Fixed a bug to ensure a second simultaneous call to the `DB_ENV->rep_elect` method does not incorrectly clear a bit flag. [#17875]
11. Fixed a bug in client-side autoremoval of log files. [#17899]
12. Removed the likelihood of dual data streams to enhance network traffic. [#17955]

13. Fixed a bug such that non-txn dup cursors are accounted for in the replication API lockout. [#18080]
14. Fixed a bug to ensure checking for other sync states for the rerequest thread. [#18126]
15. Fixed a bug to avoid getting stuck in an election forever. [#18151]
16. Fixed a bug where using client-to-client synchronization with Master Leases could have resulted in failure of a new master to get initial lease grants from sufficient number of clients, resulting in a master environment panic. [#18254]
17. Fixed a bug which had prevented Replication Manager socket operations from working on HP/UX systems. [#18382]
18. Fixed a bug where starting as a client in multiple threads after receiving dupmaster messages could have resulted in a failure to find a new log file, resulting in a panic. [#18388]
19. The default thread stack size is no longer overridden by default for Berkeley DB threads. [#18383]

Transaction Subsystem Changes

1. Fixed a bug that caused transactions to deadlock on the mutex in the sequence object. [#17731]
2. Fixed a bug to ensure that the failure checking mechanism reconstructs child transactions correctly when a process dies with active sub-transactions. [#18154]
3. Removed a memory leak during recovery related to a deleted database [#18273]

Utility Changes

1. Fixed compiler warnings in the db_sql_codegen utility. [#17503]
2. Enhanced the db_recover -v utility to display the message, "No log files found", if no logs are present. [#17504]
3. Modified the db_verify utility to verify all files instead of aborting on the first failure. [#17513]
4. Modified the db_verify utility to display a message after verification is completed. [#17545]

5. Fixed a bug in the db_sql_codegen utility where the primary key is stored in both key and data fields. Removed it from the data field. [#17925]

Example Changes

1. Fixed a bug that causes the ex_txn C# example to hang. [#17376]
2. Fixed Solaris alignment issues in Stl port test code. [#17459]
3. Added GCC 4.4 compatibility support for all examples. [#17584]
4. Added new command line arguments(-h and -d) to the env examples. [#17624]
5. Fixed configuration problems related to running java API tests. [#17625]
6. Updated the bench_001 example to include bulk testing examples. [#17766]
7. Added a new Stl example to demo advanced feature usage. The Stl test cases referred earlier are replaced by these new examples in the Stl reference document. [#18175]

Deprecated Features

1. The configuration options --disable-cryptography and --enable-cryptoraphy are being deprecated. [#18110]

Configuration, Documentation, Sample Apps, Portability and Build Changes

1. Remove build files for Windows Visual Studio 6.0. [#16848]
2. Added an API, DBENV->db_full_version, to return db full version.
3. Berkeley DB no longer supports Win9X, Windows Me (Millenium edition) and NT 4.0. The minimum supported windows platform is Win 2k.
4. Berkeley DB no longer supports Visual Studio 6.0. The earliest version supported is Visual Studio 2005.
5. Added "+u1" to CFLAGS for HP ANSI C Compiler on HP-UX(IA64) to fix the alignment issue found with the allocation functions DB->set-alloc and DB_ENV->set_alloc. [#17257]
6. Fixed a bug such that the thread local storage (TLS) definition modifier is correctly deduced from the m4 script on all platforms. [#17609][#17713]

7. Fixed a bug such that TLS key is not initialized on platforms which do not support thread local storage (TLS) keywords, such as MAC OSX, and where TLS is implemented using pthread API. [#18001]
8. Fixed a bug to ensure that when using Intel C++ compiler (icpc), the TLS code builds successfully. A stricter criteria is adopted to deduce the TLS keyword, and hence pthread API is more likely to be used to implement TLS. [#18038]
9. Adding new configuration option, `--with-cryptography={yes|no|ipp}`. Using `--with-cryptography=yes`, will give equivalent behavior to the old `--enable-cryptography` option. Using `--with-cryptography=no`, will give equivalent behavior to the old `--disable-cryptography` option. Using `--with-cryptograhpy=ipp` will enable Intel's Performance Primitive (IPP) encryption on linux. [#18110]

Known Bugs

1. The configure option `--with-uniquename` may cause macro redefinition warnings on platforms where BDB implements parts of the standard C library. These warnings (e.g., `"db_int_def.h", line 586: warning: macro redefined: strsep'`) may occur when functions in the `"clib"` directory are included during configuration. This cosmetic affect does not affect the correct operation of the library. [#17172]
2. A multithreaded application using a private environment and multi-version concurrency control could, on very rare occasions, generate an illegal pointer access error during the final steps of a clean environment shutdown. [#17507]
3. Although rare, it is possible for a partial log record header at the end of a transaction log to be erroneously accepted as if it were valid, causing the error `"Illegal record type 0 in log"` during recovery. [#17851]
4. It is possible to get the error `"unable to allocate space from the buffer cache"` when there are disk errors on the freezer files used by multi-version concurrency control . [#17902]
5. Java API does not support partitioning by keys and the C# API doesn't support partitioning. [#18350]
6. If a database is removed from an environment and it was still opened transactionally and recovery is run, then a future recovery that must process that part of the log may fail. [#18459]
7. Replication `"bulk transfer"` does not work if Berkeley DB is unable to determine, at environment open time, whether the Replication Manager will be used. To work around this problem, an application using the Replication Manager should call `DB_ENV->repmgr_set_local_site()` before opening the environment. An application

using the replication Base API should call `DB_ENV->rep_set_transport()` before opening the environment. [#18476]

8. The BTree prefix comparison function behaves slightly differently in the C API vs the C# API. In the C# API it returns a signed int and in the C API it returns an unsigned int. This can be a problem if the application needs to save more than 2^{31} bytes.

Chapter 13. Upgrading Berkeley DB 4.7 applications to Berkeley DB 4.8

Table of Contents

[Introduction](#)

[Registering DPL Secondary Keys](#)

[Minor Change in Behavior of `DB\_MPOOLFILE->get`](#)

[Dropped Support for `fcntl` System Calls](#)

[Upgrade Requirements](#)

[Berkeley DB 4.8.28 Change Log](#)

[Changes between 4.8.26 and 4.8.28:](#)

[Known bugs in 4.8](#)

[Changes between 4.8.24 and 4.8.26:](#)

[Changes between 4.8.21 and 4.8.24:](#)

[Changes between 4.7 and 4.8.21:](#)

[Database or Log File On-Disk Format Changes:](#)

[New Features:](#)

[Database Environment Changes:](#)

[Concurrent Data Store Changes:](#)

[General Access Method Changes:](#)

[Btree Access Method Changes:](#)

[Hash Access Method Changes:](#)

Queue Access Method Changes:

Recno Access Method Changes:

C-specific API Changes:

C++-specific API Changes:

Java-specific API Changes:

Direct Persistence Layer (DPL), Bindings and Collections API:

Tcl-specific API Changes:

RPC-specific Client/Server Changes:

Replication Changes:

XA Resource Manager Changes:

Locking Subsystem Changes:

Logging Subsystem Changes:

Memory Pool Subsystem Changes:

Mutex Subsystem Changes:

Test Suite Changes

Transaction Subsystem Changes:

Utility Changes:

Configuration, Documentation, Sample Application, Portability and Build Changes:

Introduction

The following pages describe how to upgrade applications coded against the Berkeley DB 4.7 release interfaces to the Berkeley DB 4.8 release interfaces. This information does not describe how to upgrade Berkeley DB 1.85 release applications.

Registering DPL Secondary Keys

Entity subclasses that define secondary keys must now be registered prior to storing an instance of the class. This can be done in two ways:

- The `EntityModel.registerClass()` method may be called to register the subclass before opening the entity store.
- The `EntityStore.getSubclassIndex()` method may be called to implicitly register the subclass after opening the entity store.

Failure to register the entity subclass will result in an `IllegalArgumentException` the first time an attempt is made to store an instance of the subclass. An exception will not occur if instances of the subclass have previously been stored, which allows existing applications to run unmodified in most cases.

This behavioral change was made to increase reliability. In several cases, registering an entity subclass has been necessary as a workaround. The requirement to register the subclass will ensure that such errors do not occur in deployed applications.

Minor Change in Behavior of `DB_MPOOLFILE->get`

DB 4.8 introduces some performance enhancements, based on the use of shared/exclusive latches instead of locks in some areas of the internal buffer management code. This change will affect how the `DB_MPOOL` interface handles dirty buffers.

Because of these changes, `DB_MPOOLFILE->get` will now acquire an exclusive latch on the buffer if the `DB_MPOOL_DIRTY` or `DB_MPOOL_EDIT` flags are specified. This could lead to an application deadlock if the application tries to fetch the buffer again, without an intervening `DB_MPOOLFILE->put` call.

If your application uses the `DB_MPOOL` interface, and especially the `DB_MPOOL_DIRTY` and `DB_MPOOL_EDIT` flags, you should review your code to ensure that this behavior change does not cause your application to deadlock.

Dropped Support for `fcntl` System Calls

Berkeley DB no longer supports mutex implementations based on the `fcntl` system call. If you have been configuring Berkeley DB to use this type of mutex, you need to either switch to a different mutex type or contact the Berkeley DB team for support.

Upgrade Requirements

The log file format changed in the Berkeley DB 4.8 release.

No database formats changed in the Berkeley DB 4.8 release.

The Berkeley DB 4.8 release does not support live replication upgrade from the 4.2 or 4.3 releases, only from the 4.4 and later releases.

For further information on upgrading Berkeley DB installations, see the [Berkeley DB Upgrade Guide](#).

Berkeley DB 4.8.28 Change Log

[Changes between 4.8.26 and 4.8.28:](#)

[Known bugs in 4.8](#)

[Changes between 4.8.24 and 4.8.26:](#)

[Changes between 4.8.21 and 4.8.24:](#)

[Changes between 4.7 and 4.8.21:](#)

[Database or Log File On-Disk Format Changes:](#)

[New Features:](#)

[Database Environment Changes:](#)

[Concurrent Data Store Changes:](#)

[General Access Method Changes:](#)

[Btree Access Method Changes:](#)

[Hash Access Method Changes:](#)

[Queue Access Method Changes:](#)

[Recno Access Method Changes:](#)

[C-specific API Changes:](#)

[C++-specific API Changes:](#)

[Java-specific API Changes:](#)

Direct Persistence Layer (DPL), Bindings and Collections API:

Tcl-specific API Changes:

RPC-specific Client/Server Changes:

Replication Changes:

XA Resource Manager Changes:

Locking Subsystem Changes:

Logging Subsystem Changes:

Memory Pool Subsystem Changes:

Mutex Subsystem Changes:

Test Suite Changes

Transaction Subsystem Changes:

Utility Changes:

Configuration, Documentation, Sample Application, Portability and Build Changes:

Changes between 4.8.26 and 4.8.28:

1. Limit the size of a log record generated by freeing pages from a database so it fits in the log file size. [#17313]
2. Fix a bug that could cause a file to be removed if it was both the source and target of two renames within a transaction. [#18069]
3. Modified how we go about selecting a usable buffer in the cache. Place more emphasis on single version and obsolete buffers. [#18114]

Known bugs in 4.8

1. Sharing logs across mixed-endian systems does not work.[#18032]

Changes between 4.8.24 and 4.8.26:

1. Fixed a bug where the truncate log record could be too large when freeing too many pages during a compact. [#17313]
2. Fixed a bug where the deadlock detector might not run properly. [#17555]

3. Fixed three bugs related to properly detecting thread local storage for DbStl. [#17609] [#18001] [#18038]
4. Fixed a bug that prevented some of our example code from running correctly in a Windows environment. [#17627]
5. Fixed a bug where a "unable to allocate space from buffer cache" error was improperly generated. [#17630]
6. Fixed a bug where DB->exists() did not accept the DB_AUTO_COMMIT flag. [#17687]
7. Fixed a bug where DB_TXN_SNAPSHOT was not getting ignored when DB_MULTIVERSION not set. [#17706]
8. Fixed a bug that prevented callback based partitioning through the Java API. [#17735]
9. Fixed a replication bug where log files were not automatically removed from the client side. [#17899]
10. Fixed a bug where code generated from db_sql stored the key in both the data and key DBTs. [#17925]
11. Fixed a bug that prevented a sequence from closing properly after the EntityStore closed. [#17951]
12. Fixed a bug where gets fail if the DB_GET_BOTH_FLAG is specified in a hash, sorted duplicates database. [#17997]

Changes between 4.8.21 and 4.8.24:

1. Fixed a bug in the C# API where applications in a 64-bit environment could hang. [#17461]
2. Fixed a bug in MVCC where an exclusive latch was not removed when we couldn't obtain a buffer. [#17479]
3. Fixed a bug where a lock wasn't removed on a non-transactional locker. [#17509]
4. Fixed a bug which could trigger an assertion when performing a B-tree page split and running out of log space or with MVCC enabled. [#17531]
5. Fixed a bug in the repquote example that could cause the application to crash. [#17547]

6. Fixed a couple of bugs when using the GCC 4.4 compiler to build the examples and the dbstl API. [#17504] [#17476]
7. Fixed an incorrect representation of log system configuration info. [#17532]

Changes between 4.7 and 4.8.21:

Database or Log File On-Disk Format Changes:

1. The log file format changed in 4.8.

New Features:

1. Improved scalability and throughput when using BTree databases especially when running with multiple threads that equal or exceed the number of available CPUs.
2. Berkeley DB has added support for C#. In addition to the new C# api, C# specific tests and sample applications were also added. [#16137]
3. Berkeley DB has added an STL API, which is compatible with and very similar to C++ Standard Template Library (STL). Tests and sample applications and documentation were also added. [#16217]
4. Berkeley DB has added database partitioning. BTree or Hash databases may now be partitioned across multiple directories. Partitioned databases can be used to increase concurrency and to improve performance by spreading access across disk subsystems. [#15862]
5. Berkeley DB now supports bulk insertion and deletion of data. Similar to the bulk get interface, the bulk put and bulk delete allow the developer to populate a buffer of key-value pairs and then pass it to the BDB library with a single API call.
6. Berkeley DB now supports compression when using BTree.
7. Berkeley DB introduces a new utility named db_sql which replaces db_codegen. Similar to db_codegen, db_sql accepts an input file with DDL statements and generates a Berkeley DB application using the C API that creates and performs CRUD operations on the defined tables. The developer can then use that code as a basis for further application development.
8. The Replication Manager now supports shared access to the Master database environment from multiple processes. In earlier versions, multiple process support on the Master required use of the Base Replication API. [#15982]
9. Foreign Key Support has been added to Berkeley DB.

10. Several enhancements were made to DB_REGISTER & DB_ENV->failchk().
11. Berkeley now supports 100% in-memory replication.
12. Berkeley DB now has the ability to compare two cursors for equality. [#16811]

Database Environment Changes:

1. Fixed a bug that could cause an allocation error while trying to allocate thread tracking information for the DB_ENV->failcheck system. [#16300]
2. Fixed a bug that could cause a trap if an environment open failed and failchk thread tracking was enabled.[#16770]

Concurrent Data Store Changes:

None.

General Access Method Changes:

1. Fixed a bug where doing an insert with secondary indices and the NOOVERWRITE flag could corrupt the secondary index. [#15912]
2. Fixed a possible file handle leak that occurred while aborting the create of a database whose metadata page was not initialized. [#16359]
3. Fixed a bug so that we now realloc the filename buffer only if we need it to grow. [#16385] [#16219]
4. Fixed a race freeing a transaction object when using MVCC. [#16381]
5. Added missing get methods for the DB and DB_ENV classes where there already was a corresponding set method. [#16505]
6. Fixed a bug to now ensure that DB_STAT_SUBSYSTEM is distinct from other stat flags. [#16798]
7. Fixed a bug related to updating multiple secondary keys (using DB_MULTIPLE). [#16885]
8. Fixed a bug so that verify (db->verify, db_verify) will now report when it cannot read a page rather than just saying the database is bad. [#16916]
9. Fixed a bug that could cause memory corruption if a transaction allocating a page aborted while DB->compact was running on that database. [#16862]

10. Fixed a bug where logging was occurring during remove of an in-memory database when the DB_TXN_NOT_DURABLE flag was set. [#16571]
11. Fixed a bug to remove a race condition during database/file create. [#17020]
12. Fixed a bug where a call to DB->verify and specifying DB_SALVAGE could leak memory when the call returned. [#17161]
13. Fixed a bug to avoid accessing freed memory during puts on primaries with custom comparators. [#17189]
14. Fixed a bug that could cause old versions of pages to be written over new versions if an existing database is opened with the DB_TRUNCATE flag. [#17191]

Btree Access Method Changes:

1. Fixed a bug which could cause DB->compact to fail with DB_NOTFOUND or DB_PAGE_NOTFOUND if the height of the tree was reduced by another thread while compact was active. The bug could also cause a page split to trigger splitting of internal nodes which did not need to be split. [#16192]
2. Fixed a bug that caused Db->compact to loop if run on an empty RECNO database when there were pages in the free list. [#16778]
3. Added a new flag, DB_OVERWRITE_DUP, to DB->put and DBC->put. This flag is equivalent to DB_KEYLAST in almost all cases: the exception is that with sorted duplicates, if a matching key/data pair exists, we overwrite it rather than returning DB_KEYEXIST. [#16803]

Hash Access Method Changes:

1. Fixed a bug to now force a group allocation that rolls forward to reinit all the pages. Otherwise a previous aborted allocation may change the header. [#15414]
2. Fixed a bug to now return the expected buffer size on a DB_BUFFER_SMALL condition. [#16881]

Queue Access Method Changes:

1. Fixed a bug that would cause the LSN reset functionality to not process queue extents. [#16213]
2. Fixed a bug that prevented a partial put on a queue database with secondaries configured. [#16460]

3. Fixed a bug to now prevent an unpinned page to be returned if a delete from a HASH database deadlocked. [#16371]
4. Fixed a bug that could cause a queue extent to be recreated if an application deleted a record that was already deleted in that extent. [#17004]
5. Added the DB_CONSUME flag to DB->del and DBC->del to force adjustment of the head of the queue. [#17004]

Recno Access Method Changes:

1. Fixed a bug which could cause DB->compact of a RECNO database to loop if the number of pages on the free list was reduced by another thread while compact was active. [#16199]
2. Fixed a bug that occurs when deleting from a Recno database and using DB_READ_UNCOMMITTED where we could try to downgrade a lock twice. [#16347]
3. Fixed a bug to now disallow passing DB_DUP and DB_RECNUM together to __db_set_flags. [#16585]

C-specific API Changes:

1. Add get functions for each set functions of DB and DB_ENV structures which didn't have one. [#16505]

C++-specific API Changes:

1. The get and set_lk_partitions methods are now available.
2. Add get functions for each set functions of Db and DbEnv classes which didn't have one. [#16505]
3. Fixed a memory leak when using nested transactions. [#16956]

Java-specific API Changes:

1. Fixed a bug where the replication finer-grained verbose flags were not available in the Java API. [#15419]
2. Fixed a bug in the BTree prefix compression API when called from the Java API. DBTs were not properly initialized. [#16417]
3. Fixed a bug so that LogCursor will work correctly from the Java API. [#16827]

4. Fixed a bug so that position(), limit() and capacity() of ByteBuffers are obeyed by DatabaseEntry objects. [#16982]

Direct Persistence Layer (DPL), Bindings and Collections API:

1. The StoredMap class now implements the standard java.util.concurrent.ConcurrentMap interface. [#15382]
2. Report a meaningful IllegalArgumentException when @Persistent is incorrectly declared on an enum class. Before, the confusing message Persistent class has non-persistent superclass: java.lang.Enum was reported. [#15623]
3. Report a meaningful IllegalArgumentException when @Persistent is incorrectly declared on an interface. Before, a NullPointerException was reported. [#15841]
4. Several validation checks have been added or corrected having to do with entity subclasses, which are @Persistent classes that extend an @Entity class. [#16077]
5. Optimized marshaling for large numbers of embedded objects improving performance. [#16198]
6. The StoredMap class now implements the Java 1.5 ConcurrentMap interface. [#16218]
7. Fix a DPL bug that caused exceptions when using a class Converter for an instance containing non-simple fields. [#16233]
8. Add EntityCursor.setCacheMode and getCacheMode. See the com.sleepycat.je.CacheMode class for more information. [#16239]
9. Fix a bug that prevents evolution of @SecondaryKey information in an entity subclass (a class that extends an @Entity class). [#16253]
10. Report a meaningful IllegalArgumentException when @Persistent or @Entity is incorrectly used on an inner class (a non-static nested class). Before, the confusing message No default constructor was reported. [#16279]
11. Improved the reliability of Entity subclasses that define secondary keys by requiring that they be registered prior to storing an instance of the class. [#16399]
12. Fix a bug that under certain circumstances causes "IllegalArgumentException: Not a key class" when calling EntityStore.getSubclassIndex, EntityStore.getPrimaryConfig, EntityStore.getSecondaryConfig, or PrimaryIndex.put, and a composite key class is used. [#16407]
13. Fixed a bug so that one can now compile DPL in the Java API on Windows. [#16570]

14. The `com.sleepycat.collections.TransactionRunner.handleException` method has been added to allow overriding the default transaction retry policy. See the javadoc for this method for more information. [#16574]
15. Fix a bug that causes an assertion to fire or a `NullPointerException` (when assertions are disabled) from the `EntityStore` constructor. The problem occurs only when the previously created `EntityStore` contains an entity with a secondary key definition in which the key name has been overridden and is different than the field name. [#16819]
16. Key cursors have been optimized to significantly reduce I/O when the `READ_UNCOMMITTED` isolation mode is used. See `EntityIndex.keys` for more information. [#16859]
17. Report a meaningful `IllegalArgumentException` when `NULLIFY` is used with a `@SecondaryKey` and the field is a primitive type. Before, the confusing message `Key field object may not be null` was reported. [#17011]
18. Enum fields may now be used as DPL keys, including primary keys, secondary keys, and fields of composite key classes. Comparators are supported for composite key classes containing enum fields. [#17140]
19. Fix a bug that prevented the use of custom key comparisons (composite key classes that implement `Comparable`) for secondary keys defined as `ONE_TO_MANY` or `MANY_TO_MANY`. [#17207]
20. The `db.jar` file now contains a `Premain` class which enables bytecode enhancement using the JVM instrumentation commands. The built-in proxy classes are also now enhanced in the `db.jar` file, which enables off-line bytecode enhancement. For more information on DPL bytecode enhancement and how to use both instrumentation and off-line enhancement, please see the `com.sleepycat.persist.model.ClassEnhancer` javadoc. [#17233]

Tcl-specific API Changes:

1. The mutex API is now available when using Tcl. [#16342]

RPC-specific Client/Server Changes:

- RPC support has been removed from Berkeley DB. [#16785]

Replication Changes:

1. Improved testing of initial conditions for `rep` and `repmgr` APIs and added heartbeat timeouts to `rep_get_timeout`. [#14977]

2. Added DB_REP_CONF_INMEM replication configuration flag to store replication information exclusively in-memory without creating any files on-disk. [#15257]
3. Added repmgr support for multi-process shared env [#15982]
4. Fixed a bug where opening a cursor from a database handle failed to check whether the database handle was still fresh. If the database handle had been invalidated by a replication client synchronizing with a new master, it could point to invalid information. [#15990]
5. Fixed a bug so that if LOG_REQ gets an archived LSN, replication sends VERIFY_FAIL. [#16004]
6. Added timestamp and process/thread id to replication verbose messages. [#16098]
7. Fixed a bug where, in very rare circumstances, two repmgr sites could connect to each other at the exact same time, the connection attempts "collide" and fail, and the same collision repeats in time synchronization indefinitely. [#16114]
8. Fixed a bug where a missing database file (FILE_FAIL error condition) can interrupt a client synchronization without restarting it. [#16130]
9. Fixed a bug by adding REP_F_INREPSTART flag to prevent racing threads in rep_start. [#16247]
10. Fixed a bug to not return HOLDELECTION if we are already in the middle of an election. Updated the egen so the election thread will notice. [#16270]
11. Fixed a bug in buffer space computation, which could have led to memory corruption in rare circumstances, when using bulk transfer. [#16357]
12. Fixed a bug that prevented replication clients from opening a sequence. The sequence is opened for read operations only. [#16406]
13. Fixed a bug by removing an assertion about priority in elections. It is not correct because it could have changed by then. Remove unused recover_gen field. [#16412]
14. Fixed a bug to now ignore a message from client if it is an LSN not recognized in a LOG_REQ. [#16444]
15. Fixed a bug so that on POSIX systems, repmgr no longer restores default SIGPIPE action upon env close, if it was necessary to change it during start-up. This allows remaining repmgr environments within the same process, if any, to continue operating after one of them is closed. [#16454]

16. After a replication client restarts with recovery, any named in-memory databases are now re-materialized from the rest of the replication group upon synchronization with the master. [#16495]
17. Fixed a bug by adding missing rep_get_config flags. [#16527]
18. Instead of sleeping if the bulk buffer is in transmission, return so that we can send as a singleton. [#16537]
19. Fixed a bug by changing __env_refresh to not hit assert on -private -rep env with an in-memory database. [#16546]
20. Fixed a bug in the Windows implementation of repmgr where a large number of commit threads concurrently awaiting acknowledgments could result in memory corruption, and leaking Win32 Event Objects. [#16548]
21. Fixed a bug by changing repmgr to count a dropped connection when noticing a lacking heartbeat; fixed heartbeat test to check for election, rather than connection drop count, and more reasonable time limit; fixed test to poll until desired result, rather than always sleeping max possible time. [#16550]
22. Fixed "master changes" stat to count when local site becomes master too. [#16562]
23. Fixed a bug where a c2c client would send UPDATE_REQ to another client [#16592]
24. Removed code to proactively expire leases when we don't get acks. Leases maintain their own LSNs to know. [#16494]
25. Fixed a bug where a client may not sync pages during internal init. [#16671]
26. Fixed a bug where a client that received and skipped a log record from the master during an election, then won the election, could then try to request a copy of the skipped log record. The result was an attempt to send a request to the local site, which is invalid: this could confuse a replication Base API application, or cause the Replication Manager to crash. [#16700]
27. Fixed a bug which could have caused data loss or corruption (at the client only) if a replication client rolled back existing transactions in order to synchronize with a new master, and then crashed/recovered before a subsequent checkpoint operation had been replicated from the master. [#16732]
28. Fixed a bug so that replication now retries on DB_LOCK_NOTGRANTED. [#16741]
29. Fixed a potential deadlock in rep_verify_fail. [#16779]

30. Fixed a bug so that an application will no longer segv if nsites given was smaller than number of sites that actually exists. [#16825]

XA Resource Manager Changes:

1. The XA Resource Manager has been removed from Berkeley DB. [#6459]

Locking Subsystem Changes:

1. Fixed a bug to prevent unlocking a mutex twice if we ran out of transactional locks. [#16285]
2. Fixed a bug to prevent a segmentation trap in __lock_open if there were an error during the opening of an environment. [#16307]
3. Fixed a bug to now avoid a deadlock if user defined locks are used only one lock partition is defined. [#16415]
4. Fixed concurrency problems in __dd_build, __dd_abort by adding LOCK_SYSTEM_LOCK() calls to __dd_build and __dd_abort. [16489]
5. Fixed a bug that could cause a panic if a transaction which updated a database that was supporting READ_UNCOMMITTED readers aborted and it hit a race with a thread running the deadlock detector. [#16490]
6. Fixed a race condition in deadlock detection that could overwrite heap. [#16541]
7. Fixed a bug so that DB_STAT_CLEAR now restores the value of st_partitions. [#16701]

Logging Subsystem Changes:

1. Fixed a bug so that the header checksum is only ignored when the log is from a previous version [#16281]
2. Fixed a bug by removing a possible race condition with logc_get(DB_FIRST) and log archiving. [#16387]
3. Fixed a bug that could cause a recovery failure of a create of a database that was aborted. [#16824]
4. An in-memory database creation has an intermediate phase where we have a semi-open DBP. If we crash in that state, then recovery was failing because it tried to use a partially open database handle. This fix checks for that case, and avoids trying to undo page writes for databases in that interim step. [#17203]

Memory Pool Subsystem Changes:

1. Fixed a bug that occurred after all open handles on a file are closed. Needed to clear the TXN_NOT_DURABLE flag (if set) and mark the file as DURABLE_UNKNOWN in the memory pool. [#16091]
2. Fixed a possible race condition between dirtying and freeing a buffer that could result in a panic or corruption. [#16530]
3. Fixed a memory leak where allocated space for temporary file names are not released. [#16956]

Mutex Subsystem Changes:

1. Fixed a bug when using mutexes for SMP MIPS/Linux systems. [#15914]
2. POSIX mutexes are now the default on Solaris. [#16066]
3. Fixed a bug in mutex allocation with multiple cache regions. [#16178]
4. Fixed MIPS/Linux mutexes in 4.7. [#16209]
5. Fixed a bug that would cause a mutex to be unlocked a second time if we ran out of space while tracking pinned pages. [#16228]
6. Fixed a bug Sparc/GCC when using test-and-set mutexes. They are now aligned on an 8-byte boundary. [#16243]
7. Fixed a bug to now prevent a thread calling DB_ENV->failcheck to hang on a mutex held by a dead thread. [#16446]
8. Fixed a bug so that __db_pthread_mutex_unlock() now handles the failchk case of finding a busy mutex which was owned by a now-dead process. [#16557]
9. Removed support for the mutex implementation based on the "fcntl" system call. Anyone configuring Berkeley DB to use this type of mutex in an earlier release will need to either switch to a different mutex type or contact Oracle for support. [#17470]

Test Suite Changes

1. Fixed a bug when using failchk(), where a mutex was not released. [#15982]
2. Added a set of basic repmgr tests to run_std and run_all. [#16092]
3. Added control wrapper for db_reptest to test suite. [#16161]

4. Fixed a bug to now skip tests if db_reptest is not configured. [#16161]
5. Changed name of run_db_in_mem to run_inmem_db, and run_inmem to run_inmem_log and made the arg orders consistent. [#16358]
6. Fixed a bug to now clean up stray handles when rep_verify doesn't work. [#16390]
7. Fixed a bug to avoid db_reptest passing the wrong flag to repmgr_start when there is already a master. [#16475]
8. Added new tests for abbreviated internal init. Fixed test not to expect in-memory database to survive recovery. [#16495]
9. Fix a bug, to add page size for txn014 if the default page size is too small. Move files instead of renaming directory for env015 on QNX. [#16627]
10. Added new rep088 test for log truncation integrity. [#16732]
11. Fixed a bug by adding a checkpoint in rep061 to make sure we have messages to process. Otherwise we could hang with client stuck in internal init, and no incoming messages to trigger rerequest. [#16781]

Transaction Subsystem Changes:

1. Fixed a bug to no longer generate an error if DB_ENV->set_flags (DB_TXN_NOSYNC) was called after the environment was opened. [#16492]
2. Fixed a bug to remove a potential hang condition in replication os_yield loops when DB_REGISTER used with replication by adding PANIC_CHECKS. [#16502]
3. Fix a bug to now release mutex obtained before special condition returns in __db_cursor_int and __txn_record_fname. [#16665]
4. Fixed a leak in the transaction region when a snapshot update transaction accesses more than 4 databases. [#16734]
5. Enabled setting of set_thread_count via the DB_CONFIG file. [#16878]
6. Fixed a mutex leak in some corner cases. [#16665]

Utility Changes:

1. The db_stat utility with the -RA flags will now print a list of known remote replication flags when using repmgr. [#15484]
2. Restructured DB salvage to walk known leaf pages prior to looping over all db pages. [#16219]

3. Fixed a problem with upgrades to 4.7 on big endian machines. [#16411]
4. Fixed a bug so that now db_load consistently returns >1 on failure. [#16765]
5. The db_dump utility now accepts a "-m" flag to dump information from a named in-memory database. [#16896]
6. Fixed a bug that would cause db_hotbackup to fail if a database file was removed while it was running. [#17234]

Configuration, Documentation, Sample Application, Portability and Build Changes:

1. Fixed a bug to now use the correct Perl include path. [#16058]
2. Updated the version of the Microsoft runtime libraries shipped. [#16058]
3. Upgraded the Visual Studio build files to be based on Visual Studio 8 (2005+). The build is now simplified. Users can still upgrade the Visual Studio 6.0 project files, if they want to use Visual Studio .NET (7.1) [#16108]
4. Expanded the ex_rep example with checkpoint and log archive threads, deadlock detection, new options for acknowledgment policy and bulk transfer, and use of additional replication features and events. [#16109]
5. Fixed a bug so that optimizations on AIX are re-enabled, avoiding incorrect code generation. [#16141]
6. Removed a few compiler warnings and three type redefinitions when using vxworks and the GNU compiler. [#16341]
7. Fixed a bug on Sparc v9 so that MUTEX_MEMBAR() now uses membar_enter() to get a #storeload barrier rather than just stbar's #storestor. [#16468]
8. Berkeley DB no longer supports Win9X and Windows Me (Millenium edition).
9. Fixed lock_get and lock_vec examples from the Java (and C#) API. Updated the Java lock example. [#16506]
10. Fixed a bug to correctly handle the TPC-B history record on 64-bit systems. [#16709]
11. Add STL API to Linux build. Can be enabled via the --enable-stl flag. [#16786]
12. Add STL API to Windows build, by building the db_stl project in the solution. There are also stl's test and examples projects in this solution. [#16786]

13. Add support to build dll projects for WinCE, in order to enable users to build DB into a dll in addition to a static library.[#16625]
14. Fixed a weakness where several malloc/realloc return values are not checked before use.[#16664]
15. Enabled DB->compact for WinCE.[#15897]
16. HP-UX 10 is no longer supported.

Chapter 14. Test Suite

Table of Contents

[Running the test suite](#)

[Running SQL Test Suite on Unix](#)

[Running SQL Test Suite on Windows](#)

[Test suite FAQ](#)

Running the test suite

[Running SQL Test Suite on Unix](#)

[Running SQL Test Suite on Windows](#)

Once you have started tclsh and have loaded the test.tcl source file (see [Running the test suite under UNIX](#) and [Running the test suite under Windows](#) for more information), you are ready to run the test suite. At the tclsh prompt, to run the standard test suite, enter the following:

```
% run_std
```

A more exhaustive version of the test suite runs all the tests several more times, testing encryption, replication, and different page sizes. After you have a clean run for run_std, you may choose to run this lengthier set of tests. At the tclsh prompt, enter:

```
% run_all
```

Running the standard tests can take from several hours to a few days to complete, depending on your hardware, and running all the tests will take at least twice as long. For this reason, the output from these commands are redirected to a file in the current

directory named `ALL.OUT`. Periodically, a line will be written to the standard output, indicating what test is being run. When the test suite has finished, a final message will be written indicating the test suite has completed successfully or that it has failed. If the run failed, you should review the `ALL.OUT` file to determine which tests failed. Errors will appear in that file as output lines, beginning with the string "FAIL".

Tests are run in the directory `TESTDIR`, by default. However, the test files are often large, and you should use a filesystem with at least several hundred megabytes of free space. To use a different directory for the test directory, edit the file `include.tcl` in your build directory, and change the following line to a more appropriate value for your system:

```
set testdir ./TESTDIR
```

For example, you might change it to the following:

```
set testdir /var/tmp/db.test
```

Alternatively, you can create a symbolic link named `TESTDIR` in your build directory to an appropriate location for running the tests. Regardless of where you run the tests, the `TESTDIR` directory should be on a local filesystem. Using a remote filesystem (for example, an NFS mounted filesystem) will almost certainly cause spurious test failures.

Running SQL Test Suite on Unix

Once the test suite is built (see [Building SQL Test Suite on Unix](#) for more information), run the entire test suite by executing the following command in the `../build_unix/sql` directory:

```
sh ../../test/sql/bdb-test.sh
```

This runs a set of tests and lists the errors each test encountered, if any. A detailed list of the test results is written to `test.log`.

To run an individual test, such as `insert.test`, execute the following command in the `../build_unix/sql` directory:

```
./testfixture ../../lang/sql/sqlite/test/insert.test
```

Running SQL Test Suite on Windows

After the test suite is built (see [Building the software needed by the SQL tests](#) for more information) and before running the entire test suite, go to `../sql/adapters/bdb-test.sh` and edit the line:

```
echo $t: `alarm $TIMEOUT ./testfixture.exe
$tpath 2>&1 | tee -a test.log | grep "errors out of"
|| echo "failed"``
```

to

```
echo $t: `alarm $TIMEOUT Win32/Debug/testfixture.exe
$tpath 2>&1 | tee -a test.log | grep "errors out of"
|| echo "failed"``
```

Running the test suite requires an Unix emulator, such as Cygwin. In a Cygwin window go to the `../build_windows` directory and execute the command:

```
sh ../sql/adaptor/bdb-test.sh
```

This runs a set of tests and lists errors that each test encountered, if any. A detailed list of the test results is written to `test.log`.

To run an individual test, such as `insert.test`, execute the following command in the `../build_windows` directory:

```
Win32/Debug/testfixture.exe ../sql/sqlite/test/insert.test
```

Test suite FAQ

1. The test suite has been running for over a day. What's wrong?

The test suite can take anywhere from some number of hours to several days to run, depending on your hardware configuration. As long as the run is making forward progress and new lines are being written to the `ALL.OUT` files, everything is probably fine.